

movemental

MOVEMENTAL FIELD GUIDE · VOLUME ONE

It Starts With Safety.

A field guide for building
your AI Organizational Guidebook.

BRAD BRISCO · ALAN HIRSCH · JOSHUA SHEPHERD

INTRODUCTION

Why we begin here.

AI in its present form is a massive organizational threat and a massive organizational opportunity. The reason we begin with Safety — other than that safety is always first, and never last — is that the threats are slightly more urgent than the opportunities. The threats have to be prevented before the search for opportunities begins.

Comprehensively building your organization's AI Guidebook is no guarantee of safety. People will still make mistakes. They will require training. Vendors will still pitch tools your Guidebook will tell you to refuse. New threats will arrive that your Guidebook does not yet anticipate. But not doing the work this Field Guide describes is a guarantee that an AI failure or breach of trust is inevitable. Because AI is already here, inside your organization, being used by your staff in ways no one has decided are acceptable. Our only choice is to respond to it wisely.

That begins with Safety.

How to use this Field Guide.

This guide walks you through the five layers of your organization's AI Guidebook — Statement, Policy, Context, Rules, and Response Plans — in the order they need to be built. Each layer is broken into a checklist of what a complete version contains. Your leadership team can work through this Field Guide together over two weeks and produce a complete AI Organizational Guidebook your board can ratify.

If you would rather have facilitation, Movemental's SafeStart engagement builds the Guidebook with you and delivers it inside a private organizational dashboard — retrievable, live, role-based, and integrated with everything that comes next on the Movemental Path. The two paths section later in this guide compares both options.

AUTHORS' NOTE

Where this work comes from, and what we are inviting you into.

AI will mirror and amplify the humanity it finds sitting before it in any given interaction.

That is the sentence beneath everything in this Field Guide, and the sentence beneath the company we have built. The question facing mission-driven organizations is not whether AI can be used for good. It is whether your organization's already-existing goodness — the convictions you hold, the people you have formed, the trust you have built, the credibility you have earned over decades — will be responsibly translated into the uninvited age of AI that has arrived inside your work whether you were ready for it or not.

Fragmented organizations — including the kind of fragmentation that results from a lack of a cohesive, organized response to AI — inherently produce AI slop, data risk, security vulnerabilities, and the gradual or sudden loss of hard-earned credibility and trust. Cohesion is the only defense. The Guidebook this Field Guide describes is one form that cohesion takes. Safety is where cohesion begins. But it does not end there, and that matters for why this Field Guide exists at all.

How we got here.

Brad and Alan have spent decades doing missional and ecclesial work — planting churches, training leaders, writing the frameworks that have shaped how a generation of practitioners think about movement, formation, and the recovery of mission. Send Network, Forge, The Forgotten Ways, 5Q, APEST, 100Movements, Movement Leaders Collective: that body of work predates AI, and its credibility predates AI.

Josh came to this work from inside it — years as a Methodist pastor, then as founder of a neomonastic, communal nonprofit where he and his wife lived alongside young adults in extended residency — and from inside that practice began building the digital infrastructure for what comes next.

We began collaborating in 2024 around a shared concern: that the present moment in AI would happen to mission-driven organizations rather than be navigated by them. That the organizations who have spent decades earning trust would be the ones who lose it fastest if no one named what was happening and built the infrastructure to respond. That the answer would not come from AI consulting firms, because the answer is not primarily about AI. The answer is about whether organizations whose work depends on human covenant can carry that covenant into an age of machine fluency without compromising it.

What you are reading is the first public release of work we have been developing together for more than a year. It is the on-ramp to the four-stage Movemental Path — Safety, Sandbox, Skills, Solutions — the methodology and infrastructure we have built for mission-driven organizations who want to navigate

this season wisely rather than reactively. The Field Guide you are reading is Safety, made portable, so you can begin even if you never engage us beyond this document.

What we are inviting you into.

Everything starts with Safety. It does not stop there. The Guidebook this Field Guide helps you build is the foundation for everything else — for the bounded exploration of AI in your actual work that comes next, for the formation of leaders who can steward AI wisely, for the eventual integration of AI into how your organization runs in a way that strengthens rather than erodes the trust your work depends on.

We are inviting you into a path. A wise and mature path through what AI is asking of organizations like yours. We have walked enough of it ourselves to know it works, and we have built the infrastructure to walk it with you if you want company. But the invitation comes first as a document you can read, a checklist you can work through, a Guidebook you can ratify, and a posture you can hold publicly. That much is yours, freely, whether we ever speak again or not.

CONTENTS

What you'll find here.

Introduction — Why we begin here	01
Authors' Note	02
Why Safety — the urgency, the cost	05
How this work gets done — process, roles, tension	07
The Data — by the numbers	12
The five layers of your AI Guidebook	14
01 AI Organizational Statement	15
02 AI Organizational Policy	16
03 AI Organizational Context	18
04 AI Organizational Rules	20
05 AI Organizational Response Plans	22
Self-assessment	24
Common mistakes	26
Two paths into Safety — SafeGuide vs. SafeStart	27
About Movemental	30
Citations	31

WHY SAFETY

Safety is always first. Never last.

Safety is the first stage of the Movemental Path for two reasons that work together. The first is structural: AI is now creating both new threats and new opportunities, and the threats are dire enough that they have to be addressed before any exploration of opportunities can responsibly occur. The second is timing: not addressing the threats is the surest way to ensure they happen. There is no neutral middle position. Inaction is a decision, and in this case it is the decision that produces the breach.

We are not fearmongers. We do not believe AI is uniquely evil or that mission-driven organizations should refuse it. We believe the opposite — that organizations whose work depends on trust have something distinctive to offer in this moment, and that AI responsibly used can free those organizations to do more of the work they exist to do. But the urgency of this first step is real, and the cost of skipping it is also real.

The urgent cost.

These are not predictions. They are the documented present, drawn from research published in the last twelve months. Each one is something that has already happened, at scale, across organizations like yours.

Voice cloning attacks rose 400 percent year over year in 2025. Voice cloning is now layered into business email compromise schemes targeting senior leaders. A pastor in Alabama, a development director in Florida, an executive director in Missouri, an institutional president in Nebraska — all have had their voices cloned and used to request wire transfers from staff. The IC3 documented \$893 million in adjusted losses tied to AI-enabled scams in 2025 alone, with \$352 million of that sustained by Americans 60 and over. A mid-sized nonprofit losing a six-figure wire transfer to a voice clone of its executive director is not a hypothetical. It is the average case in 2025.

Forty-three percent of Protestant churchgoers disagree with their pastor using AI for sermon preparation. Twenty-four percent disagree strongly. Sixty-two percent of pastors are concerned about non-disclosure of AI involvement. The trust gap is not theoretical — it is sitting in your pews, your donor base, your member directory, your student body. A single undisclosed AI-generated communication that gets discovered does damage that takes years to reverse.

Donor data is the most common breach vector for mission-driven organizations. The 2020 Blackbaud ransomware attack exposed donor records at more than 13,000 organizations and produced a multistate settlement exceeding \$49.5 million. The Identity Theft Resource Center documents a 50 percent increase in cyberattacks on nonprofits between 2020 and 2024. AI-tool data leakage is a new vector inside an old problem: every time a staff member pastes a donor list, a member directory, or a financial record into a consumer AI tool without enterprise terms, your organization's data is potentially training a model you do not control.

Forty-seven percent of nonprofits have no AI governance policy at all. Five percent of churches have an AI policy in any form. If your organization is in either of those majorities, the absence of policy is not

a position — it is exposure. When a journalist, a denomination, an accreditor, or a major donor asks what your organization's AI position is, "we haven't written one down" is not an answer they will accept.

The long-term cost.

The acute risks are real. The slower, structural cost is harder to see but more consequential over time. An organization without a shared Guidebook — without a single source of truth for what AI is, what it is for, what it is not for, and what it is not allowed to touch — cannot govern AI use. Not because the leadership team is negligent, but because AI is so new, so powerful, and so pervasive that informal norms cannot keep up with it. Every staff member becomes a unilateral decision-maker by default. Every vendor pitch lands against a leadership team with no shared frame to evaluate it. Every edge case produces a fresh argument.

The eventual outcome is one of two crises. Either a discrete incident — a leaked donor list, a voice-cloning fraud, a disclosed AI-generated sermon, a chatbot that represented the organization without authorization — produces a public trust event that the organization spends years recovering from. Or, more commonly, the organization experiences a gradual erosion of credibility as constituents, donors, and stakeholders notice the absence of a coherent position and quietly lose confidence. Neither outcome happens because leaders are stupid. They happen because the tool is so new and so pervasive that being thoughtful in retrospect is not enough.

Cohesion is the only defense. A shared Guidebook is the form cohesion takes.

Why we are responding the way we are.

This is the logic behind everything that follows. AI presents new threats and new opportunities. The threats are dire enough that they must be addressed before any exploration of opportunities can responsibly occur. Not addressing them is also a way of ensuring they happen. The urgency is real.

That is why we have built two paths into Safety, both of which produce the same AI Organizational Guidebook. SafeGuide is this Field Guide — free, comprehensive, and sufficient for any organization with the internal capacity to do the work in one to two months. SafeStart is the facilitated path — \$1,000, two weeks, with Movemental drafting your Guidebook against your specific context and delivering it inside a private organizational dashboard. The two paths section later in this document compares them in detail.

Whichever path you choose, choose now. The next section names what that work actually looks like — the roles, the process, and the tension this work surfaces inside the organizations you lead.

PROCESS & ROLES

How this work gets done.

Before the data, before the architecture of the Guidebook itself, a word on how this work actually moves through an organization. The framework that follows is only useful if your organization can run it. The most common reason Safety stalls is not a missing concept; it is a missing process.

This is adaptive work, not technical work.

The work of building your AI Organizational Guidebook is adaptive work, not technical work. The distinction matters. Technical problems have known solutions and identifiable experts; you call the right specialist and they solve it. Adaptive problems require the organization itself to learn, to weigh trade-offs, and to make decisions that no outside expert can adjudicate from outside.

AI governance for a mission-driven organization is unmistakably adaptive. No AI consultant can tell you what your theology says about AI in pastoral care. No vendor can tell you what your specific donors will tolerate. No board governance handbook can tell you what your particular leadership team needs to feel ownership of the outcome. The work has to be done by the organization, not delivered to it.

What adaptive work requires is a trustworthy process — one with clear roles, clear pacing, clear inputs, and a credible owner responsible for the integrity of the process rather than the rightness of the answer. The leader of the organization is not the person who should draft the Guidebook unilaterally. But the leader is responsible for the legitimacy and pace of the process by which it is drafted.

Visible structure is what makes the process trustworthy. People need to know what stage they are in, who is contributing, who decides, and when. The opposite — invisible, unstructured, ad-hoc deliberation — is what produces the fragmentation that leads to AI slop, breach, and the gradual loss of credibility named in the previous section. A good process can produce a defensible document. A bad process can produce a good document and still fail, because the organization will not live by it.

Who contributes. Who decides.

Two roles are non-negotiable. **Contributors** are the people whose work and perspective need to shape the substance of the Guidebook. **Deciders** are the people with authority to ratify what contributors produce. Most organizations conflate these roles, which is one of the most common reasons the process fails. Naming them separately, in writing, before drafting begins, is the first act of an adaptive process.

There is no single right configuration. What matters is that yours is named, that everyone in the process knows which role they hold, and that the deciding body has ratified the role design itself before the work begins. Below are the viable options across the organizational types we work with most.

FOR A CHURCH.

Contributors typically include the senior pastor, an executive or operations pastor, the communications lead, the finance or operations lead, an elder or board representative, and sometimes a technology or

data lead. A respected lay leader with relevant professional background can add ballast. Deciders are typically the senior pastor (who signs the Statement) and the elder board or governing body (which ratifies). Four to six contributors is workable; more than seven slows the process without strengthening the outcome.

FOR A NONPROFIT.

Contributors typically include the executive director, development director, program director, communications director, finance or operations director, and sometimes a board member or major donor with relevant expertise. Deciders are typically the executive director (signing) and the board (ratifying). Same numerical guidance applies.

FOR A THEOLOGICAL INSTITUTION, SEMINARY, OR CHRISTIAN COLLEGE.

Contributors typically include the president or senior administrator, the provost or academic dean, the communications lead, the chief information officer, the development lead, and a faculty representative from theology, ethics, or a relevant department. Deciders are typically the president and the board of trustees, with faculty senate consultation if institutional culture requires it. These environments often require longer ratification cycles; plan accordingly.

FOR A DENOMINATION, NETWORK, OR MOVEMENT.

Contributors typically include the network or denominational executive, a communications lead, a theological or doctrinal lead, and representatives from member organizations or churches. Deciders are typically the executive and the governing council. The added complexity is that the Guidebook may need to function both as the network's own internal Guidebook and as a template member organizations can adapt for themselves.

In every case: contributors are four to seven people from across functions; deciders are one to three people with executive authority plus the governing body for ratification; both roles are named in writing before any drafting begins. Surprises about who has decision authority, surfacing midway through the work, are what derail otherwise sound efforts.

What the SafeStart process looks like.

The SafeStart engagement runs on a two-week timeline with clearly defined steps. Each step is paced so the organization has time to think but no time to drift. Below is what the engagement looks like end to end.

ENGAGEMENT.

You sign the engagement letter at movemental.ai/safety/sign. Your private organizational dashboard is provisioned within 24 hours. Your contributor and decider roles are confirmed in writing before the kickoff Zoom.

ONE-HOUR KICKOFF ZOOM WITH STAKEHOLDERS AND JOSHUA SHEPHERD.

Held within the first three business days of engagement. The contributor team and the lead decider meet with Josh by video conference. The hour covers organizational context (mission, polity, current

AI usage if known), specific concerns the leadership team is carrying, contributors and deciders confirmed, organizational artifacts that should inform the drafting (existing policies, theological documents, board governance materials), and pacing expectations for the remainder of the engagement.

DRAFTING AND ASYNCHRONOUS REVIEW (WEEK ONE).

Movemental drafts your Guidebook across all five layers — Statement, Policy, Context, Rules, Response Plans — customized to your specific organizational context. The draft is published to your private dashboard within seven days of the kickoff. Contributors review asynchronously, leaving comments and proposed edits in the dashboard. This is where the substantive feedback happens.

ONE-HOUR REVISION ZOOM (WEEK TWO).

Held in the second week. The contributor team meets with Josh to walk through significant comments, resolve open questions, and lock in revisions. Movemental publishes the revised Guidebook to the dashboard within 48 hours of the revision Zoom.

INTERNAL RATIFICATION MEETING (WEEK TWO, END).

Your deciding team meets — internally, without Movemental in the room — to ratify the Guidebook through whatever governance process your organization requires. Josh is available by phone if the deciding team has questions during the meeting, but the ratification is yours. The Guidebook is yours when your decision body signs it.

PUBLICATION AND ROLLOUT (WEEK THREE).

The board ratification packet exports from the dashboard. The constituent-facing communication template is ready to deploy. Staff onboarding materials are available. The Guidebook is now live, retrievable, and integrated with everything that comes next on the Movemental Path.

Six steps, two weeks, one ratified Guidebook. The SafeGuide path mirrors this structure across one to two months, with your team holding all the roles Movemental holds in the facilitated version. Both paths produce the same outcome; the discipline of the process is what produces it.

What your own process needs to look like.

Some of the work in a SafeStart engagement is Movemental's. The rest is yours, and we cannot do it for you. Naming what is yours, up front, is part of the trust we are asking you to extend. Each of the items below is something your leadership team needs to carry, whether you engage SafeStart or take the SafeGuide path on your own.

PRE-ENGAGEMENT COMMUNICATION TO STAFF.

Before the kickoff (in SafeStart) or before drafting begins (in SafeGuide), the leadership team should communicate to staff that the organization is engaging in this work, why, and what to expect. Quiet, internal-only engagements produce suspicion. Openly named engagements produce buy-in. The transparency framework on the next page is the script for this communication.

DESIGNATION OF CONTRIBUTORS AND DECIDERS.

This happens within your governance structure, with whatever consultation your bylaws or polity require. It cannot be assigned by Movemental. The legitimacy of the process depends on these roles being filled by your decision, not ours.

CALENDAR TIME FOR CONTRIBUTORS TO ACTUALLY ENGAGE WITH THE DRAFTS.

The asynchronous review week is not a formality. Contributors who do not actually read the draft and submit substantive feedback are contributors in name only, and the document that results will be weaker for it. Block calendar time. Make participation expected, not optional.

PRAYER, DISCERNMENT, AND THEOLOGICAL REFLECTION.

For faith-based organizations, this is non-negotiable. A Statement that has not been prayed over by the leadership team will not carry the moral weight it needs to carry under pressure. The Care Boundaries document, in particular, requires theological reflection that staff time alone cannot produce. Build space for it. The two-week SafeStart timeline accommodates this, but only if you intentionally make room.

FINAL RATIFICATION MEETING THAT YOU OWN.

The decision to ratify must be made by your governance body, in your space, on your terms. Movemental is not in that room. The ratification is the organizational moment that converts a document into a covenant.

POST-RATIFICATION COMMUNICATION TO CONSTITUENTS.

The Guidebook is not finished when it is signed. It is finished when the people you serve know it exists. The constituent-facing communication is the move that converts internal governance into a public posture, and it is one of the highest-leverage moments in the engagement. Do not rush it. Do not skip it.

SETTING THE CADENCE FOR FUTURE REVIEW.

The Statement is reviewed annually. The Inventory is refreshed quarterly. The Incident Response Plan is updated after any actual incident. Your governance calendar needs to absorb these cadences before the engagement ends, or the Guidebook will go stale within twelve months.

Naming the tension before it names you.

Mission-driven organizations contain people with profoundly different views on AI. Some staff are enthusiastic, already deeply using tools, frustrated that the organization has not moved faster. Some are deeply skeptical, theologically or ethically concerned that AI use compromises the organization's witness. Some have not thought about it. Some have strong views they have not yet shared. The same range exists in your board, your major donors, your constituents.

The act of beginning Safety surfaces these differences. They were always there; the process makes them visible. That surfacing can either destroy the engagement or strengthen it, depending on how it is named. The single most important move is to be transparent and clear, up front, about exactly what is being done, what is not being done, and why.

THE TRANSPARENCY FRAMEWORK, ADAPTED TO YOUR CONTEXT, LOOKS SOMETHING LIKE THIS:

We are addressing AI Safety in our organization. AI is already affecting our work in specific ways — [name them: staff using ChatGPT for drafting, vendors selling us AI features inside existing software, donors asking what our position is, members hearing voice-cloning fraud stories from peer organizations].

If we do nothing, specific things become likely — [name them: voice cloning fraud targeting our staff and donors, undisclosed AI use in pastoral or constituent communication, data exposure through tools without enterprise terms, loss of credibility when we are asked about our position and have no answer].

We are choosing one specific process. Contributors are [named]. Deciders are [named]. Timeline is [two weeks for SafeStart, one to two months for SafeGuide]. Deliverable is the publication of our AI Organizational Guidebook — the Statement, Policy, Context, Rules, and Response Plans that govern how AI is and is not used in our work.

This step does not address every question we will eventually need to answer about AI. It does not address every specific tool we may or may not adopt. It does not address individual use cases your staff may want to explore. It does not address broader theological or ethical debates about AI in society. Those questions are real, and we will address them. They cannot be addressed responsibly until we have done this foundational work first. We are starting with Safety because the urgent threats AI poses to an unorganized organization must be prevented before broader exploration is responsible.

Why this framing works.

This framing accomplishes four things at once.

The skeptics know we are not handing over the keys. The boundary is the Guidebook, and the Guidebook contains Named Refusals. Skeptics often agree more vigorously with the refusals than the enthusiasts do.

The enthusiasts know we are not banning anything. The Policy sanctions categories of use; the Rules clarify what is allowed; the engagement clears the ground so use can be governed rather than ad hoc. Enthusiasts often find the Guidebook releases them to do work they were already doing more safely.

The undecided know there is a coherent process they do not have to participate in directly. The decision is being made by named contributors and a named deciding body. The undecided can trust the process even if they cannot yet form a view on the substance.

The deciding body knows exactly what they are being asked to ratify. The scope is named. The exclusions are named. The deliverable is named. Ratification becomes a discrete act rather than an open-ended one.

The tension is real. Named, it becomes the energy that produces a Guidebook your organization will live by. Unnamed, it becomes the friction that prevents the work from happening at all.


THE DATA

What the research shows.

Six independent studies, released in the last twelve months, converge on the same pattern. Adoption is broad. Capability gain is rare. Governance is largely absent. The harm landscape is documented.

The capability gap.

Across every sector studied, the cohort that turns AI adoption into measurable capability sits between five and seven percent. The rest adopt without governing, without redesigning workflows, and without producing material value.



The trust gap.

Constituent trust is conditional on disclosure. Most organizations have no Disclosure Rules.



81%

of nonprofits describe their AI use as individual and ad hoc.

VIRTUOUS · 2026 · N=346

The harm landscape.

These are not projections. They are documented losses and incident counts from the FBI's 2025 IC3 Annual Report, the AI Incident Database, and supporting industry research.

\$893M

in adjusted losses tied specifically to AI-enabled scams in 2025.

FBI INTERNET CRIME
REPORT · 2025

400%

year-over-year increase in voice cloning fraud attacks in 2025.

INDUSTRY TRACKING / FBI ·
2025

22,364

complaints filed with the FBI carrying an AI-related descriptor in 2025.

FBI IC3 ANNUAL REPORT ·
2025

The governance gap.

AI governance is now an expansion of the board's fiduciary duties of Care, Loyalty, and Obedience. Most boards have not yet acted on that.

47%

of nonprofits have no AI governance policy at all.

VIRTUOUS · 2026 · N=346

5%

of churches have an established AI policy.

BARNA / PUSHPAY · 2026 ·
N=1,306

64%

of church leaders believe an AI use policy is important. Most do not have one.

BARNA / PUSHPAY · 2026 ·
N=1,306

THE ARCHITECTURE

The five layers of your AI Guidebook.

Your organization's AI Guidebook is a single document with five layers. Each layer depends on the one before it. You cannot write Policy without a Statement that grounds it. You cannot write Rules without knowing the Context. You cannot draft Response Plans without Rules that name what shouldn't happen in the first place. The layers progress from abstract to concrete, from identity to action.

01	Statement	What we believe about AI in relation to our mission.
02	Policy	What we will do and what we refuse to do, operationally.
03	Context	What is actually true in our environment right now.
04	Rules	What governs AI use in specific domains.
05	Response Plans	What we do when something goes wrong.

The rest of this guide walks through each layer in order, with a checklist of what a complete version contains.

LAYER 01

AI Organizational Statement.

What we believe about AI in relation to our mission.

The Statement is the highest-level layer of your Guidebook. It names what your organization believes about AI: the values that will hold under pressure, the convictions that ground every operational decision that follows, and the connection between your AI use and the mission your organization already exists to serve. Two pages is enough. Longer and it stops being readable in a single sitting; shorter and it has not done the theological or ethical work the moment requires.

The Statement is what makes every layer that follows coherent rather than arbitrary. Staff who read the Acceptable Use Policy need to recognize the principles behind it. Boards reviewing the Vendor Inventory need to know what the inventory is for. Stakeholders asking about your position need a single document that names where the organization stands. The Statement is that document. It is signed by senior leadership and ratified by the board.

A COMPLETE STATEMENT CONTAINS:

- ☐ A mission alignment clause naming the relational and mission commitments AI use must serve rather than displace
- ☐ A posture statement naming the organization's stance toward both fear and uncritical enthusiasm about AI
- ☐ A category-of-care clause naming the kinds of work where the organization will not let efficiency be the deciding criterion
- ☐ A theological or ethical grounding paragraph (for faith-based and mission-driven organizations) connecting AI posture to existing convictions
- ☐ A scope clause naming which parts of the organization the Statement applies to (staff, volunteers, contractors, vendors, board)
- ☐ A review cadence (typically annual) naming when and how the Statement will be revisited
- ☐ Signatures from senior leadership
- ☐ Board or governing body ratification, with date of vote

What is at stake without it: every layer below reads as arbitrary administrative compliance. The first hard case produces an internal argument with no grounding to settle it.

LAYER 02

AI Organizational Policy.

What we will do, and what we refuse to do, operationally.

The Policy translates the Statement into operational position. Where the Statement names what the organization believes, the Policy names what the organization does — the categories of AI use that are sanctioned, the categories that are prohibited, the categories that require disclosure, and the categories that the organization commits to refuse on principle regardless of pressure or demonstration.

The Policy is the document a new staff member reads in their first week. It is also the document a journalist or major donor reads if they ask what your organization's AI position is. It must serve both audiences: concrete enough for staff to apply, public enough for stakeholders to evaluate.

The Policy contains a section called **Named Refusals** — the public commitment to refuse specific applications of AI on principle. The Named Refusals are not predictions about what your organization will avoid; they are covenants. They can be cited, defended, and referenced when a vendor pitches the use anyway.

A COMPLETE POLICY CONTAINS:

- ☐ A sanctioned-use clause naming the categories of work where AI is welcomed (drafting, research, summarization, scheduling, formatting, analysis)
- ☐ A constrained-use clause naming the categories where AI is permitted only with conditions (disclosure, supervision, editorial review)
- ☐ A prohibited-use clause naming the categories where AI is not permitted regardless of efficiency or demonstration
- ☐ A consequences clause naming what happens when the Policy is violated
- ☐ A scope clause naming who the Policy applies to (employees, volunteers, contractors, vendors using AI on the organization's behalf)
- ☐ A delegation clause naming who has authority to interpret edge cases and who can grant exceptions

THE NAMED REFUSALS SECTION OF THE POLICY CONTAINS:

- ☐ A refusal of AI use in eligibility determinations for services, programs, or membership
- ☐ A refusal of AI-driven behavioral surveillance of constituents that differentiates how they are treated
- ☐ A refusal of AI-generated personal voice without disclosure (appeals, sympathy letters, pastoral correspondence written in someone's voice)

- ☐ A refusal of AI substitution in pastoral, clinical, counseling, or formative care relationships that require irreducible human presence
- ☐ A refusal of AI deepfakes, voice cloning, and undisclosed chatbots representing the organization
- ☐ Additional context-specific refusals your organization commits to (institutional, theological, programmatic)

What is at stake without it: shadow use. Staff use whatever is convenient because the organization has not named what is acceptable. By the time the question arrives, the practice is already entrenched.

LAYER 03

AI Organizational Context.

What is actually true in our environment right now.

The Context layer converts unknowns into knowns. Before your organization can write Rules that govern AI use, it has to know what AI is currently in use, what data your organization holds, and what vendor terms govern the tools your staff are touching. The Context layer is the situational ground on which the Rules will stand.

Two documents constitute the Context layer: the **Vendor and Tool Inventory** and the **Data Classification**. Both are living documents, reviewed quarterly, updated as the environment changes.

A COMPLETE VENDOR AND TOOL INVENTORY CONTAINS:

- ☐ A list of every AI tool currently in use across the organization, including consumer tools (ChatGPT, Claude, Gemini), enterprise tools (Microsoft Copilot, Google Workspace AI), embedded AI features in existing software, and tools used by volunteers or contractors
- ☐ The names of the staff using each tool, by name or by role
- ☐ The categories of work each tool is being used for
- ☐ The data each tool is permitted to touch under your Data Handling Rules
- ☐ The vendor terms governing each tool: training rights on inputs, data residency, retention policies, enterprise vs. consumer tier
- ☐ A procurement gate naming the review process any new AI tool must pass before entering the inventory
- ☐ A review cadence (typically quarterly) for keeping the inventory current

A COMPLETE DATA CLASSIFICATION CONTAINS:

- ☐ Defined tiers of data sensitivity (typically: public, internal, sensitive, confidential)
- ☐ Examples within each tier specific to your organization (member directories, donor records, pastoral notes, financial details, employee data, beneficiary information, clinical notes)
- ☐ A clear statement, for each tier, of which categories of AI tool may process data in that tier
- ☐ An explicit list of data categories that may not be processed by any external AI tool under any circumstance
- ☐ An owner for each data tier (who decides edge cases, who reviews exceptions)

What is at stake without it: the organization cannot answer basic questions about its own AI exposure, which means it also cannot govern it.

LAYER 04

AI Organizational Rules.

What governs AI use in specific domains.

The Rules layer is where Policy meets reality. The Policy says what the organization does. The Rules say how, specifically, in three domains where the operational stakes are highest: data, disclosure, and care. Each rule is a substantive document. Together they cover the operational surface area where most AI incidents in mission-driven organizations actually originate.

A COMPLETE DATA HANDLING RULES DOCUMENT CONTAINS:

- ☐ The mapping of data tiers from the Data Classification to permitted AI tools and contexts
- ☐ A clear answer to the question “can I put this in this tool?” for the most common categories of work staff perform
- ☐ A set of pre-approved workflows that comply with the Rules, so staff have a default path that does not require case-by-case judgment
- ☐ A procurement-review process for new tools, tied to the Vendor Inventory
- ☐ An exception process for cases not covered, naming who decides and how the decision is recorded
- ☐ A breach-reporting process tied to the Incident Response Plan

A COMPLETE DISCLOSURE & ATTRIBUTION RULES DOCUMENT CONTAINS:

- ☐ The categories of work where AI assistance is unremarkable and disclosure is not required (formatting, scheduling, summarizing internal documents)
- ☐ The categories where AI involvement must be disclosed (personal communication, pastoral correspondence, fundraising appeals in a leader's voice, public-facing content)
- ☐ The specific language or marker used for disclosure (e.g., “AI-assisted under editorial review by [name]”)
- ☐ A tiered disclosure framework if your organization distinguishes between levels of AI involvement
- ☐ A statement about attribution when AI generates substantive content the organization publishes
- ☐ Edge cases and how to handle them (AI used in sermon preparation, AI used in donor segmentation, AI used in board materials)

A COMPLETE CARE BOUNDARIES DOCUMENT CONTAINS:

- ☐ The specific categories of work where the presence of a human being is the irreducible input (initial pastoral conversation after grief or trauma, crisis-line first contact, clinical counseling intake, end-of-life conversation support, programmatic eligibility determinations, disciplinary conversations)
- ☐ The secondary cases where AI may prepare or support a human but may not substitute for one (sympathy letters, intake forms, follow-up communication)
- ☐ Named workflows for each secondary case making the human role explicit
- ☐ A theological, clinical, or ethical rationale for each boundary (drawn from the Statement)
- ☐ A review process for new categories of work that emerge as AI capabilities expand

What is at stake without it: a slow displacement of human presence in exactly the places where the organization's reason to exist resides. Disclosure violations that erode trust before anyone notices. Data exposure that becomes a breach.

LAYER 05

AI Organizational Response Plans.

What we do when something goes wrong.

The Response Plans layer is what your organization activates when normal operation breaks down. Where the Rules describe steady-state behavior, the Response Plans describe what happens when steady state fails: when AI output causes harm, when AI-enabled fraud targets staff or constituents, when AI-generated content is published under the organization's name without authorization, when data is exposed through a tool that did not meet the Rules.

The **Incident Response Plan** is the primary Response Plan. Additional Response Plans may emerge over time — a Migration Plan for retiring deprecated tools, a Renewal Plan for the annual Statement and Policy review — but the Incident Response Plan is what every organization must have from day one.

A COMPLETE INCIDENT RESPONSE PLAN CONTAINS:

- ☐ A clear definition of what counts as an incident (harmful or inaccurate AI output published under the organization's name; AI-enabled impersonation of staff, leaders, or constituents; voice cloning or deepfake attacks; data exposure through a non-compliant tool; significant Policy violations by staff)
- ☐ Named owners for each incident category, with backup owners named for after-hours and weekend coverage
- ☐ An escalation path naming when an incident moves from staff-level to senior-leadership to board-level
- ☐ Pre-arranged verification protocols for senior staff (code words, callback procedures, dual-authorization for fund transfers above a threshold)
- ☐ A constituent-communication template ready to deploy when an incident affects the people the organization serves
- ☐ A law-enforcement coordination workflow including the IC3 reporting process at ic3.gov for AI-enabled fraud
- ☐ A post-incident review process to update the Rules, the Inventory, and the Plan itself based on what was learned
- ☐ A communication plan for what gets reported to the board and on what timeline

OTHER RESPONSE PLANS THE GUIDEBOOK MAY INCLUDE OVER TIME:

- ☐ A Migration Plan for retiring AI tools that fall out of compliance or are deprecated by vendors
- ☐ A Renewal Plan naming the cadence for annual Statement and Policy review, quarterly Inventory updates, and trigger-based reviews after any incident

- ☐ A Crisis Communication Plan for high-severity incidents that reach public attention

What is at stake without it: improvised response during an active incident, slower public communication, larger financial loss, and a longer trust-recovery period.

TAKE STOCK

Self-assessment.

Five questions a leadership team can answer together in thirty minutes. One question per layer of your AI Guidebook. Each answer is yes, partial, or no. The pattern of answers tells you where to begin. The assessment is most useful when taken with the senior leadership team in the room — the disagreement is the most useful output.

LAYER 01 — STATEMENT

Does your board or governing body have a written, ratified position on AI?

A strong answer is a signed Statement or equivalent document of record. A weak answer is “we have talked about it.” Forvis Mazars frames AI governance as a fiduciary expansion of the board’s duties of Care, Loyalty, and Obedience; “we have talked about it” does not satisfy that expansion.

LAYER 02 — POLICY

Have you named, in writing, the specific applications of AI your organization commits to refuse on principle?

A strong answer is a Named Refusals section signed by senior leadership. A weak answer is an internal sense of unease that has not been put on paper. The refusals you can name are the convictions you can defend under pressure.

LAYER 03 — CONTEXT

Can your senior leadership name every AI tool currently in use across your staff, with what data each tool is touching?

A strong answer is a written Vendor and Tool Inventory paired with a Data Classification. A weak answer is a description of what tools have been “approved” without verification of actual use. The Virtuous data places 81 percent of nonprofits in the second category.

LAYER 04 — RULES

Do you have written Rules that tell a staff member, before they have to guess, what data may go into which tool and when AI involvement must be disclosed?

A strong answer is a Data Handling Rules document, a Disclosure & Attribution Rules document, and a Care Boundaries document — each ratified and in active use. A weak answer is reliance on individual staff discretion.

LAYER 05 — RESPONSE PLANS

If a senior leader were impersonated tomorrow by an AI voice clone or deepfake targeting your staff or donors, would your team know what to do in the first hour?

A strong answer is an Incident Response Plan with named owners, escalation paths, pre-arranged verification protocols, and a constituent-communication template ready to deploy. A weak answer is “we would handle it as it came up.”

How to read your pattern.

No to three or more layers. Begin Safety now. Your Guidebook is largely unwritten and your organization is exposed to documented harms. SafeStart is the lowest-cost facilitated path; SafeGuide is sufficient if you have internal capacity and alignment.

Yes or partial on most layers. Begin Safety this quarter. Finishing the Guidebook is what gets you into the five-to-seven percent cohort of organizations whose AI use is governed rather than reactive.

Yes on all five. You are likely already in the high-performer cohort. The question facing you is sustainability: revisiting the Statement annually, refreshing the Inventory quarterly, updating the Incident Response Plan after any actual incident.

WHAT GOES WRONG

Common mistakes.

Three patterns of failure recur across the organizations we work with. Each one is preventable. Each one is endemic.

The unilateral staff member.

One person, often a senior one, makes AI decisions for the whole organization without authority or process. The decisions are usually defensible in isolation but produce a precedent the organization has not consented to. The fix is governance, not personnel: until there is a Statement and a Policy and a Rules layer, every staff member is making unilateral decisions by default.

The technology-first trap.

The organization picks a tool before deciding what they want it to do or what they want it not to do. The tool then dictates the policy by being the thing the policy has to accommodate. The fix is sequence: Statement before Policy, Policy before Rules, Rules before tool selection.

The fiduciary blind spot.

The board has not engaged AI governance because no one has framed it as a fiduciary question. Forvis Mazars now frames it that way. The fix is a single agenda item at the next board meeting, framed by the Statement and the Policy and the Vendor Inventory, with a vote of record.

YOUR TWO PATHS

Two paths into Safety.

There are two paths to a complete AI Organizational Guidebook. Both produce the same five layers and the same ratifiable documents. The difference is who guides you through the work, what the finished artifact looks like, and what kind of organizational discipline the path requires from you. We are not trying to discourage the free route. We want you to see both paths honestly.

How the two paths compare.

	SafeGuide SELF-DIRECTED	SafeStart FACILITATED BY MOVEMENTAL
COST	Free.	\$1,000.
TIMELINE	One to two months, depending on internal capacity and alignment.	Two weeks. Fixed.
DIFFICULTY	High. Requires roughly twenty hours of senior leadership work, sustained focus, and the discipline to finish what you start.	Low. We do the drafting; your team reviews, revises, and ratifies inside the dashboard.
WHO DRAFTS	Your team writes the Guidebook from the framework, with no outside facilitation.	Movemental drafts the Guidebook customized to your organization. You revise and ratify.
FINISHED ARTIFACT	Whatever format your team produces (Word, Google Doc, PDF, Notion). Quality varies with your team's capacity.	Publishable internal documentation platform plus print-quality PDF plus rollout supports and messaging for each Guidebook area. No extra charge.
LIKELY CHALLENGES	Internal disagreement that has no third-party adjudication. Pace slowing without external structure. Incomplete sections that never get finished. Format inconsistencies that make the Guidebook hard to update later.	None at the Guidebook level. Some organizations underestimate the post-engagement work of operationalizing the Guidebook — which Sandbox addresses.

SafeGuide

SELF-DIRECTED

SafeStart

FACILITATED BY MOVEMENTAL

BEST FIT

Leadership teams with time, alignment, internal capacity, and a discipline of finishing what they start.

Leadership teams that would benefit from external pacing, third-party adjudication, and a finished artifact that is alive rather than filed away.

What SafeStart actually delivers.

When we say SafeStart delivers the Guidebook better than the SafeGuide path in every measurable way, we mean specifically these things. Each one is included in the \$1,000 fee. No upsell. No add-on tier.

A PUBLISHABLE INTERNAL DOCUMENTATION PLATFORM.

Your Guidebook lives inside a private organizational dashboard built for the engagement. Staff search by question and get the relevant clause. New hires open the dashboard on day one and see the organization's full AI posture. The board sees role-appropriate views. The dashboard updates in place when the Guidebook is revised.

A PRINT-QUALITY PDF OF YOUR FINISHED GUIDEBOOK.

Typeset and exportable at any time, for board packets, denomination filings, accreditor submissions, or whatever your existing governance documentation expects.

SUPPORTS AND MESSAGING FOR EACH GUIDEBOOK AREA.

Pre-built rollout materials for each layer of the Guidebook: a board ratification packet, staff announcement templates, a constituent-facing communication draft, FAQ for the people you serve, and an incident-readiness briefing for your senior leadership. The materials are yours to adapt and deploy.

INTEGRATION WITH THE REST OF THE MOVEMENTAL PATH.

When your organization moves to Sandbox, the Use Case Portfolio is adjudicated against your Guidebook's Rules automatically. When your team enters Skills, the formation curriculum is anchored to the Statement and Policy you have already ratified. The Guidebook is not a one-time artifact — it is the foundation everything else builds on.

If you do this on your own.

If you read this Field Guide and decide to build your Guidebook with your own team, we hope it serves you. The framework is genuinely sufficient. The trust dividend your Guidebook exists to produce arrives the moment your board ratifies it, regardless of who drafted it. The risk of SafeGuide is not that the framework is missing something — it is that the discipline of finishing the work without external

pacing is harder than most leadership teams expect. Some organizations do it well. Others start, stall, and end up with three layers drafted and two unfinished sitting in a folder no one opens.

If you want company on the walk, that is what SafeStart is for. We have walked in the shoes of the leaders we serve long before we started building digital infrastructure and AI platforms. The facilitation is not consulting overlay — it is practitioners helping practitioners move at the pace that produces a finished, defensible, durable Guidebook your organization can stand behind.

WHO WE ARE

About Movemental.

Movemental is the company that built the path that walks mission-driven organizations through AI governance, formation, and deployment without compromising the trust their work depends on. We are not an AI consulting firm. We are a small team of practitioners and engineers who put together the framework, the infrastructure, and the facilitated engagements that make the four-stage Movemental Path — Safety, Sandbox, Skills, Solutions — actually runnable inside an organization on a real Tuesday afternoon.

The four-stage Movemental Path.

Safety	What this Field Guide is about. Two weeks, the AI Organizational Guidebook with its five layers. The foundation everything else is built on.
Sandbox	The next stage. Roughly ten hours of in-person training across four to six weeks. Your team explores AI against real organizational work, using a library of pre-tested use cases for churches, nonprofits, and theological institutions. The output is a Use Case Portfolio adjudicated as green, yellow, or red; a Discernment Memo for your board; and a Sandbox Readiness Assessment that names whether Skills is the right next step.
Skills	The eight-week formation cohort that produces leaders capable of stewarding AI inside their organizations — not staff who have learned to use a tool. The capacities are discernment, authorship, and stewardship.
Solutions	Where AI is integrated into how your organization actually runs, owned by formed people, governed by your working Guidebook. Scoped per engagement.

How to start a conversation.

If you read the Field Guide and want to talk about your organization, email Joshua Shepherd at josh@movemental.ai or visit movemental.ai/contact. We respond personally to every inbound message.

If you decide to do the work yourselves, that is a good outcome. The document is freely shareable in its entirety with attribution; passing it to a peer at another organization is welcomed.

SOURCES

Citations.

- [1] Virtuous and Fundraising.AI. *The 2026 Nonprofit AI Adoption Report*. Released February 16, 2026. Benchmark survey of 346 nonprofit organizations. Findings: 92 percent of nonprofits use AI in some capacity; 7 percent report major improvements in organizational capability; 81 percent report using AI individually and on an ad hoc basis; 47 percent have no AI governance policy. virtuous.org/resource/the-2026-nonprofit-ai-adoption-report-download.

- [2] Barna Group and Pushpay. *Technology for Missional Impact: State of Church Tech 2026*. Released March 9, 2026. Survey of 1,306 U.S. church leaders. Findings: 60 percent of church leaders personally use AI at least once a month; 33 percent report their church is using AI in operations; 64 percent believe an AI use policy is important; 5 percent of churches have an established AI policy.

- [3] Lifeway Research. *Pastors' Views on Artificial Intelligence*. Released April 21, 2026. Phone survey of 1,003 Protestant senior pastors. Findings: 42 percent of U.S. Protestant pastors are using AI for ministry; 62 percent are concerned about non-disclosure. Companion churchgoer survey: 43 percent of Protestant churchgoers disagree with their pastor using AI for sermon preparation, 24 percent strongly.

- [4] McKinsey and Company. *The State of AI in 2025*. November 2025. Online survey, 1,993 respondents across 105 nations. Findings: 88 percent organizational AI adoption; 6 percent qualify as AI high performers. Boston Consulting Group. *The Widening AI Value Gap*. September 2025. 1,250 firms; 60 percent generate no material value from AI. MIT NANDA Initiative. *The GenAI Divide*. July 2025. 95 percent of enterprise generative AI pilots fail to generate measurable financial returns.

- [5] Forvis Mazars. *AI Governance for Nonprofit Boards*. By Karen Bodach and Dan Prater. February 13, 2026.

- [6] FBI Internet Crime Complaint Center. *2025 IC3 Annual Report*. Published 2026. Findings: 22,364 complaints with an AI-related descriptor in 2025; \$893 million in adjusted losses tied to AI-related scams; \$352 million sustained by Americans 60 and over. Industry tracking sources: voice cloning fraud increased more than 400 percent in 2025.

- [7] AI Incident Database Incident 1355: multi-state pattern of AI impersonation of pastors in Alabama, Florida, Missouri, Nebraska, New York, and a Philippines megachurch. incidentdatabase.ai/cite/1355. Vatican Dicastery for Communication reporting on Pope Leo XIV deepfakes.

- [8] Identity Theft Resource Center: 50 percent increase in cyberattacks on nonprofits between 2020 and 2024. The 2020 Blackbaud ransomware attack exposed donor data at more than 13,000 organizations; multistate settlement exceeded \$49.5 million plus a separate \$6.75 million California settlement in 2025.

- [9] Nonprofit Quarterly. "How Nonprofits Can Resist the AI Efficiency Trap." October 28, 2025. Council of Nonprofits. *Board Roles and Responsibilities*, 2025.

movemental

COLOPHON

It Starts With Safety.

A Field Guide for Building Your AI Organizational Guidebook.

EDITION 1.1 · MOVEMENTAL · 2026

Brad Brisco · Alan Hirsch · Joshua Shepherd

© 2026 Movemental. This document may be freely shared in its entirety with attribution. For permission to excerpt, contact josh@movemental.ai.

movemental.ai