

Pre-Sandbox AI Assessment

Full question & answer catalog (code-backed sources)

Generated: 2026-05-12

This document consolidates every organizational assessment item in this repository that informs Safety and pre-Sandbox readiness, including Likert scales and multiple-choice option text exactly as implemented in application code. Submissions for the Path Integrity Diagnostic are stored in Postgres (`integrity_diagnostic_submissions`), but the question bank itself is defined only in TypeScript — not in the generic `assessment_questions` table.

Where the lists live

In this repository (canonical)

- Safety Self-Assessment (7) — /start-with-safety — src/components/sections-mock/start-with-safety/safety-self-assessment.tsx
- Path Integrity Diagnostic (22 questions) — /assess — src/lib/integrity-diagnostic/questions.ts (version 2026-05-02)
- SSSS Path Integrity diagnostic (18 items) — /api/assess/ssss-integrity — src/lib/ssss-integrity-assessment/items.ts + options.ts (version ssss-i-v1)

Supabase (Movemental project)

The `assessment_questions` table holds hundreds of rows for other instrument types (mDNA, APEST, organizational maturity, etc.) keyed by `assessments.slug` — not the public web Integrity Diagnostic or Safety Self-Assessment banks above. Query `assessments.assessment_type` for `mdna` | `apest` | `organizational` | `personality` | `current_reality_map`.

Part A — Safety Self-Assessment

Public route: /start-with-safety. Four ordered choices per question (stored client-side only; nothing is POSTed).

Answer scale (all seven questions)

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 01 · tool-inventory

Do you know which AI tools your staff are already using?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 02 · written-guidance

Do you have written guidance for acceptable and unacceptable use?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 03 · data-boundaries

Have you defined what information should never enter AI tools?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 04 · human-review

Do leaders agree on when human review is required?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 05 · sensitive-handling

Do staff know how to handle confidential, pastoral, donor, student, or client information?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 06 · use-case-review

Do you have a process for reviewing new AI use cases?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Question 07 · articulation

Can your team explain why your AI boundaries exist?

Choices

1. Yes
2. Mostly
3. Partially
4. Not yet

Part B — Path Integrity Diagnostic (full bank)

Version 2026-05-02. Submitted answers: array of 22 integers (0–3) in question order via POST /api/assess.

Subset: "pre-Sandbox truth" dimensions only (01–03)

Dimensions 01–03 (Path integrity, Posture clarity, Refusal capacity) contain twelve questions that do not ask about sandbox environments. Dimension 04 (Sandbox truth) and later dimensions still concern pre-production discipline but explicitly reference sandboxes.

Dimension 01 — Path integrity (sequence-integrity)

Are we building solutions before we have established safety?

Q1 (1/4 in this dimension)

Has your organization established an internal AI safety policy?

Answer options (indices 0–3 for API)

1. Yes, fully implemented
2. In draft form
3. Under discussion
4. Not yet considered

Q2 (2/4 in this dimension)

Do you know which third-party tools may train on your organizational data?

Answer options (indices 0–3 for API)

1. Yes, we have full visibility
2. We know about official tools, not shadow IT
3. We assume they all do
4. We do not know

Q3 (3/4 in this dimension)

Before scaling an AI pilot, do you require explicit approval tied to documented safety readiness?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q4 (4/4 in this dimension)

Do you maintain a written record of workflows where AI use is categorically disallowed?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Dimension 02 — Posture clarity (posture-clarity)

Do our people know what we believe about AI, or just what tools we bought?

Q5 (1/4 in this dimension)

If asked, could most staff articulate your organization's ethical or theological stance on AI?

Answer options (indices 0–3 for API)

1. Yes, consistently
2. Some could
3. Only senior leadership
4. No

Q6 (2/4 in this dimension)

Is your AI approach documented separately from general IT strategy?

Answer options (indices 0–3 for API)

1. Yes, it has its own framework
2. It is a subset of IT
3. They are one and the same
4. We have neither

Q7 (3/4 in this dimension)

Is there a single narrative leaders use when explaining AI posture to boards, donors, or members?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q8 (4/4 in this dimension)

Do volunteers and contractors receive practical guidance aligned with staff AI posture?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Dimension 03 — Refusal capacity (refusal-capacity)

Do we have criteria for saying no?

Q9 (1/4 in this dimension)

Have you explicitly decided not to use AI for at least one sensitive operational or ministerial task?

Answer options (indices 0–3 for API)

1. Yes, multiple times
2. Yes, once or twice
3. We are open to anything that helps
4. We haven't evaluated enough to say no

Q10 (2/4 in this dimension)

Does your tool evaluation include theological or missional constraints, not only utility?

Answer options (indices 0–3 for API)

1. Yes, as a primary filter
2. Yes, as a secondary consideration
3. We focus mainly on utility
4. No

Q11 (3/4 in this dimension)

Is there a documented path for staff to challenge or pause an AI deployment they believe is unsafe?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q12 (4/4 in this dimension)

Do you review vendor claims about AI features with the same skepticism you apply to security claims?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Part B (continued) — remaining dimensions

Dimension 04 — Sandbox truth (sandbox-truth)

Are we experimenting in contained spaces, or risking real data?

Q13 (1/4 in this dimension)

Are you using isolated environments (sandboxes) to test AI tools before production use?

Answer options (indices 0–3 for API)

1. Yes, across the organization
2. In specific departments
3. Informally
4. No, we test in production

Q14 (2/4 in this dimension)

Before copying production-like data into a sandbox, do you review de-identification or synthetic substitutes?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q15 (3/4 in this dimension)

Do sandbox experiments have named owners and time-boxed sunset dates?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q16 (4/4 in this dimension)

After a sandbox sprint, do you archive or delete artifacts according to policy?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Dimension 05 — Formation, not training (formation-not-training)

Are we forming our people to use AI, or just teaching them prompts?

Q17 (1/3 in this dimension)

Does AI-related learning emphasize formation, discernment, and critical thinking—not only prompts?

Answer options (indices 0–3 for API)

1. Yes, deeply integrated
2. We address it somewhat
3. Mostly focused on prompt engineering
4. We don't do AI training yet

Q18 (2/3 in this dimension)

Do managers coach teams on situations where reaching for AI first would be unwise?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q19 (3/3 in this dimension)

When misuse occurs, do you treat it as a formation conversation—not only a policy strike?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Dimension 06 — Solution restraint (solution-restraint)

Are we trying to solve every problem with AI, or only the right ones?

Q20 (1/3 in this dimension)

Are AI initiatives mapped to documented organizational bottlenecks—not tool catalogs alone?

Answer options (indices 0–3 for API)

1. Rigorous mapping
2. Rough alignment
3. Adopting available features
4. Reactive adoption

Q21 (2/3 in this dimension)

Do you limit concurrent AI experiments so operational attention and governance keep pace?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Q22 (3/3 in this dimension)

At least annually, do you review AI initiatives against measurable mission outcomes?

Answer options (indices 0–3 for API)

1. Strongly yes
2. Somewhat
3. Rarely
4. No

Part C — SSSS Path Integrity diagnostic (Likert)

Version ssss-i-v1. Each item uses the same five-point truth scale (scores map to positions 0–4). Optional metadata: audience context and email — see POST /api/assess/ssss-integrity.

Likert labels (positions 0–4)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Audience context options (metadata, not scored)

- movement_leader — Movement leader, author, or teacher
- church_ministry — Church or ministry
- nonprofit_mission — Nonprofit or mission-driven organization
- institution_other — Institution, network, or seminary

Item bank (all stages)

Items with stage "Safety" are the strict pre-Sandbox stage gate in the four-stage Movemental Path framing. Later stages include Sandbox, Skills, Solutions, and Cross.

Item 01 — Q01 · Safety

category: boundaries_and_authority · weight: 1

We can state, without notes, what is forbidden in external-facing work and what requires human review.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 02 — Q02 · Safety

category: governance_artifact · weight: 1

We have a published map of decision rights—not only informal habit.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 03 — Q03 · Safety

category: conviction_lines · weight: 1

Our deepest convictions (theological or ethical) are explicit enough to say no to plausible shortcuts.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 04 — Q04 · Safety

category: operational_spread · weight: 1

Data sensitivity tiers and escalation paths are understood across departments, not only legal/IT.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 05 — Q05 · Sandbox

category: learning_artifact · weight: 1

We keep a shared dated log of experiments, surprises, and failures—not only success stories.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 06 — Q06 · Sandbox

category: environment_compliance · weight: 1

Our experiments run in environments that comply with Safety, not primarily as shadow individual use.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 07 — Q07 · Sandbox

category: portfolio_discipline · weight: 1

We have a prioritized use-case portfolio screened by governance constraints.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 08 — Q08 · Sandbox

category: evidence_quality · weight: 1

We can point to evidence of what “good” and “not us” look like for our voice—not only opinions.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 09 — Q09 · Skills

category: distributed_judgment · weight: 1

Mid-level staff can describe good AI-assisted work without reading policy verbatim.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 10 — Q10 · Skills

category: culture_of_correction · weight: 1

We see public self-correction when outputs drift (voice, facts, ethics).

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 11 — Q11 · Skills

category: verification_norms · weight: 1

Verification habits are social norm, not heroics by one reviewer.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 12 — Q12 · Skills

category: formation_vs_training · weight: 1

Training time is spent on judgment, not only buttonology.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 13 — Q13 · Solutions

category: workflow_infrastructure · weight: 1

We deploy workflows with clear owners, gates, and failure modes—not tool brands as substitutes for design.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 14 — Q14 · Solutions

category: procurement_gates · weight: 1

Procurement conversations are shortened by pre-baked constraints and graduated use cases.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 15 — Q15 · Solutions

category: measurement_legibility · weight: 1

We measure workflow outcomes, not only licenses activated.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 16 — Q16 · Solutions

category: tool_independence · weight: 1

We could swap tools without losing the practice (documentation + skill).

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 17 — Q17 · Cross

category: honest_location · weight: 2

We know where we are on the Movemental Path—and where we skipped—without self-deception.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

Item 18 — Q18 · Cross

category: incident_posture · weight: 1

When something goes wrong, we adjust proportionately rather than panic-freeze or ban everything.

Likert scale (same labels for every item)

1. Not at all true
2. Slightly true
3. Partly true
4. Mostly true
5. Fully true

