

AI Has Already Arrived

A documentation of the AI reality facing churches, nonprofits, and institutions in 2026

The minimum viable Safety checklist that responds to it

And the working source-of-truth for productizing the work

Extended in this edition to Sandbox Discovery, Skills Development, and Solutions Deployment

A Movemental field paper and internal source document · May 2026

How to use this document

This document serves three audiences and three purposes simultaneously, and is structured so that pieces of it can be extracted for each.

For organizational leaders considering the path, Parts One and Two are the reading. They document the AI reality facing organizations like theirs and define the minimum work that responsibly addresses that reality. These two parts can be lifted out as a public-facing field guide.

For Movemental internal use, Part Three is the working layer for the Safety stage. For each of the seven Safety items, Part Three contains the actual boilerplate text that can be templated, the decisions the organization must make to customize the boilerplate, and Movemental's specific role at each productization level. This is the source of truth from which the free toolkit, the AI-assisted product, the facilitated engagement, and any future iteration are built.

When the team builds the free downloadable toolkit, Part Three section A1 (boilerplate) and A2 (decisions) for each item are the source. When the team builds the AI-assisted product, Part Three section A2 (decisions) becomes the intake structure and section A1 (boilerplate) becomes the generation template. When a facilitator runs a \$5,000 engagement, Part Three section A3 (Movemental role at facilitated level) is the working brief. The document is meant to be a single point of update — change boilerplate here, and downstream artifacts update from this source.

Part Four contains the productization strategy and a reference cross-walk that maps each item to each level of delivery, with notes on what work is required where.

Parts Five through Eight extend the document to the remaining three stages of the Movemental AI Path. Part Five covers Sandbox Discovery — the four-week exploratory engagement with the traffic-light evaluation system, the Use Case Portfolio, the Discernment Memo, and the Readiness Assessment. Part Six covers Skills Development — the eight-module formational course built on Movemental's LMS infrastructure. Part Seven covers Solutions Deployment — the technological deployment with five

configurations and the readiness threshold for proceeding. Part Eight is the cross-stage productization view that ties Safety, Sandbox, Skills, and Solutions into a single coherent operational and commercial structure.

The pattern across all four stages is the same: each stage produces specific artifacts; each artifact has boilerplate text, decision questions, and a Movemental role-by-level specification. The path is one path. This document keeps it coherent.

Part one: the reality of AI inside organizations in 2026

Movemental tells the organizations we serve that AI built on top of a fragmented human and informational layer amplifies dysfunction rather than reducing it. We tell them that the work has a sequence — Safety, Sandbox, Skills, Solutions — and that Safety is the non-negotiable first stage. Underneath that language is a claim about the world. The claim is that AI is already inside the organizations we serve, that it is already creating consequences, and that the absence of governance is producing measurable harm to mission, trust, and operations right now. This part tests the claim against current research.

Churches

The most authoritative current data on AI in churches comes from the 2026 State of AI in the Church Survey, conducted by Exponential AI NEXT and AIforChurchLeaders.com, drawing responses from church leaders across denominations, geographies, staff seniority, and functions. A separate Lifeway Research study of 1,003 Protestant pastors and 1,200 churchgoers, conducted in late 2025 and published in April 2026, provides a complementary view of pastor and congregant attitudes. Together they document a sector that has crossed a threshold in adoption while remaining almost entirely without governance.

Adoption is widespread and accelerating.

The 2026 Exponential survey found that 91% of church leaders support AI use in ministry. Sixty-one percent now use it weekly or daily, with 25% using it daily. That compares to a 43% weekly-or-daily adoption rate in 2024. Adoption nearly doubled in two years. Seventy-eight percent of church leaders use AI weekly or more.

Sermon preparation is the most common use case. Sixty-four percent of pastors now use AI for sermon preparation. Other primary uses include text content creation (36%), research (22% to 26% across surveys), image generation (20%), and administrative work (16%). The Lifeway study found a more conservative adoption picture among rank-and-file Protestant pastors, but with a clear demographic

skew: pastors under 55, in urban settings, with graduate degrees, and at larger churches are far more likely to be active adopters.

Governance has not kept pace with use.

Only 9% of churches have a formal AI policy of any kind. Only 7% have an AI disclosure statement that informs their congregation about AI use. Seventy-three percent have nothing written at all. Fifty-one percent have not addressed AI use with their congregation in any way.

The risks are not theoretical. They are arriving.

Sixty percent of church leaders said they are "very concerned" about AI voice cloning being used to defraud their congregants. Twenty-five percent — one in four — said AI-generated scams or misinformation have already affected their church community. The pattern is documented in incident reports across multiple states (Alabama, Florida, Missouri, Nebraska, New York) and internationally. Scammers harvest publicly available sermon footage, livestream archives, and social media content to clone pastor voices and likenesses, then deploy those clones to solicit emergency donations from congregants. High-profile cases include Father Mike Schmitz and several megachurches. The FBI in late 2025 warned that AI-supercharged holiday scams had produced more than 9,000 AI-related complaints.

Theological misalignment is the top concern.

Seventy-five percent of church leaders named theological misalignment as their top ethical concern. The concern is well-founded. Generative AI tools are trained on enormous corpora of internet text, including theological material from traditions across the spectrum. The output can cite scripture, use pastoral language, and still be doctrinally inconsistent with a specific church's convictions. Sixty-seven percent of church leaders also worry that AI will gradually replace the human connection that pastoral ministry depends on.

Congregations are watching.

Eighty-seven percent of church leaders said their congregation brings up AI regularly. Only 13% said the topic has not surfaced. Congregants are forming opinions about church AI use whether or not the church communicates about it. The question for church leadership is not whether the congregation will think about AI. They are already thinking about it. The question is whether the church's leadership shapes that conversation or absorbs whatever conclusions form in its absence.

Nonprofits

The most authoritative current data on AI in nonprofits comes from the 2026 Nonprofit AI Adoption Report, a benchmark study of 346 organizations conducted by Virtuous and Fundraising.AI and published in February 2026. A complementary view comes from the TechSoup 2025 State of AI in

Nonprofits Benchmark Report and a February 2026 Forvis Mazars analysis framing AI governance as a fiduciary responsibility for nonprofit boards.

Adoption is near universal.

Ninety-two percent of nonprofits use AI in some form. Seventy-nine percent of nonprofits report small to moderate efficiency gains from AI use — faster drafts, quicker research, improved content quality. The early-adopter question has been answered. The remaining questions are about how AI is being used and at what cost.

The use is fragmented and individual.

Eighty-one percent of nonprofits use AI individually rather than through shared workflows. The dominant pattern is one staff member experimenting with ChatGPT or a comparable tool — while the rest of the team continues with manual processes. Only 7% of nonprofits report major improvements in organizational capability through AI use, despite the 92% adoption rate.

Governance is broadly absent.

Forty-seven percent of nonprofits have no AI governance policy at all. The TechSoup benchmark found a more pessimistic figure: 76% of nonprofits have no formal AI strategy. Without a regulatory deadline or a major-funder requirement, AI policy work has been perpetually deprioritized. Staff find tools that help them, use them, and the organization catches up — often years later, sometimes only after an incident.

Donor and beneficiary data is at risk.

Nonprofit AI use overwhelmingly involves the organization's most sensitive data: donor giving histories, beneficiary records, program intake information, and sometimes protected health or financial details. When staff paste this data into consumer AI tools like the free tier of ChatGPT, the data may be retained by the model provider or used for model training. The Forvis Mazars February 2026 analysis explicitly frames this as a fiduciary issue. A board's traditional duties of care, loyalty, and obedience now include technological oversight.

Mission integrity is in tension with AI output.

AI-generated content reflects the statistical center of its training data, which is often misaligned with the mission, voice, theological commitments, or community standards of the nonprofit using it. Examples include AI-drafted donor appeals that subtly mischaracterize the program; AI-generated impact reports that round inconvenient details into a tidier narrative; AI-translated communications that lose the cultural specificity the original program was designed to serve. Each of these is an erosion of mission, often invisible to the staff producing it.

Institutions: seminaries, training networks, and denominational bodies

Institutions face a different version of the AI question because they operate on two timelines simultaneously: their own institutional adoption, and what they teach the next generation of leaders about AI.

Students are using AI heavily, with or without guidance.

Generative AI use in higher education is functionally universal among students under 30. International benchmarks show 63% of students using AI chatbots for text generation, 58% for translation, 55% for explaining difficult concepts, and 52% for summarizing academic articles. Seminary students are using AI for sermon drafting, exegesis assistance, paper writing, theological research, and increasingly as a kind of pastoral conversation partner.

Institutional policy varies wildly.

Some seminaries have published thoughtful policies (Columbia Theological Seminary, Saint Meinrad, Puritan Reformed, Austin Presbyterian). Others have nothing. There is no association-wide standard from the Association of Theological Schools. Faculty within the same seminary often hold different policies because no institutional policy exists to align them. Curriculum has not yet caught up to the formative reality.

Network responsibility is the unique institutional weight.

Whatever an institution adopts internally becomes the implicit template their member churches, nonprofits, and partner organizations look to. A denominational AI charter influences hundreds or thousands of congregations. A seminary's AI policy influences every church its graduates lead. Institutional inaction is not neutral. It is a choice that propagates through the network as the absence of guidance.

What is true across all three segments

Reading the church, nonprofit, and institutional data side by side, four patterns appear in all three: adoption has crossed a threshold no organization can opt out of; governance is the universal gap; the audiences these organizations serve are not unified about AI; and AI is already inside the organization in shadow form, below the line of formal organizational awareness.

AI built on top of a fragmented human and informational layer amplifies dysfunction rather than reducing it. The data confirms the claim. The governance gap is the fragmentation. The amplification is already underway. The work of Safety is to close the gap before further amplification compounds it.

Part two: the minimum viable Safety checklist

The Safety stage exists to do a specific, concrete, urgent piece of work: to bring an organization's AI use out of the shadows and into governed daylight, in the smallest set of artifacts that responsibly address the documented reality. It is not a full risk-management program. It is not a multi-month consulting engagement. It is the minimum that protects the organization from the harms the data shows are already arriving.

The MVP Safety checklist contains seven items. Six are written artifacts. The seventh is a single act of organizational communication. Together they take roughly two weeks to draft and ratify when an experienced facilitator is involved, and somewhat longer when the organization works through them independently.

The seven items in summary

- **Item 1.** AI Use Charter — the principles and posture document
- **Item 2.** Acceptable Use Policy — the operational rules for staff
- **Item 3.** Data Handling Standards — what data may touch which tools
- **Item 4.** Disclosure and Attribution Standard — when AI involvement is named publicly
- **Item 5.** Pastoral / Programmatic / Educational Care Boundaries — where AI must not go
- **Item 6.** Voice Cloning and Impersonation Response Plan — the protocol for the active criminal threat
- **Item 7.** Communication to the Constituent Base — the public act

What this list deliberately does not include

Several items that appear on more elaborate AI governance frameworks are not on this MVP list: vendor evaluation frameworks, formal AI risk registers, full incident response protocols beyond impersonation, vendor data processing agreements, and annual review processes. These are appropriate for organizations that have completed the MVP work and want to mature beyond it. They are not minimum-viable.

Part three: the working source of truth

For each of the seven items, this section provides three things: the actual boilerplate text that can be templated and customized, the decisions the organization must make to customize the boilerplate, and Movemental's specific role at each of the three productization levels.

The three productization levels referenced throughout this section are:

Level 1: Free Toolkit. Downloadable field guide on movemental.ai. The organization completes the work itself using the boilerplate and questions provided. Free in exchange for email.

Level 2: AI-Assisted Self-Service Product. A web product or Claude skill that walks an organization through structured intake (the decisions/questions) and generates draft artifacts (the boilerplate, customized) for the leadership to review and ratify. Priced at approximately \$1,997.

Level 3: Facilitated MVP Engagement. Two-week facilitated engagement with a Movemental facilitator who runs the leadership conversations, customizes the artifacts, and supports ratification. Priced at \$5,000.

Each item below documents the exact role Movemental plays at each level. This is the working layer for product build, facilitator training, and AI-assisted product configuration.

Item 01: AI Use Charter

What this is

A short document — typically two to four pages — that names the organization's posture on AI. It declares the principles guiding AI use, the lines AI should not cross in this organization, and the values the leadership wants the staff to apply when AI questions arise that the charter does not anticipate. The charter is not a policy. It is the page a board ratifies and a senior leader can hand to a new staff member to communicate "this is how we think about this."

Cost of not doing it. *Without a charter, the organization has no shared frame for AI decisions. Every staff member operates from their own theology, ethics, or vendor enthusiasm. The communications director, the pastoral care lead, and the executive director all hold different assumptions, and the inconsistency surfaces as small mistakes that erode trust over time. In the 2026 data, this is the gap behind the 91% of churches and 47% to 76% of nonprofits with no governance: there is no shared starting point for the rest of the work.*

A1. BOILERPLATE TEXT

The following is the templated language for an AI Use Charter. Square-bracketed fields are organization-specific and require completion. Italicized parenthetical notes are guidance for the customizer and are removed in the final document.

AI Use Charter for [Organization Name]

Adopted by [Board / Elder Council / Leadership Team] on [Date].

Why this charter exists. Artificial intelligence tools have entered the daily work of [Organization Name] whether we have authorized them or not. Members of our staff, our volunteers, and our community are using AI tools to draft communications, conduct research, generate content, and increasingly to seek counsel. We believe these tools can serve our mission well when used with discernment, and we believe they can damage our mission when used without it. This charter names the principles by which we will use AI in our work, the lines we will not cross, and the values we expect our staff and partners to apply when situations arise that this charter does not anticipate.

Our mission is [insert mission statement or short framing of organizational purpose]. AI is a tool. Our mission is the measure. Where AI serves the mission, we use it. Where AI threatens the mission, we do not.

Principles guiding our use of AI.

First: the human comes first. AI does not replace human discernment, presence, or relationship in the work that matters most. We use AI to serve people, not to substitute for the humans those people came to be in relationship with.

Second: integrity is non-negotiable. [For churches: The theological integrity of our teaching, preaching, and pastoral care.] [For nonprofits: The integrity of our mission, our beneficiary relationships, and our donor trust.] [For institutions: The integrity of our formation work and our accountability to the leaders we are forming.] AI-generated or AI-assisted content that touches these areas is reviewed by a human responsible for that integrity before it reaches the people we serve.

Third: data is held with care. The data of our [members / donors / beneficiaries / students] is sacred. We do not expose it to AI tools that do not meet our data handling standards. We treat our data with the same care our people have a right to expect.

Fourth: transparency builds trust. We tell our [congregation / donors / community / students] what role AI plays in our work. We do not pretend AI is not present. We do not pretend AI is more central than it is. We are honest about how we use it, and we invite questions.

Fifth: discernment over efficiency. When efficiency and discernment pull in different directions, we choose discernment. AI tempts every organization toward speed. Our calling is not speed. Our calling is faithfulness in the work we have been given, and we will not let AI rush us past it.

Lines we will not cross.

We will not use AI as a substitute for [pastoral care / direct beneficiary care / spiritual formation / the relational work that defines our mission].

We will not expose [member records / donor data / beneficiary information / student records / minor-related data] to AI tools that have not been vetted under our data handling standards.

We will not publish AI-generated content under [our pastoral / leadership / institutional / staff] name without a human author taking responsibility for the substance of what is published.

We will not use AI in ways that are deceptive to the people we serve.

Living with this charter. This charter is a starting point, not a finished policy. Specific operational decisions are made in our Acceptable Use Policy and our Data Handling Standards, which derive from this charter. When a situation arises that none of these documents anticipate, we expect our staff to apply the principles above and to err on the side of [discernment / caution / consultation with leadership]. When in doubt, ask. When the question is significant, escalate.

We will review this charter [annually / every two years]. We will revise it as our understanding deepens and as the technology changes.

Ratified by [Names of Board or Leadership]. Signed [Date].

DECISIONS THE ORGANIZATION MUST MAKE

To customize the AI Use Charter, the organization must answer the following questions. The answers populate the bracketed fields in the boilerplate and inform the language choices throughout.

- What is our organization's name and the formal body that will ratify this charter (board, elder council, leadership team)?
- What is our mission statement, in one or two sentences, framed in the language we already use about ourselves?

- What is the integrity standard most central to our work? (For churches: theological tradition and pastoral integrity. For nonprofits: mission, donor, and beneficiary integrity. For institutions: formation and accountability integrity.) Name it specifically.
- What language do we use to describe the people we serve? (Members, congregants, donors, beneficiaries, students, partners, community.) Use our actual language consistently.
- What are the three to five non-negotiable lines we want named explicitly? Beyond the defaults provided, what additional lines does our specific tradition or mission require?
- What is our default disposition when in doubt? Discernment, caution, consultation, prayer, escalation? Pick the language that fits how we actually want staff to behave.
- Who specifically signs and ratifies this charter, and on what date? Name the people who carry the institutional weight.
- What is our review cadence — annually, every two years, or some other interval?

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

We provide the boilerplate above as a downloadable Word and Markdown document. We provide the eight decision questions as a worksheet. The organization fills in the brackets and adapts the language themselves. We provide a one-page guide on common pitfalls (writing it too long, making it too theological for the staff to use, drafting it without leadership in the room). The organization owns the work end to end. We provide no human support.

At Level 2 (AI-Assisted Product, \$1,997):

The eight decision questions become the structured intake. The organization completes the intake in roughly fifteen to twenty minutes. The product generates a customized AI Use Charter draft using the boilerplate language tuned to their answers, their tradition, and their voice. The organization reviews the draft, edits in the product interface, and exports the final ratification-ready Word document. The product also generates a one-page board briefing memo that summarizes the charter for ratification. We provide a single asynchronous question-and-answer thread (handled by a Movemental team member or by the AI itself, escalating to human as needed) to address questions during ratification.

At Level 3 (Facilitated MVP, \$5,000):

A Movemental facilitator runs a 90-minute leadership conversation that surfaces the answers to the decision questions in the room. The facilitator drafts the charter within 72 hours, returns it for one round of leadership feedback, and supports the board ratification meeting (attending if helpful). The Charter is one of seven artifacts produced in the two-week engagement. The facilitator handles voice, language, theological texture, and the political navigation of getting the document ratified. The Charter is finalized by end of week one of the engagement.

Item 02: Acceptable Use Policy

What this is

A working document that names which AI tools are approved for use in the organization, what they are approved for, what they are not approved for, and what the consequences are for violating the boundaries. This is the operational artifact staff actually consult. It should be specific enough to answer real questions: Can I paste donor data into ChatGPT? Can I use AI to draft a sermon and not disclose it? Can a youth pastor use an AI counseling app with a student?

Cost of not doing it. Without an acceptable use policy, staff make individual judgments under time pressure, and those judgments accumulate into organizational practice without anyone deciding it should. The most common failure mode in the 2026 nonprofit data — 81% of nonprofits using AI individually without shared workflows — is downstream of this gap. Without a policy, donor data ends up in consumer AI tools. Pastoral notes get processed through models that retain training data. Beneficiary records get pasted into chatbots for "summarization." Each is a real incident pattern, and each is preventable with one written page.

A1. BOILERPLATE TEXT

Acceptable Use Policy for AI Tools at [Organization Name]

Effective [Date]. Reviewed [Cadence].

Purpose. This policy governs the use of artificial intelligence tools by all staff, contractors, and volunteers of [Organization Name]. It exists to operationalize the principles in our AI Use Charter and to give our team clear guidance on what is and is not appropriate.

Approved tools. The following AI tools are approved for use in our work, subject to the limitations below.

[List approved tools — typically: ChatGPT (paid tier), Claude (paid tier), Microsoft Copilot, Google Gemini, Grammarly, transcription tools like Otter or Fireflies, image generation tools as applicable. Specify whether free tier or paid tier is approved — typically only paid tier for any tool that handles organizational data.]

Tools not on this list require approval from [Designated Leader] before use for organizational work.

Approved use cases. Staff may use approved AI tools for the following purposes:

- Drafting and editing internal communications (subject to review before sending externally)
- Research and information gathering (with verification of any factual claims before use)
- Brainstorming and idea generation
- Translation drafts (subject to review by a fluent speaker before publication)
- Administrative tasks: scheduling, meeting summaries, agenda drafting
- [Sermon preparation / proposal drafting / curriculum drafting / etc., as applicable] — subject to the disclosure standards in our Disclosure and Attribution Standard
- Image generation for [internal use / approved external surfaces]

Use cases requiring additional review. The following uses require sign-off from [Designated Leader] before proceeding, and may have additional documentation requirements:

- Any use involving [member / donor / beneficiary / student] data
- [For churches: any use that touches pastoral care, theological content, or worship preparation]
- [For nonprofits: any use that touches program delivery, beneficiary intake, or donor communications at scale]
- [For institutions: any use in formative settings, accredited course delivery, or institutional communications]
- Any use that produces content published under our organizational name

Prohibited uses. Staff may not use AI tools for the following purposes:

- Substitute for direct [pastoral care / beneficiary care / formation / spiritual direction / human relationship]
- Processing of [member directory information / donor giving records / beneficiary case records / student records / financial records / minor-related data] in tools not approved under our Data Handling Standards
- Generation of content presented as the personal voice of [a pastor / executive / faculty member] without that person's explicit authorization
- Deceptive practices, including impersonation of any individual or organization
- Bypass of any review or approval process required by this policy or by other organizational standards

Review and approval workflow. Content produced with substantial AI assistance and intended for external publication is reviewed by [Specific Role] before release. Theologically sensitive content [for churches] is reviewed by [a pastor or designated theological reviewer]. Content involving [donors / beneficiaries / students] is reviewed by [the leader responsible for that constituency].

Personal use. Staff may use AI tools for personal purposes outside of organizational work without restriction by this policy. However, staff may not use organizational accounts for personal purposes, and may not feed organizational data into personal-account AI tools.

Violations and consequences. Violations of this policy are addressed through [Organization's standard staff accountability process]. Significant violations — particularly those involving sensitive data or reputational risk — may result in [disciplinary action up to and including termination / restoration of trust through a defined process / consultation with the board].

Questions. Questions about this policy should be directed to [Designated Leader]. When in doubt about a specific use case, ask before proceeding.

Adopted by [Authority] on [Date]. Reviewed [Cadence].

DECISIONS THE ORGANIZATION MUST MAKE

- What is the title and effective date of this policy, and what is our review cadence?
- Which specific AI tools have we already authorized, and which tier of each? (Free vs. paid matters — paid tiers typically have stronger data handling protections.)
- Who is the designated leader who approves new tools and handles questions?
- What constituency-specific language do we use? (Members, congregants, donors, beneficiaries, students, partners, community.)
- What are our segment-specific sensitive use cases that require sign-off? List the three to five categories that matter most to us.
- Who reviews content involving each category? (For churches: who reviews theological content? For nonprofits: who reviews donor-facing content? For institutions: who reviews formative content?)
- What is our staff accountability process, and what consequence language fits our culture? (Some organizations are direct about termination; others prefer restoration language; some require board consultation for serious matters.)
- How do we want to handle personal use of AI by staff? (Most organizations explicitly do not regulate personal use, but draw the line at using organizational accounts or feeding org data into personal tools.)

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

We provide the boilerplate as a Word and Markdown document. We provide the eight decision questions as a worksheet. We include a one-page reference card listing the most common AI tools and their data-handling characteristics (e.g., free tier of ChatGPT trains on your data; ChatGPT Team and Enterprise do not). The organization completes the policy themselves.

At Level 2 (AI-Assisted Product):

Structured intake captures the eight decision answers. The product generates a customized Acceptable Use Policy draft. The product also generates a one-page staff-facing summary card (the policy is for the document; the summary card is for the wall of the staff break room). The organization reviews and edits in-product, then exports finalized versions. The product maintains an updated reference list of AI tools with current data-handling characteristics, refreshed quarterly, so the policy reflects current reality.

At Level 3 (Facilitated MVP):

A facilitator runs a 60-minute working session that surfaces the eight decision answers and produces the policy draft live. The facilitator handles the politically sensitive piece — the consequences-and-violations section — with care for the organizational culture. Final policy delivered within 72 hours. The facilitator also customizes the staff summary card and a five-minute talking script that the senior leader uses to introduce the policy to staff.

Item 03: Data Handling Standards

What this is

A specific document — usually one to two pages — that names what categories of organizational data may and may not be exposed to AI tools, and under what conditions. The categories are sector-specific. The handling rules are mostly standard. The mapping of which data the organization actually has and where it lives is genuinely organization-specific.

Cost of not doing it. Without data handling standards, staff treat all data as equivalent and make the worst kind of mistake: pasting the most sensitive information into the least controlled tool. The 2026 nonprofit data documents this directly — donor records and beneficiary data routinely flowing through consumer AI tools without contractual data protection. Major-donor confidence, legal exposure, and beneficiary safety can be compromised in seconds, and the organization often does not know until much later.

A1. BOILERPLATE TEXT

Data Handling Standards for AI Tools at [Organization Name]

Effective [Date]. Reviewed [Cadence].

Purpose. This document classifies the data our organization holds, defines which categories of data may be processed by AI tools, and specifies the conditions under which any such processing may occur. It operationalizes the data-protection commitments in our AI Use Charter.

Data classification. Our organization holds data in the following categories. Each category has a specific handling standard.

Category A — Public Information. Information that is already publicly available about our organization: published sermons, public newsletters, public-facing program descriptions, public board members, public financial summaries. Handling: may be processed in any approved AI tool (paid tier) without further restriction.

Category B — Internal Operational Information. Information used in the course of organizational work but not generally protected: internal meeting notes, draft communications, project plans, internal training material. Handling: may be processed in approved AI tools with paid-tier accounts that do not retain inputs for training. Free-tier consumer AI tools are not approved for this data.

Category C — Constituent Identifying Information. [Member directory data / donor names and giving histories / beneficiary intake records / student records and academic information]. Handling: may only be processed in AI tools with explicit contractual data protection (Enterprise-tier accounts with signed Data Processing Agreements, on-premises tools, or organizational AI tools we have specifically vetted). Consumer AI tools — even paid tiers — are not approved unless their data terms have been verified by [Designated Leader].

Category D — Sensitive Personal Information. [Pastoral care notes / beneficiary trauma disclosures / financial records of individuals / minor-related data / health information / immigration status / counseling records]. Handling: not approved for AI processing in any consumer or general-purpose tool. May only be processed in tools we have specifically vetted, with a documented use case approved by [Designated Leader]. When in doubt, do not use AI.

Category E — Privileged or Protected Information. [Attorney-client communications / executive session board materials / specific personnel records / data subject to FERPA, HIPAA, or other regulatory protection where applicable]. Handling: never processed in AI tools without explicit authorization from [Designated Leader] and review of applicable legal and regulatory requirements.

Specific prohibitions.

- Never paste [donor giving histories / beneficiary records / student records] into the free tier of ChatGPT, Claude, or any consumer AI tool.
- Never use personal AI accounts to process organizational data in any category above Public.
- Never use AI transcription tools for [pastoral care conversations / counseling sessions / sensitive board discussions] without explicit authorization and informed consent of all parties.
- Never feed [member directory / mailing list / donor database] in bulk to AI tools for "analysis" without explicit authorization.

What to do when you are unsure. If you are not sure which category a piece of data belongs to, treat it as the higher category. When in doubt, ask [Designated Leader]. The cost of asking is small. The cost of guessing wrong can be substantial.

Mapping our actual data. Appendix A of this document is an internal-only mapping of where each data category lives in our organization (which systems, which staff have access, which third-party services hold copies). This appendix is reviewed annually and updated as our systems change.

Adopted by [Authority] on [Date].

DECISIONS THE ORGANIZATION MUST MAKE

- What constituency-specific data do we hold? List the actual categories — for churches: member directory, giving records, pastoral care notes, sermon archive, prayer requests, minor-related data. For nonprofits: donor records, beneficiary intake, program records, financial records, sensitive case files. For institutions: student records, faculty research, financial data, denominational records.
- What sensitive personal information do we hold that requires Category D treatment? List specifically — counseling notes, immigration data, health information, minor records, etc.
- What regulatory categories apply to us? (FERPA for institutions handling student records, HIPAA-adjacent obligations if handling health-related data, donor confidentiality commitments under state law.)
- Which specific AI tools have we approved for which categories? Map the tools we actually use to the categories they handle safely.
- Who is our designated leader for data handling questions, and what is the escalation path?

- Where does each data category actually live in our organization? Which systems, which staff have access, which third parties hold copies? (This is the mapping work that requires real organizational knowledge and is the part Movemental cannot do without the organization's direct input.)
- What are our existing data protection commitments to constituents? (Donor confidentiality promises, member privacy expectations, student rights statements.) The data handling standards must be consistent with these.

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

We provide the boilerplate plus a sector-specific data category template (one for churches, one for nonprofits, one for institutions). We provide a Data Mapping Worksheet that helps the organization complete the appendix work. We provide a reference table of common AI tools and their current data-handling characteristics. The organization performs the mapping work themselves, which is the part that takes the most time.

At Level 2 (AI-Assisted Product):

The intake structure walks the organization through data categories specific to their segment. The product generates a customized Data Handling Standards document. The product cannot complete the appendix mapping without organizational input — this requires the organization to actually inventory their systems. The product provides a guided mapping workflow but the organization does the work of identifying their own data locations. The product maintains the AI-tool reference table updated quarterly.

At Level 3 (Facilitated MVP):

A facilitator runs a 90-minute data inventory session with the operations leader and IT-equivalent role. This is the most labor-intensive of the seven items because it requires actual mapping. The facilitator produces the standards document plus the appendix mapping. The facilitator also identifies any high-risk data flows that need immediate attention (e.g., a staff member who has been pasting donor records into the free tier of ChatGPT) and produces a remediation note. The mapping work is the irreducible human element of this item — even at this level, the organization must produce the inventory; the facilitator structures and codifies it.

Item 04: Disclosure and Attribution Standard

What this is

A statement of when and how the organization discloses AI involvement in its public-facing work. The standard need not be elaborate. It needs to answer one question clearly: when is AI involvement disclosed, and in what form?

Cost of not doing it. Without a disclosure standard, the organization is making an implicit promise it has not examined: that what appears in the organization's name was produced by the organization's human voice. When that promise is later revealed to have been broken, the credibility damage is significantly larger than if the practice had been disclosed from the start. Only 7% of churches have any disclosure statement at all. The 87% of churches whose congregations are bringing up AI questions are forming opinions about church AI use without information from the church. This is the most common path to trust erosion and the cheapest to prevent.

A1. BOILERPLATE TEXT

Disclosure and Attribution Standard for AI-Assisted Content at [Organization Name]

Effective [Date].

Our commitment. We tell our [congregation / donors / community / students] when AI plays a substantive role in the content we publish under our name. We do not pretend AI is absent. We do not pretend AI is more central than it is. Honesty about our practices is part of how we maintain trust.

What requires disclosure. The following types of content disclose AI involvement when AI has played a substantive role in their creation:

- [Sermons / homilies / teaching content] published in writing or recorded for distribution
- [Pastoral letters / executive director letters / institutional communications] published under a named leader's name
- [Donor appeals / fundraising letters / grant applications] presented as the organization's direct communication
- [Curriculum / formation content / educational materials] that bear our institutional name
- [Public-facing reports / impact statements / annual reports]

What does not require disclosure. The following types of work do not require disclosure of AI assistance, because AI use in these contexts is so common that disclosure would be noise rather than information:

- Routine administrative communications (scheduling, logistics, internal coordination)

- Spelling and grammar tools (Grammarly, spell-check) used to refine human-authored work
- Translation drafts that are then reviewed and finalized by a fluent human
- Internal-only documents not published outside the organization
- Brainstorming and research support that informs but does not author the final work

What "substantive role" means. AI played a substantive role when the AI generated material that is included in the final published work without being substantially rewritten in human voice. Editing AI output for grammar is not substantive use. Using AI to draft a paragraph that appears in the final work, with light editing, is substantive use. The test: if the AI had not been involved, would the published content be meaningfully different in substance? If yes, disclose.

How we disclose. Disclosure takes one of the following forms, depending on context:

- A standing notice on our website that names our AI use practices in general terms
- A brief footnote or end-note on long-form published content noting AI involvement (e.g., "This sermon was prepared with the assistance of AI tools for research and outline development. The content and theological positions are the author's own.")
- A statement at the beginning of formal published content where appropriate

What we do not do. We do not publish AI-generated content under a named leader's personal voice without that leader's direct authorship and authorization. AI cannot be the ghostwriter that the public believes is the leader. The leader's name on the work means the leader has done the work.

Adopted by [Authority] on [Date].

DECISIONS THE ORGANIZATION MUST MAKE

- What public-facing content categories do we produce that require disclosure consideration? (Sermons, newsletters, donor appeals, social media, formal reports, etc.)
- What is our disposition on disclosure: tell-everything, tell-when-substantive, or tell-rarely? Most organizations land at tell-when-substantive. Some traditions push toward tell-everything. Pick the disposition that matches the trust commitments we have made to our constituents.
- What is our standing notice — the language we put on our website about general AI use practices? Draft a two-or-three-sentence statement.
- What footnote language do we use for long-form content? Draft the standard form.

- What is our position on AI ghost-writing under named leader voice? (Most organizations explicitly prohibit this. Some will allow it with the named leader's explicit authorization and review.) Pick our line.
- Who reviews content for disclosure compliance, and at what point in the workflow?

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

Boilerplate plus the six decision questions plus three sample standing notices (a more open one, a more cautious one, and a middle option) for the organization to choose between or adapt.

At Level 2 (AI-Assisted Product):

Intake captures the six decision answers. Product generates a customized Disclosure and Attribution Standard. Product also generates the standing notice copy ready to paste on the website, the footnote template, and a one-page staff reference. Product offers three preset disclosure dispositions (transparent, balanced, conservative) that pre-configure the language tone.

At Level 3 (Facilitated MVP):

A facilitator runs a 30-minute conversation focused on the trust dimension — what have we promised our constituents, and what does honesty require given that promise? The facilitator drafts the standard plus the standing notice, the footnote template, and the staff reference. This item is among the lighter-touch items at the facilitated level.

Item 05: Pastoral / Programmatic / Educational Care Boundaries

What this is

A document specific to the segment that names where AI must not go in the work the organization considers most sacred. For churches, this is pastoral care. For nonprofits, it is programmatic boundaries — direct service delivery to beneficiaries in vulnerable situations. For institutions, it is formative boundaries — spiritual formation work, mentorship, direct theological discernment. The document does not have to enumerate every case. It has to draw a clear line that staff can apply.

Cost of not doing it. Without this document, AI quietly enters the spaces the organization considers most sacred. The 2026 data shows 67% of church leaders worry that AI will gradually replace the human connection pastoral ministry depends on. Staff under time pressure are already turning to AI for tasks that involve confidential disclosure. The cost is twofold: the immediate confidentiality risk to the person whose disclosure is being processed, and the slower erosion of the trust that distinguishes the organization's most distinctive work.

A1. BOILERPLATE TEXT

Care Boundaries for AI Tools at [Organization Name]

Effective [Date].

Our conviction. The most sacred work of [Organization Name] is the [pastoral / formative / direct care] relationship between a human carrying organizational responsibility and a human who has come to us in need, in formation, or in spiritual seeking. AI does not replace this relationship, and we do not allow AI tools to enter spaces where their presence would compromise it.

Where AI does not go.

[FOR CHURCHES — adapt language as needed]

- Confessional and pastoral counseling conversations are not transcribed by AI tools, summarized by AI tools, or processed by AI tools in any form. The confidentiality of these conversations is part of the pastoral covenant.
- Crisis pastoral care — including bereavement support, marriage crisis, suicidal disclosure, mental health crisis, and acute spiritual distress — is not delegated to AI tools, in any form. AI is not consulted as a substitute for human pastoral judgment in these moments.
- Pastoral notes following confidential conversations are not generated by AI, processed by AI, or stored in systems that expose them to AI. They are written by the pastor, stored securely, and treated as sacred records.
- Conversations with minors that involve any disclosure of vulnerability, family circumstance, or personal struggle are not processed by AI tools.
- AI does not preach. A sermon may be prepared with AI assistance for research and outlining (subject to our Acceptable Use Policy and Disclosure Standard), but the act of preaching — the words spoken in the moment to the gathered people — is the work of a human who has prayed, prepared, and stands accountable for the words.

[FOR NONPROFITS — adapt language as needed]

- Direct beneficiary care conversations — intake involving trauma disclosure, advocacy work involving sensitive personal circumstance, services delivered in person to vulnerable individuals — are not processed by AI tools without explicit beneficiary consent and clear authorization.
- Beneficiary case notes are not generated by AI or processed by AI without explicit authorization and contractual data protection.

- AI is not used to make decisions about service eligibility, resource allocation to specific beneficiaries, or other determinations that affect individuals' access to our programs without human review and accountability.
- Donor stewardship conversations — major-donor relationship work that involves personal disclosure, family circumstance, or estate planning — are not transcribed or summarized by AI tools without donor consent.

[FOR INSTITUTIONS — adapt language as needed]

- Spiritual formation work, including direction, discernment, and mentorship, is not delegated to AI tools. AI is not the spiritual director.
- Confidential conversations between students and faculty regarding vocational discernment, personal struggle, or formative concerns are not processed by AI.
- Faculty assessment of student formative readiness — for ordination, for ministry deployment, for advanced placement — is not generated or substantially informed by AI evaluation. AI may inform research; it does not assess persons.
- Theological assessment of student work that touches on questions of orthodoxy, fidelity, or doctrinal formation requires human faculty judgment and is not delegated to AI evaluation.

Where AI may serve. AI may serve the work that supports these sacred relationships: scheduling, logistics, administrative coordination, research that informs human judgment, drafting of communications that the responsible human reviews and authorizes. The boundary is not between AI and our work generally. The boundary is between AI and the specific moments in which presence, judgment, and human responsibility define the work.

When in doubt. If a staff member is unsure whether a specific use case crosses a care boundary, the answer is not. Ask before proceeding. The cost of pausing is small. The cost of crossing the boundary, even once, can be substantial and is often impossible to undo.

Adopted by [Authority] on [Date].

DECISIONS THE ORGANIZATION MUST MAKE

- What is the most sacred work our organization does? Name it specifically in our own language. (Pastoral care, formation, direct service, spiritual direction, theological judgment, etc.)
- What are the three to seven specific scenarios where AI must not enter? Be concrete. The document is more useful when scenarios are named than when only categories are described.

- What is our position on AI in sermon preparation, formation work, beneficiary care, and other lighter-touch supporting tasks? We need to distinguish "where AI may not go" from "where AI may serve."
- Who in our organization holds the responsibility for these sacred relationships, and how does the document name their authority?
- What is our theological or mission-based reason for these boundaries? The boundaries should not feel arbitrary; they should connect to what we believe about the work.
- What is the staff guidance for cases that arise that this document does not anticipate? (Most organizations land on "ask before proceeding" or "err on the side of human presence." Pick the language that fits our culture.)

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

We provide three sector-specific boilerplate documents — one for churches, one for nonprofits, one for institutions. We provide the six decision questions. We include three short reflections from sector-specific theologians or practitioners on why these boundaries matter (sourced from Hirsch, Halter, or other movement-network voices, where permission is granted). The organization customizes the document themselves. This item is the one most resistant to productization, because the lines are theologically and contextually specific. The free toolkit gets the organization to a strong starting point but cannot draw the lines for them.

At Level 2 (AI-Assisted Product):

Intake captures the six decision answers, with thoughtful prompting that draws out the organization's actual theological convictions about what is sacred. The product generates a customized boundaries document. The product flags this document as one that should be reviewed by senior leadership and (for churches) a pastor or theological reader before ratification. The product does not promise that its draft is theologically complete; it promises a serious starting point that the organization's leadership owns finalizing. The product offers an optional 30-minute consultation with a Movemental facilitator as an add-on for this item specifically (priced at \$300), because this is the item where AI assistance most needs human review.

At Level 3 (Facilitated MVP):

This is the most theologically and pastorally substantive item in the engagement. A facilitator (ideally one with pastoral or theological training) runs a 90-minute conversation focused specifically on this question: where does AI not belong in this organization's most sacred work? The facilitator produces the boundaries document with deep customization to the organization's tradition, theology, and pastoral practice. This item is where the \$5,000 facilitated engagement most clearly

justifies its price. The facilitator may recommend additional theological review by the organization's own pastoral or theological resources before ratification.

Item 06: Voice Cloning and Impersonation Response Plan

What this is

A specific protocol for what the organization does when its leaders' voices, faces, or names appear in AI-generated content used to defraud or manipulate constituents. The protocol names the people responsible for monitoring, the channel through which constituents can report suspicious communications, the verification standard the organization will use, the platforms the organization will pre-emptively notify, and the message to the constituent base that establishes the verification norm in advance.

Cost of not doing it. *This is the only item on the list that addresses a documented, currently-active criminal threat. Twenty-five percent of churches have already experienced AI-generated scams or misinformation in their community. Sixty percent of church leaders are very concerned about voice cloning. The FBI has issued public warnings. Without a response plan, the first incident is the disaster. With a plan, the first incident is the activation of a protocol the organization has practiced. The cost of not having a plan is paid by the most vulnerable members of the organization — typically older constituents — who lose money to a voice they trusted.*

A1. BOILERPLATE TEXT

Voice Cloning and Impersonation Response Plan for [Organization Name]

Effective [Date]. Practiced [Cadence].

Why this plan exists. AI-generated impersonation of religious and nonprofit leaders is a documented and active threat. Scammers harvest publicly available audio, video, and likeness from sermons, livestreams, social media, and recorded events to create convincing synthetic content used to solicit donations, request emergency wire transfers, and manipulate constituents. Our [pastors / executives / leaders] are at risk because of the public nature of our work. This plan defines how we prevent, detect, and respond to such incidents.

What we tell our constituents in advance. We commit to telling our [congregation / donor base / community / students] the following, repeatedly and through every channel we use:

- Our [pastors / leaders / staff] will never contact you through unsolicited calls, texts, or messages asking for money, gift cards, wire transfers, or urgent personal financial assistance.
- Real giving opportunities are always communicated through our official channels: [our website, our official email list, our published giving page, in-service announcements].
- If you receive a surprising message that appears to be from one of our leaders requesting urgent financial action, assume it is fraudulent. Verify by contacting our office directly at [Official Number] or [Official Email] before acting.
- You can report suspicious communications to [Designated Contact / Reporting Channel].

We will publish this commitment on our website, include it in our regular communications, address it from the [pulpit / platform / staff meeting] at least quarterly, and reinforce it whenever a new wave of scam activity is reported in the broader sector.

Detection: how we monitor for impersonation.

- [Designated Person] reviews social media mentions and direct reports of suspicious content monthly.
- We have established a reporting channel — [specific email or phone] — where staff and constituents can report suspected impersonation.
- Our staff is trained to recognize the signs of synthetic content: unnatural pauses, slight visual inconsistencies, requests for urgent action, requests for nonstandard payment methods (gift cards, wire transfers, cryptocurrency), and messages from new or unverified accounts.

Response: what we do when an incident is detected.

Step 1 (within 24 hours of confirmed incident): [Designated Leader] is notified and assesses scope. The incident is logged.

Step 2 (within 48 hours): A communication is sent to our [congregation / donor base / community / students] through all primary channels naming the incident, identifying the impersonated leader, and reaffirming our verification standard.

Step 3 (within 72 hours): Reports are filed with the relevant platforms (Facebook, Instagram, TikTok, YouTube, etc.) requesting takedown of the impersonating content. Reports are filed with the FBI Internet Crime Complaint Center (IC3) at ic3.gov for any incident involving financial fraud.

Step 4 (within 7 days): If any constituent has been financially harmed, we offer pastoral or organizational support and connect them with [law enforcement / consumer protection / counseling resources] as appropriate.

Step 5 (ongoing): The incident is reviewed in our next leadership meeting, and any improvements to this plan are noted and implemented.

Practice. We practice this plan annually through a tabletop exercise. The practice ensures that the people responsible know what to do, the communication templates are ready, and the reporting channels are functional.

Designated roles. [Specific person] is responsible for monitoring. [Specific person] is responsible for incident response. [Specific person] is responsible for the constituent communication. These roles may be held by one person in a small organization or by different people in a larger one.

Adopted by [Authority] on [Date]. Practiced [Cadence].

DECISIONS THE ORGANIZATION MUST MAKE

- Which of our leaders are at highest risk because of public exposure? (Senior pastor, executive director, named institutional leaders with public-facing roles, faculty with public profiles.)
- Through what channels do we currently reach our constituents, and which do we consider "official"? (Website, email list, in-service announcements, member portal, etc.)
- Who specifically holds the monitoring, response, and communication roles? In a small organization, this may be one person. In a larger one, it should be different people with clear handoffs.
- What is our official reporting channel for constituents to flag suspicious communications? (Specific email, specific phone number, web form.)
- What is our official phone number and email for constituents to verify communications? Make sure these are easy to find and stable over time.
- How often will we communicate the verification commitment to constituents? (Quarterly is typical. More frequently during periods of heightened scam activity.)
- What is our practice cadence — annually is recommended. Schedule the first tabletop exercise on the calendar now.

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

We provide the boilerplate plus a one-page Constituent Verification Notice ready to paste on a website. We provide the seven decision questions plus a Tabletop Exercise script the organization

can run themselves. We provide a list of relevant reporting URLs (FBI IC3, major platform reporting, BBB) updated quarterly.

At Level 2 (AI-Assisted Product):

Intake captures decision answers. Product generates customized response plan including: the master document, the constituent verification notice ready for website paste, the constituent communication templates pre-drafted (one for the proactive announcement, one for the post-incident communication), and the tabletop exercise script. The product also includes a calendar reminder for annual practice. This is one of the items where AI generation produces excellent output because the structure is highly templatable.

At Level 3 (Facilitated MVP):

A facilitator runs a 45-minute working session that walks through the decisions and produces the plan. The facilitator also runs the first tabletop exercise (a 30-minute additional commitment, included in the engagement) so that the organization has practiced the plan once before the engagement ends. This item is fast at the facilitated level because the structure is so templatable; the facilitated value-add is the practice exercise.

Item 07: Communication to the Constituent Base

What this is

A single act of communication — a letter, a sermon, a board statement, a member email — in which the organization tells the people it serves what its position on AI is, what governance it has put in place, and what they can expect. This is not a document. It is the act of speaking publicly about the work the prior six items represent. It can be five paragraphs long. It can take fifteen minutes to read aloud. What matters is that it happens, that it is signed by the senior leadership, and that it establishes the organization as having a position rather than a void.

Cost of not doing it. *Without this communication, the prior six artifacts exist privately and the constituent base continues to form its own opinions about AI in the organization without information. The 2026 data shows 87% of church congregations bringing up AI regularly and 51% of churches not addressing AI use with their congregation at all. That gap is filled by rumor, suspicion, and projection. The cheapest, fastest, highest-leverage governance move on this list is the public communication that the work has been done.*

A1. BOILERPLATE TEXT

A letter to our [congregation / community / donors / students] on artificial intelligence

From [Senior Leader Name and Title]

[Date]

Dear [friends / family in Christ / community / partners],

I want to write to you about something we have been thinking about carefully as a [church / organization / institution]: how we will and will not use artificial intelligence in our work together.

Like you, we have watched AI move quickly into the daily life of every kind of organization. We have seen the benefits — research that takes minutes instead of hours, communications drafted faster, translations made accessible, content reaching people who could not encounter it before. We have also seen the costs — content that sounds authoritative while being wrong, voices cloned to defraud trusting people, the slow erosion of the human relationships that define the work we are called to.

We want you to know that we have done the work of thinking carefully about this. Over the past [weeks / month], our [board / leadership / elder council] has adopted [an AI Use Charter / a set of governance documents / a framework] that names how we will use AI in our work and where we will not allow it to go. The full documents are available on our website at [URL]. The short version is this:

First, we believe AI can serve our work, and we will use it carefully — for research, for drafting, for administrative support, for the kinds of tasks that free our staff to do the work that requires their full presence.

Second, we believe there are places AI does not belong in our work. [Pastoral care / direct beneficiary work / spiritual formation / the relationships at the center of what we do] is human work. We do not let AI substitute for the human presence and judgment those relationships require.

Third, we will be honest about how we use AI. When AI plays a substantial role in something we publish under our name, we will tell you. We do not want you to discover something is true that we should have told you ourselves.

Fourth, we want to protect you from a specific threat. Scammers are using AI to clone the voices and faces of pastors, executive directors, and other leaders to manipulate the people who trust them. Please know: I will never contact you out of the blue asking for money, gift cards, or urgent financial help. If you ever receive a message that appears to be from me requesting

urgent financial action, please assume it is fraudulent and verify by contacting our office at [Official Number]. Tell your friends and family in our community. The most vulnerable members of any community are the most likely to be targeted, and our protection of them depends on a shared verification habit.

Finally, we want to invite your questions. AI is going to keep changing. The questions it raises are not going to settle quickly. If you have concerns, observations, or ideas about how we should engage these questions, please bring them to us. Our work together is stronger when our constituents speak into it.

Thank you for your trust. We do not take it lightly.

[Senior Leader Signature, Title, Organization]

DECISIONS THE ORGANIZATION MUST MAKE

- Who is the senior leader signing this communication? Their voice and authority carry the message.
- Through what primary channel will this communication reach the constituent base? (Email, mailed letter, sermon, video message, member portal post, board statement.) Some organizations will use multiple channels.
- What is the appropriate length and tone for our culture? (Some communities expect formal letter format; others prefer conversational sermon-style; others want a brief executive memo.) Pick the form that fits.
- What is the URL where the full governance documents will be available? Set this up before the communication goes out.
- What official phone number and email do we want constituents to use for verification?
- What is our follow-up plan for questions that come back? (A dedicated email address, a town hall meeting, a Q&A session, designated staff to handle inquiries.)
- When will the communication go out? Set a specific date that is after ratification of the other six artifacts but before too much time has passed.

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Free Toolkit):

Boilerplate letter plus the seven decision questions plus three sample alternate forms (formal letter, conversational sermon-style, executive memo). The organization customizes the letter in their senior leader's voice.

At Level 2 (AI-Assisted Product):

Intake captures decisions. Product generates a customized letter in the form selected. The product offers an optional voice-tuning step where the senior leader provides a sample of their writing (a recent sermon, a recent newsletter, a recent letter) and the product re-tunes the draft to match the leader's actual voice. This voice-tuning capability is the highest-value AI use in the entire product, because it solves a real problem (the leader has to sign something that sounds like them) at a level no template can.

At Level 3 (Facilitated MVP):

A facilitator works directly with the senior leader on this letter — typically a 45-minute conversation in which the leader speaks the heart of what they want to say and the facilitator captures it, then drafts the letter in the leader's voice within 24 hours. The facilitator also helps the leader rehearse the in-person reading if the letter will be delivered as a sermon or address. This item closes the engagement, and is the moment the work becomes public. The facilitator's role is to ensure that moment is well-shaped.

Part four: productization strategy and cross-walk

This part documents the strategic logic that places each of the seven items at the right productization level, and provides a cross-walk reference for how Movemental delivers Safety across the three levels.

Strategic logic

Movemental's position is that AI built on a coherent foundation compounds rather than corrodes. The Safety stage operationalizes that position: it is the stage where the foundation gets built. This creates a strategic alignment opportunity. Movemental should productize Safety in ways that demonstrate the company's thesis rather than contradict it. Specifically: an AI-assisted product that helps organizations build their AI foundation is the most coherent possible expression of the company's philosophy. It says with action what we say with words.

At the same time, Safety is the most accessible-by-design stage of the path. The documented harms apply to every organization in the sector. Withholding the framework behind a paywall would be inconsistent with the missional posture of the leaders we serve and the communities they lead. Open-sourcing the framework, while productizing the implementation, is the right balance: framework freely given, work delivered at multiple price points.

The three levels of productization correspond to three buyer segments and three trust thresholds. The Free Toolkit serves organizations who cannot afford engagement and who would not pay for it; it builds top-of-funnel credibility and aligns Movemental with the missional generosity expected in the sector.

The AI-Assisted Product serves organizations with budget but a preference for self-service; it captures revenue from a segment significantly larger than facilitation can serve and demonstrates the agentic infrastructure as a customer-facing product. The Facilitated MVP serves organizations who want speed, external coordination, and a relationship with Movemental for downstream stages; it is the relationship-building tier that converts into Sandbox, Skills, and Solutions engagements.

Cross-walk reference table

The following describes Movemental's role for each of the seven items at each of the three levels. This is the operational reference for product configuration, facilitator training, and pricing decisions.

Item 1 — AI Use Charter

Free Toolkit: Boilerplate document plus decision worksheet, no human support. AI-Assisted Product: Structured intake plus AI generation plus async Q&A support. Facilitated MVP: 90-minute leadership conversation plus drafted charter plus board ratification support.

Item 2 — Acceptable Use Policy

Free Toolkit: Boilerplate plus decision worksheet plus AI tools reference card. AI-Assisted Product: Intake plus AI generation plus staff summary card plus quarterly-updated tools reference. Facilitated MVP: 60-minute working session plus drafted policy plus customized staff summary card plus leader talking script.

Item 3 — Data Handling Standards

Free Toolkit: Sector-specific boilerplate plus mapping worksheet (organization completes inventory themselves). AI-Assisted Product: Intake plus AI generation plus guided mapping workflow (organization still completes inventory). Facilitated MVP: 90-minute data inventory session plus drafted standards plus completed appendix mapping plus high-risk flow remediation note.

Item 4 — Disclosure and Attribution Standard

Free Toolkit: Boilerplate plus decision worksheet plus three sample standing notices. AI-Assisted Product: Intake plus AI generation plus standing notice paste-ready copy plus footnote template plus staff reference. Facilitated MVP: 30-minute trust conversation plus drafted standard plus all customized supporting copy.

Item 5 — Care Boundaries

Free Toolkit: Three sector-specific boilerplate documents plus decision worksheet plus sourced reflections. (Most resistant to productization.) AI-Assisted Product: Intake plus AI generation plus optional add-on \$300 facilitator consultation specifically for this item plus theological-review flag for senior leadership. Facilitated MVP: 90-minute pastorally-framed conversation plus deeply customized boundaries document. (Most theologically substantive item.)

Item 6 — Voice Cloning Response Plan

Free Toolkit: Boilerplate plus verification notice plus tabletop exercise script plus reporting URLs. AI-Assisted Product: Intake plus AI generation plus pre-drafted communication templates plus tabletop script plus annual practice calendar reminder. Facilitated MVP: 45-minute working session plus drafted plan plus first tabletop exercise run by facilitator.

Item 7 — Communication to Constituent Base

Free Toolkit: Boilerplate letter plus decision worksheet plus three sample alternate forms. AI-Assisted Product: Intake plus AI generation plus optional voice-tuning capability that adapts the letter to the senior leader's actual voice from a writing sample. Facilitated MVP: 45-minute conversation with senior leader plus drafted letter in leader's voice plus rehearsal support.

Build sequence and effort estimate

If Movemental commits to building all three levels, the recommended build sequence and effort estimate is:

Phase 1 (week 1): Refine Parts One, Two, and Three of this document into a designed downloadable Field Guide. Publish on movemental.ai as the Free Toolkit, gated by email signup. Total effort: roughly 30 hours including design.

Phase 2 (weeks 2-4): Build v1 of the AI-Assisted Product. The seven intake structures are derived from the decision questions in Part Three. The seven generation templates are derived from the boilerplate text in Part Three. The product is implemented as either a structured web form connected to a Claude or comparable model, or as a Claude Skill, depending on the platform that fits Movemental's existing infrastructure. Soft launch to existing leader network as beta. Total effort: roughly 80 hours of focused build given existing agentic infrastructure.

Phase 3 (ongoing): Continue offering the \$5,000 Facilitated MVP using Part Three as the working brief. Train any additional facilitators using Part Three as the source of truth. Iterate on boilerplate and decision questions based on real engagement data.

Update protocol

This document is the single source of truth. When boilerplate language is improved, when decision questions are refined, when sector-specific data shifts, or when Movemental's offering changes, the update happens here first. Downstream artifacts — the Field Guide, the AI-Assisted Product templates, the facilitator brief — are regenerated from this source. The document is reviewed quarterly and updated as the field evolves.

Stage 02: Sandbox

Part five: Sandbox Discovery

The Sandbox stage is the exploratory, disciplined, locally-tested discovery of where AI actually creates value in this organization, what use cases are wise to pursue, what use cases need revision before pursuit, and what use cases are off the table. It is a four-week engagement priced at \$15,000 in its facilitated form.

The premise of the Sandbox stage is that AI is not a technological challenge first. The interface to current AI tools is human communication. Cutting-edge AI tools do not present as conventional technology — they present as human simulation. The way to evaluate AI in any organization is therefore not to study it from a distance but to put real recipes in the hands of real staff doing real work, observe what happens, and adjudicate the findings against the values and constraints documented in the Safety stage. The Sandbox is not a training program. It is a structured exploration with documented outcomes.

Five-A: What Sandbox produces

Sandbox Discovery produces three connected deliverables.

Deliverable one: a Use Case Portfolio.

A working document containing every use case explored during the four-week engagement, with for each: a description of the workflow, the AI tool or recipe applied, the staff member or team that tested it, the time the test took, the value the use case promises (time saved, revenue generated, work improved), the human or ethical concerns surfaced, the integrity considerations relative to the organization's Safety documentation, and a final adjudication of green light, yellow light, or red light. The portfolio is the document the organization keeps and uses to make decisions about which use cases to deploy, modify, or set aside.

Deliverable two: a Discernment Memo.

A short narrative document — typically four to eight pages — that names the patterns the Sandbox surfaced. What kinds of use cases consistently produced value in this organization. Which kinds of use cases consistently raised concerns. What the team learned about itself in the process. What the leadership now sees that they did not see before. The Discernment Memo is the document that informs the decision about whether to proceed to Skills, to Solutions, or to neither.

Deliverable three: a Readiness Assessment.

A brief evaluation of whether the organization is ready to proceed to Skills (deepening human capacity to lead this work) or directly to Solutions (technological deployment). Not every organization needs every stage. Some are well-served by Sandbox and the Discernment Memo, then a focused Solutions deployment. Some need substantial Skills work before any further technology is added. The Readiness Assessment names which path is right for this organization based on what the Sandbox surfaced.

Five-B: The traffic-light system

Every use case tested in the Sandbox receives one of three signals. The signals are only assigned in the context of discovered potential value — the Sandbox does not adjudicate use cases the team did not pursue, did not test, or did not surface. Adjudication happens for use cases that have demonstrated some promise, with the integrity question evaluated against that promise.

GREEN LIGHT. The use case produces clear value, raises no significant concerns under the organization's Safety documentation, can be performed by current staff with reasonable training, and serves the mission rather than warping it. Green light means the use case is ready for deployment in Solutions or for immediate productive use within the boundaries already documented.

YELLOW LIGHT. The use case produces value but requires documented revision before deployment. The revision could be technical (the recipe needs adjustment), procedural (the workflow needs review or oversight added), formational (the staff using it need additional preparation), or ethical (a concern surfaced that requires a leadership decision). Yellow light means promising but not yet ready. The Discernment Memo captures what specifically must change before the use case is deployed.

RED LIGHT. The use case cannot be pursued under the organization's current Safety framework, mission integrity, or constituent commitments. The reason is documented and named clearly. Red light is not a negotiation. It is the sober recognition that not every promising-on-paper use case is wise in practice for this specific organization. Red light use cases are valuable to document because they protect the organization from later attempts to deploy what should have been declined.

Five-C: The Sandbox process, week by week

Week 0: Pre-engagement assessment.

Before the engagement formally begins, the organization completes a broad assessment that maps where each staff member currently stands. The assessment is delivered as a Claude artifact — a structured interactive intake — and asks each participant about their current workflow, their team context, their feelings about AI, the tools they currently use, their experience with AI to date, and the

friction points in their work that they wish were different. The output is a personalized starting point for each participant. This is the moment Movemental meets each person where they are, not where the organization wishes they were. The assessment respects that AI work is most usefully tested as locally as possible — down to the user — because each staff member's relationship to AI, their tasks, their fears, and their existing tools is specific to them.

Week 1: Recipe deployment and value mapping.

Each participant is given a small set of pre-built Claude skills — recipes — adapted to their specific role. The recipes are configured against the organization's Safety documentation and are pre-screened to operate within Data Handling Standards and Care Boundaries. Participants run the recipes against real workflows: a communications director drafts donor appeals with a Donor Letter recipe; a pastor researches a sermon series with a Sermon Research recipe; an executive director summarizes board materials with a Board Briefing recipe; an operations lead drafts policies with a Policy Drafting recipe. Each participant logs their experience: time taken, time saved relative to baseline, output quality, concerns observed.

The point of this week is value discovery. Not training. Not theory. Real recipes against real work, with the friction and the value both made visible. Participants do not need to understand prompt engineering or model selection to run a recipe well. The recipe handles the technical complexity. The participant brings the contextual judgment. This honors what the AI moment actually is — a frontier we need to navigate immediately rather than a topic we need to study before approaching.

Week 2: Concern surfacing and use case expansion.

Participants reconvene, having now lived with the recipes for a week. The conversation in week two surfaces what worked well, what produced unease, what surprised the participant about their own response to AI assistance, and what new use cases the participant would now like to explore based on what they have already tried. New recipes are configured for the second-week tests. Participants who saw promising patterns extend their use; participants who encountered concerns slow down and explore those concerns more deliberately.

The facilitator (Movemental staff) attends to several specific signals during week two. Where is AI-assisted work landing back in the participant's human workflow well? Where is it creating new friction? Where are participants beginning to over-trust the output? Where are participants under-trusting it and missing real value? The facilitator's job is not to push participants harder toward AI use but to help each participant calibrate accurately.

Week 3: Adjudication and revision.

The facilitator works with participants and leadership to adjudicate each tested use case. Some receive green-light status immediately. Others need a working session to determine what revision would

convert them from yellow to green, or whether the revision is itself worth doing. A small number receive red-light status with documentation of why. The Use Case Portfolio takes shape during this week.

Week 4: Discernment Memo and Readiness Assessment.

The facilitator drafts the Discernment Memo and the Readiness Assessment based on what the four weeks revealed. The leadership reviews and refines. The engagement concludes with a leadership conversation that names what the organization has learned, what it now believes is true about its relationship with AI, and what it will do next. The conversation is structured but not scripted. It is the moment the organization moves from exploring to deciding.

Five-D: Sandbox technological supports

The Sandbox is the first stage at which Movemental's technological capability becomes directly visible to the organization. Three technological supports are central.

Pre-built Claude skills (recipes). Movemental maintains a growing library of Claude skills configured for the work organizations in our segments actually do. These are not generic prompts. They are skills configured against the organization's Safety documentation, adapted to the role of the staff member running them, and pre-screened for the constraints documented in Data Handling Standards and Care Boundaries. A Sermon Research recipe knows the church's tradition. A Donor Letter recipe knows the nonprofit's donor language. A Curriculum Outline recipe knows the institution's formative commitments. The recipes carry the organizational context the participant does not have to re-create from scratch.

The assessment artifact. A Claude artifact delivered in week zero that surfaces each participant's starting point. The artifact is structured but conversational. It asks the questions that reveal where the person actually is, and produces a personalized starting profile that informs which recipes are appropriate for that person and what the facilitator should attend to during the engagement.

The Use Case Portfolio interface. A web interface (built on Movemental's existing React/Next.js/Tailwind/Supabase stack) that captures use case data as it is generated. Each participant logs their tests; the facilitator adjudicates each entry; the portfolio accumulates as a living document throughout the engagement and is exported as the deliverable at week four.

Five-E: Sandbox boilerplate templates

A1. USE CASE PORTFOLIO ENTRY TEMPLATE

Use Case [Number]: [Short Name]

Workflow tested. [Description of the actual workflow the staff member was performing — e.g., "drafting the monthly donor newsletter," "researching exegesis for a four-week Romans series," "summarizing the quarterly board financial report for the executive team."]

Recipe applied. [Name of the Claude skill or recipe used, with version reference.]

Tested by. [Staff member name and role.]

Time invested in test. [Hours or days.]

Time savings observed. [Quantitative estimate against baseline workflow, with caveats about whether the test was representative of typical conditions.]

Value promised. [Specific value the use case offers if deployed: time savings per occurrence, revenue impact if applicable, work improvement description. Quantify where possible; describe qualitatively where not.]

Frequency of underlying workflow. [How often this work happens — weekly, monthly, annually, ad hoc. Multiplies the value estimate.]

Human and ethical concerns surfaced. [What the staff member or team noticed: places where the AI output felt off, places where the workflow now requires more rather than less human attention, places where the integrity of the work is potentially affected, places where the recipient of the work might feel differently if they knew AI was involved.]

Safety alignment. [Reference to specific Safety documents that govern this use case: Acceptable Use Policy clauses that apply, Data Handling Standards relevant to the data involved, Care Boundaries that bear on this work, Disclosure and Attribution Standard implications.]

Adjudication. [Green / Yellow / Red, with reasoning.]

If yellow: required revision. [What specifically must change before this use case can move to green status. Could include: technical refinement of the recipe, addition of human review steps, formational preparation of the staff using it, leadership decision on a specific concern, additional documentation in Safety files.]

If red: rationale. [Why this use case is not pursued in this organization. The rationale should be specific enough that a future staff member who suggests the same use case can read this entry and understand the prior decision.]

A2. DISCERNMENT MEMO STRUCTURE

Discernment Memo: [Organization Name] Sandbox Discovery

Engagement dates. [Start and end dates of the four-week Sandbox.]

Participants. [Roles and number of staff who participated.]

What we tested. [Brief overview of the use cases explored, organized by domain — communications, formation, operations, etc.]

What surfaced as patterns of value. [Two to four paragraphs naming the kinds of work where AI assistance produced consistent value in this organization. Be specific: in this organization, AI was genuinely helpful for X, Y, and Z. The pattern is not generic. It reflects this team, this work, this season.]

What surfaced as patterns of concern. [Two to four paragraphs naming the kinds of work where AI assistance produced consistent unease, friction, or integrity questions. Again be specific. Concerns are honored as data, not dismissed as resistance.]

What we learned about ourselves. [One to two paragraphs naming what the team noticed about its own relationship to AI through the four weeks. People discover surprising things about themselves when they actually use these tools rather than discussing them in the abstract. Capture what surfaced.]

What we now believe is true. [Two or three short statements that the leadership endorses as the organization's current position, derived from this engagement. These statements may evolve. They are the organization's honest read of its own situation as of the end of the Sandbox.]

What we are choosing to do next. [Concrete decisions: which green-light use cases will deploy immediately, which yellow-light use cases the organization will revise and reconsider, which red-light use cases are documented as off-limits, and whether the organization will pursue Skills, Solutions, both, or neither in the next season.]

A3. READINESS ASSESSMENT TEMPLATE

Readiness Assessment: [Organization Name]

At the conclusion of Sandbox Discovery, the following observations inform our recommendation regarding next stages.

Observed staff capacity for AI work. [Description: where staff demonstrated genuine competence, where staff hit ceilings of comfort or capability, where formational work would most benefit the organization.]

Observed leadership clarity. [Description: where leadership has clear convictions about AI use, where leadership is still working through its position, where the team is awaiting leadership clarity to proceed.]

Observed technological readiness. [Description: where the organization's existing tools and infrastructure are sufficient for the green-light use cases, where infrastructure changes would unlock substantial additional value, where premature technology investment would compound problems rather than solve them.]

Recommended next stage.

[Choose one and elaborate:]

Option A: Skills Development is the priority. The Sandbox surfaced significant capacity-building opportunities; the organization's green-light use cases require deeper formation before they will be sustained. Recommend Skills Development before any further technology investment.

Option B: Solutions Deployment is appropriate now. The Sandbox surfaced clear, valuable, deployable use cases; the team is ready to operate them; existing technology with focused customization is sufficient. Recommend a focused Solutions engagement, with Skills as an optional later addition.

Option C: Both Skills and Solutions in parallel. The organization has both the capacity opportunity and the deployment opportunity; running them in sequence would lose momentum. Recommend overlapping engagement with clear coordination.

Option D: Neither at this time. The Sandbox surfaced that the organization's current state is well-served by the green-light use cases now operating, and that further investment would not currently produce proportional value. Recommend operating the current state, revisiting in 12-18 months.

DECISIONS THE ORGANIZATION MUST MAKE

- Who from our organization will participate in the Sandbox? (Typically 4-10 people across roles. Choose for diversity of role and openness to learning rather than pre-existing AI enthusiasm.)
- What time can each participant commit weekly to the Sandbox? (Realistic estimate: 3-5 hours per week per participant during the four-week engagement.)
- What workflows are we most curious about testing? (Initial list, knowing the Sandbox will surface use cases we did not anticipate.)
- What concerns or hesitations do we want to actively explore? (The Sandbox is as valuable for naming concerns as for finding value.)

- Who in our leadership owns the adjudication decisions? (Yellow and red light decisions ultimately need leadership ratification, not just facilitator recommendation.)
- What is our intended posture toward concerns that surface? (Some organizations want to push through concerns; others want to honor them as data; the right posture is a leadership decision.)
- How will the Use Case Portfolio be maintained after the engagement ends? (Living documents need owners; without one, the portfolio becomes a snapshot rather than an ongoing reference.)

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Self-Service Sandbox Toolkit, ~\$497):

Movemental publishes a downloadable Sandbox toolkit containing the assessment artifact (delivered as a Claude artifact link), a starter library of generic Claude skills the organization can adapt to its own context, the Use Case Portfolio template, the Discernment Memo template, the Readiness Assessment template, and a four-week facilitator guide that a senior leader can use to run the Sandbox internally. The organization runs its own Sandbox. Movemental provides the materials and a single 60-minute orientation call. This level is appropriate for organizations with internal capacity to facilitate their own discovery.

At Level 2 (AI-Assisted Sandbox, ~\$3,997):

Movemental provides a configured set of Claude skills tuned to the organization's segment and Safety documentation, the assessment artifact pre-customized for the organization, the Use Case Portfolio interface (web-based), and asynchronous facilitator support throughout the four weeks (typically 2-3 hours of facilitator time per week, delivered through the portfolio interface and email). The organization runs the engagement; Movemental provides the technological infrastructure and asynchronous expertise. This level scales to organizations that want strong support without full facilitation.

At Level 3 (Facilitated Sandbox Discovery, \$15,000):

A Movemental facilitator runs the full four-week engagement. The facilitator delivers the assessment, configures the Claude skills against the organization's Safety documentation, conducts weekly working sessions with the participant team, adjudicates use cases collaboratively with leadership, drafts the Discernment Memo and Readiness Assessment, and concludes with a leadership conversation. The Use Case Portfolio interface is used throughout. This is the standard offering for organizations engaging Movemental for Sandbox.

Network Engagement Sandbox (Institutions only, from \$60,000):

For denominations, training networks, and multi-site organizations, the Sandbox can be run as a Network Engagement that includes the institution itself plus 3-10 member organizations, with shared discovery, shared portfolio, and a teaching artifact published to the broader network at the conclusion. Pricing scales with scope.

Stage 03: Skills

Part six: Skills Development

The Skills stage is the formational deepening of human capacity to lead AI work in this organization with wisdom. It is an eight-week engagement priced at \$15,000 (cohort facilitation) or \$4,800 per year per organization (self-paced license).

The premise of the Skills stage is that running recipes is a starting point, not an end point. Organizations that complete Sandbox have demonstrated they can use AI productively. Organizations that complete Skills have developed the discernment to lead AI work — to evaluate new use cases, modify existing recipes, recognize when AI is the wrong tool, and form other staff in the same wisdom. The shift is from practitioner to discerner. From running well-designed recipes to making well-considered judgments about what to run, when, how, and whether.

Six-A: The metaphor and the shift

We have used the language of recipe and chef in conversation about this stage. The metaphor captures something true — the shift from following someone else's instructions to making your own informed decisions — but a stronger frame for what Skills actually accomplishes is the shift from practitioner to discerner.

A practitioner runs the recipe well. They follow the steps, attend to the inputs, produce the output. Their competence is measured by the quality of their execution. In the AI context, a practitioner uses Claude skills configured by someone else, applied to workflows they have been trained on, against data their Safety documentation has approved. They do good work.

A discerner does all of that, and also knows when the recipe should not be run. When the workflow has shifted in ways the recipe no longer fits. When the data carries weight the recipe does not account for. When the constituent on the other side of the work would experience the recipe differently than the practitioner does. When the organization's mission is served by saying no rather than yes. The discerner reads the situation, weighs it against the values and convictions documented in Safety, considers the people involved, and chooses. Their competence is measured by the quality of their judgment.

The Skills stage forms discerners. The work is not technological — though technology is the context. The work is formational — though competence is the result. Every module builds judgment as much as capability.

Six-B: Learning outcomes for Skills Development

At the conclusion of Skills Development, participants are able to:

- Read an AI use case in context — understanding the workflow, the data, the people, the integrity considerations, and the organizational values that bear on the case — and articulate a position on whether and how to proceed, with reasoning their leadership can review and ratify.
- Modify an existing Claude skill or recipe to better fit a workflow that has evolved or a context the original recipe did not anticipate, including adjusting prompts, context windows, output structure, and integration with downstream systems where applicable.
- Identify the difference between a use case that needs technical refinement, a use case that needs procedural change, a use case that needs formational preparation of the staff using it, and a use case that should not proceed at all — and recommend the appropriate response for each.
- Evaluate AI output for quality, fit-to-context, alignment with the organization's voice and convictions, and absence of subtle errors that the staff member must catch before publication or deployment.
- Form other staff in beginning AI competence by selecting appropriate recipes, supervising initial use, providing feedback, and recognizing when a colleague has crossed from supervised practice into independent capability.
- Hold the human and theological dimensions of AI work — including questions of presence, formation, vocation, voice, and integrity — with sufficient depth to navigate situations the Safety documentation does not explicitly anticipate.
- Lead a use case adjudication conversation with peers and leadership — proposing a position, hearing pushback, modifying the proposal, and arriving at a decision the team can implement and the leadership can defend.
- Recognize the limits of one's own discernment and know when to escalate — to leadership, to Movemental, or to a peer with more relevant expertise — rather than proceeding past one's actual capacity.

Six-C: Course architecture

Skills Development is delivered as an eight-module course over eight weeks. Each module is structured around the same three-part rhythm: a formational reading or input, a practical exercise grounded in the participant's actual work, and a cohort or asynchronous discussion that sharpens the judgment the module is forming. Each module is approximately three to five hours of participant time per week.

The course architecture is built on Movemental's existing LMS infrastructure (React, Next.js, Tailwind, Supabase, with agentic tutoring and feedback capabilities). The course can be delivered in three formats: cohort-facilitated (eight weeks with a Movemental facilitator running weekly sessions of a

small group), self-paced (eight weeks of asynchronous progression with agentic tutoring and quarterly cohort moments), or hybrid (self-paced core with monthly facilitated cohort sessions for deeper discernment work).

Module 01: AI as human simulation: what we are actually working with

LEARNING OUTCOMES

- Articulate why current AI tools present as human simulation rather than as conventional technology, and why this changes the nature of the competence required to use them well.
- Describe the specific ways AI output mirrors human language, judgment, and authority — and where this mirroring is reliable, where it is misleading, and where it is dangerous.
- Distinguish between using AI as a tool, using AI as a collaborator, and using AI as a substitute — and recognize which mode is operating in any given workflow.

Module one establishes the conceptual frame for everything that follows. Most AI training treats AI as a sophisticated tool. This module argues — and demonstrates — that current AI is better understood as human simulation: it speaks our languages, mimics our reasoning, produces our cultural artifacts, and presents itself through the same interface as another human collaborator. The competence required to work with human simulation is not technical training. It is discernment about when the simulation is helpful, when it is misleading, and when something other than simulation is needed.

Practical exercise: each participant brings a recent piece of AI-generated work from their own role and analyzes it against this frame. Where did the AI succeed at simulating what was needed? Where did it fail? What did the participant's human judgment add that the simulation could not?

Module 02: Reading a use case in context

LEARNING OUTCOMES

- Identify the seven dimensions of context that bear on any AI use case — workflow, data, people, integrity, values, frequency, and downstream effect — and apply the framework to a use case from the participant's own work.
- Articulate why context is the irreducible human contribution to AI work, and what this means for how staff at every level should be formed.
- Recognize when a use case proposal is missing context the practitioner has assumed and the proposer has not surfaced.

This module trains the practitioner-to-discerner shift directly. A use case is never just "use AI to do X." It is always "use AI to do X, with this data, in this workflow, affecting these people, against these integrity standards, at this frequency, with these downstream effects." The participant who can read all seven

dimensions of context produces dramatically better adjudications than the participant who reads one or two.

Practical exercise: each participant adjudicates three use cases from a curated library — one green-light, one yellow-light, one red-light — and writes the reasoning that supports each adjudication. Cohort discussion compares adjudications and surfaces where participants read context differently.

Module 03: Modifying a recipe with discernment

LEARNING OUTCOMES

- Modify an existing Claude skill or recipe to fit a workflow that has evolved, by adjusting prompts, context, output structure, and integration points.
- Distinguish between modification that improves the recipe and modification that compromises the recipe's alignment with Safety documentation.
- Recognize the limits of one's own modification authority and know when to escalate to a more capable collaborator or to Movemental.

Module three is the most technically substantive in the course, and is also the module where the formational dimension matters most. Modifying a recipe well requires both the technical capacity to make the change and the judgment to know whether the change is wise. Participants learn to modify recipes responsibly — making them better at the work the organization actually does, while keeping them aligned with the values the organization has documented.

Practical exercise: each participant takes one recipe from their Sandbox engagement and produces a modified version tuned to a workflow that has shifted since the original Sandbox. The modification is reviewed by the cohort and the facilitator.

Module 04: Evaluating AI output: catching what the simulation misses

LEARNING OUTCOMES

- Evaluate AI output for quality, voice fit, contextual accuracy, subtle factual errors, and alignment with the organization's convictions.
- Describe the categories of error that current AI produces consistently, and the patterns that indicate output should not be used as-is.
- Develop a personal review discipline that catches the errors most relevant to one's own role.

AI output is fluent, plausible, and frequently wrong in ways that are difficult to catch without active discernment. Module four trains participants to read AI output with the kind of attention that catches what the simulation misses. The discipline is teachable, and the teaching is most effective when grounded in the specific kind of work the participant does.

Practical exercise: each participant is given a set of AI-generated outputs from their domain (for a pastor, AI-generated sermons and pastoral communications; for a development director, AI-generated donor appeals; for a faculty member, AI-generated theological writing). Each participant identifies the errors, the missing context, the voice mismatches, and the points where the output should not be used. Cohort discussion sharpens the eye.

Module 05: The human dimensions: presence, voice, vocation, integrity

LEARNING OUTCOMES

- Articulate the specific human dimensions of one's own work that AI cannot substitute for, and the reasons rooted in the organization's tradition and theology.
- Recognize the temptations created by AI assistance — toward efficiency over presence, output over relationship, scale over fidelity — and develop personal practices that resist them.
- Hold a conversation with a colleague who is over-using or under-using AI in ways that affect the work, in a way that deepens rather than strains the relationship.

This module is the formational core of the course. It is also the module where the segment-specific texture matters most — pastors process this material differently than development directors, who process it differently than seminary faculty. The course supports segment-specific cohorts where possible. The aim is not to argue against AI use but to develop the personal practices and theological clarity that make AI use sustainable in this kind of work over time.

Practical exercise: each participant writes a one-page personal commitment articulating where they will and will not use AI in their work, and why, grounded in their own theology, vocation, and the people they serve. The commitment is shared with the cohort and revisited at the end of the course.

Module 06: Forming others: from learning to teaching

LEARNING OUTCOMES

- Select an appropriate recipe for a colleague who is beginning AI work, and supervise the initial use in a way that builds the colleague's discernment alongside their capability.
- Provide feedback on AI-assisted work in a way that develops the colleague's judgment rather than simply correcting the output.
- Recognize when a colleague has moved from supervised practice into independent competence, and adjust supervision accordingly.

Module six trains participants to do for others what the course is doing for them. This is the multiplier in the Skills stage: each participant becomes capable not only of their own competent practice but of forming others in the same competence. Organizations that complete Skills do not depend on Movemental for ongoing development. They develop their own people.

Practical exercise: each participant identifies a colleague (real or hypothetical, depending on organizational context) and develops a one-month formation plan for that colleague — recipes to start with, supervision rhythm, feedback approach, signs of readiness for independent work. The plan is reviewed by the cohort.

Module 07: Leading adjudication conversations

LEARNING OUTCOMES

- Lead a use case adjudication conversation that produces a decision the team can implement and the leadership can defend.
- Hear pushback on a proposed adjudication without becoming defensive, and modify the proposal in response to substantive challenge while maintaining substantive convictions.
- Document an adjudication decision in a way that future staff can read, understand, and apply.

Most AI work in organizations does not break because of bad technology. It breaks because of unmade decisions, unspoken disagreements, and undocumented adjudications. Module seven trains participants in the conversational and decisional discipline that keeps AI work on track at the team level. This is leadership work, regardless of the participant's formal role.

Practical exercise: each participant facilitates a 30-minute adjudication conversation with their cohort on a use case the cohort has not yet considered. The conversation is recorded (or observed by the facilitator) and reviewed for the leadership skills the module names.

Module 08: Knowing your limits: escalation and ongoing formation

LEARNING OUTCOMES

- Recognize the situations that exceed one's current discernment, and know whom to escalate to and how.
- Develop a personal plan for ongoing formation in AI wisdom over the next 12-24 months as the technology and the field continue to evolve.
- Identify the resources, relationships, and rhythms that will sustain the participant's judgment as the conditions change.

Module eight closes the course by naming what the course cannot teach — and what the participant must develop on their own going forward. The conditions of AI work will continue to shift. The recipes will change. New use cases will emerge. New concerns will surface. The participant who has completed Skills is competent at this moment; staying competent requires ongoing formation. Module eight equips participants to lead their own formation rather than depending on Movemental indefinitely.

Practical exercise: each participant writes a personal ongoing formation plan — what they will read, who they will stay in conversation with, what rhythms of reflection they will maintain, when they will return to advanced learning. The plan is shared with the cohort and the facilitator. The course concludes with a closing conversation in which each participant names what they are taking with them.

Six-D: Skills technological supports

LMS infrastructure. The course is delivered through Movemental's existing LMS (React, Next.js, Tailwind, Supabase). Each module is structured with reading inputs, video or written explanations, practical exercises with submission, cohort or asynchronous discussion, and assessment. The infrastructure supports cohort-facilitated, self-paced, and hybrid delivery from the same source.

Agentic tutoring and feedback. Self-paced and hybrid participants receive agentic tutoring throughout the course. The agent reads participant submissions, provides initial feedback, surfaces patterns the participant should attend to, and escalates to a Movemental facilitator when human judgment is required. The agent does not replace the facilitator; it ensures that participants in self-paced learning still receive substantive engagement at scale.

Recipe sandbox environment. Module three (recipe modification) and several practical exercises require participants to actually modify Claude skills. The course provides a configured sandbox environment in which modifications can be tested without risk to live organizational systems. The sandbox connects to the participant's organization's Safety documentation so modifications are pre-screened for alignment.

Cohort interaction tools. For cohort-facilitated and hybrid delivery, the LMS supports synchronous video sessions, asynchronous threaded discussion, peer review of practical exercises, and the seventh-module facilitated adjudication conversations. The infrastructure is built to handle cohorts of 8-15 participants with one facilitator.

Six-E: Skills boilerplate templates

B1. PERSONAL COMMITMENT TEMPLATE (Module 5)

Personal Commitment Regarding AI Use in My Work

My role is [role] at [organization]. The work I am responsible for is [brief description of the work that defines my vocation in this season].

Where I will use AI in this work. [Specific list of use cases the participant has discerned are appropriate to their role, grounded in the organization's Safety documentation and the

participant's own theological and vocational clarity. Be specific. "Drafting initial outlines for routine communications" is more useful than "writing assistance."]

Where I will not use AI in this work. [Specific list of use cases the participant has discerned are not appropriate to their role. Again be specific. "Pastoral counseling conversations, including any summarization of these conversations" is more useful than "personal ministry."]

My reasoning. [One to three paragraphs. Why these lines? What is the participant's theological, vocational, or relational reasoning? What does it mean for the participant's work that AI is and is not allowed in these specific places?]

My personal practices. [What rhythms or disciplines will the participant maintain to sustain this commitment? Examples: reviewing this commitment quarterly; bringing AI-related questions to a specific colleague or supervisor; pausing before any new use case to evaluate against this commitment; etc.]

Date and signature. [Participant signs and dates. The commitment is shared with cohort and facilitator. The participant returns to it at the end of the course and considers what to revise.]

B2. ADJUDICATION DECISION DOCUMENTATION TEMPLATE (Module 7)

Adjudication Decision: [Use Case Short Name]

Date of decision. [Date.]

People involved in the decision. [Names and roles. The decision is documented as a team decision, not an individual one.]

Use case as proposed. [Description of the use case being adjudicated, including the workflow, the data, the people involved, and the value it would create.]

Context surfaced during adjudication. [What dimensions of context the team considered. Reference to the seven-dimension framework from Module 2.]

Concerns surfaced. [What concerns came up during the conversation, and how each was addressed.]

Adjudication. [Green / Yellow / Red.]

Reasoning. [Two to four paragraphs. Why this adjudication? What were the decisive considerations? What would the team say to a future staff member who proposes a similar use case and reads this document?]

If yellow: required revision and timeline. [What specifically must change before this use case can move to green status, and by when.]

If red: documentation for future reference. [Specific enough that a future staff member who proposes the same or a similar use case can read this entry, understand the prior decision, and either accept it or surface new information that warrants revisiting.]

Review cadence. [When will this adjudication be revisited? AI conditions change; adjudications should not be assumed to be permanent. Typical: yellow-light revisited in 90 days, green-light revisited annually, red-light revisited if the proposing party brings new information.]

DECISIONS THE ORGANIZATION MUST MAKE

- Which delivery format fits our organization — cohort-facilitated (highest engagement, fixed schedule), self-paced (highest flexibility, more individual responsibility), or hybrid (balanced)?
- Who from our organization will participate? (Skills is most valuable for staff who will lead AI work in the organization. Typically 5-12 participants per cohort. Choose for leadership trajectory and capacity for formation, not just current AI enthusiasm.)
- What is our intended outcome for the cohort? (Some organizations want capacity-building across many staff; others want a small group of deep practitioners who will lead the rest. The intent shapes who participates and how the course is delivered.)
- How will the cohort's learning be integrated back into the organization? (Skills is wasted if it stays inside the participants. The integration plan — how participants share learning, lead adjudications, form others — is itself a leadership decision.)
- What investment of participant time is realistic? (3-5 hours per week for 8 weeks per participant. Confirm before enrollment.)
- What is our cadence for ongoing formation after the course? (The eighth module produces a personal formation plan, but the organization should also have a collective rhythm — quarterly use case reviews, annual cohort reunions, ongoing facilitator engagement, etc.)

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Self-Paced License, \$4,800/year per organization):

The full course is licensed to the organization for self-paced delivery. Up to 25 participants per organization per year. Agentic tutoring throughout. Quarterly cohort moments where licensed-organization participants from across the network gather virtually for facilitated discussion. Annual update of course content as the field evolves. This is the volume tier — how Movemental serves organizations that want substantive formation at scale.

At Level 2 (Cohort Facilitation, \$15,000):

A Movemental facilitator runs an eight-week cohort for 8-15 participants from a single organization (or for multi-organization cohorts at the institution level). Weekly synchronous sessions, asynchronous discussion between sessions, facilitator review of practical exercises, and a closing conversation that integrates the cohort's learning into organizational planning. The standard offering for organizations engaging Movemental for Skills.

At Level 3 (Network Engagement Skills, from \$60,000 for institutions):

Skills can be delivered as part of a Network Engagement for institutions whose member organizations need formation across the network. Typical structure: institutional cohort plus 3-10 member-organization cohorts, with shared resources, cross-cohort learning moments, and a culminating institution-level integration. The institution becomes a teaching artifact for its network.

Add-on: Train-the-Facilitator certification, \$25,000:

For institutions and large organizations that want to deliver Skills internally, Movemental certifies internal facilitators through a 12-week certification program. Certified facilitators can run the Skills course inside their own institution under license, with annual recertification. This is the path for institutions that want Skills capacity to live inside their walls rather than depending on Movemental indefinitely.

Stage 04: Solutions

Part seven: Solutions Deployment

The Solutions stage is the technological deployment of what the upstream stages have made the organization ready to deploy. It is the building of agentic ecosystems, the optimization of existing tools like Claude Enterprise, the integration of AI capability into the organization's actual workflows, the construction of custom infrastructure where warranted, and the operational rhythm that sustains the deployed capability over time. It begins at \$30,000 for focused engagements and scales by scope.

The premise of the Solutions stage is that technology built on a coherent foundation compounds, while technology built on a fragmented foundation corrodes. The organization that arrives at Solutions has done the foundational work: Safety has been ratified, Sandbox has surfaced and adjudicated specific use cases, Skills has formed the staff who will operate and lead the deployed capability. Solutions is therefore not a build-from-scratch engagement. It is the deployment of clarity that has already been generated.

Seven-A: What Solutions actually deploys

Solutions Deployment is not one offering. It is a family of offerings calibrated to what Sandbox surfaced as the right next investment for this specific organization. Five typical Solutions configurations cover the majority of engagements.

Configuration A: Optimization of existing tools.

The organization's green-light use cases are deployed through the AI tools the organization already has — typically Claude Enterprise, Microsoft Copilot, or the equivalent enterprise-tier configuration of a major model provider. Movemental builds custom skills, prompts, and integrations within the existing tool environment. Documentation, staff training, and a 90-day operational support window are included. Pricing typically lands in the \$30,000-\$50,000 range. This is the right configuration for organizations whose Sandbox surfaced strong use cases that do not require new infrastructure.

Configuration B: Custom agentic ecosystem.

The organization's green-light use cases require integration across multiple tools, custom data handling, or sustained agentic workflows that exceed what off-the-shelf enterprise tools deliver well. Movemental builds a custom agentic infrastructure — typically combining Claude or comparable models, the organization's data systems (CRM, member database, donor system, content libraries), and custom workflow orchestration. The build is hosted on Movemental's infrastructure or the organization's, depending on data sovereignty requirements. Pricing typically lands in the \$80,000-\$150,000 range. This is the right configuration for organizations with significant scale or substantial multi-system workflows.

Configuration C: Content and intelligence platform.

For institutions, networks, and content-heavy organizations (large churches with substantial sermon archives; nonprofits with extensive program documentation; seminaries with faculty writing and curriculum), the deployment is a content and intelligence platform — the organization's archive made queryable, searchable, and active rather than static. The platform supports staff workflows that draw on the archive, member or constituent workflows that access the archive in defined ways, and ongoing additions that keep the archive growing. Pricing typically lands in the \$60,000-\$120,000 range plus annual platform licensing. This is the configuration that converts an organization's accumulated content into an active asset.

Configuration D: Network deployment (Institutions only).

For denominational bodies, training networks, and multi-site institutions, the deployment extends across multiple member organizations under a shared governance framework and shared platform. Each member organization receives a configured deployment; the institution holds the framework and the platform; the cost is shared across the network. Pricing begins at \$60,000 for the institutional

engagement and scales with the number of member organizations. This is the configuration that lets an institution lead its network into AI capability rather than watching members navigate independently.

Configuration E: Hybrid and bespoke.

Some organizations need a combination — optimization of existing tools for some workflows, custom agentic infrastructure for others, content platform integration for the archive, and network distribution for member organizations. These are scoped per conversation. Pricing reflects scope. Movemental does not over-engineer; we deploy what Sandbox actually surfaced as needed and supported by the organization's Skills capacity.

Seven-B: The readiness threshold

Movemental does not deploy Solutions for organizations that have not crossed a threshold. The threshold is named explicitly because most consulting and technology firms do not name it, and the absence of clarity is a primary cause of the AI deployment failures the broader sector is currently experiencing.

An organization is ready for Solutions Deployment when:

- Safety Documentation is in place — ratified by leadership, communicated to staff, and in active use as a reference for adjudications. Whether this work was done with Movemental or independently does not matter, but it must be done.
- Sandbox Discovery has surfaced specific, named, adjudicated use cases. Solutions does not build for hypothetical value. It deploys for documented value with documented integrity assessment. Without Sandbox, Solutions is building on speculation.
- Skills Development has produced — or the organization already possesses — the staff capacity to operate and lead what is being deployed. Technology deployed without the human capacity to lead it does not produce value; it produces dependency on the vendor that built it. Movemental does not want this dependency, and we do not configure engagements that produce it.
- Leadership has clarity about what is being deployed and why. The Discernment Memo from Sandbox is the typical artifact that establishes this clarity. Without it, Solutions becomes an exercise in shifting goals during the engagement — expensive for the organization and frustrating for everyone involved.

Organizations that have done some but not all of this upstream work are not refused; they are routed. Movemental will help an organization complete the upstream work first, then deploy Solutions. The path is the path. The threshold is real.

Seven-C: The Solutions process, by phase

Phase 1: Scoping (2-3 weeks).

Movemental works with leadership to translate the Sandbox findings and Discernment Memo into a specific Solutions configuration. The scoping phase produces a Solutions Brief that names the configuration (A, B, C, D, or E), the specific deliverables, the integration points with the organization's existing systems, the staff who will operate the deployed capability, the timeline, and the price. The brief is signed before build begins.

Phase 2: Build (4-12 weeks depending on configuration).

Movemental engineers the deployment. For Configuration A (optimization), this is typically 4-6 weeks. For Configuration B (custom agentic), 8-12 weeks. For Configuration C (content platform), 8-12 weeks plus ongoing migration. For Configuration D (network), 12-16 weeks plus per-member rollouts. The build phase includes weekly check-ins with leadership and key staff, milestone reviews, and progressive testing against real workflows rather than artificial demonstrations.

Phase 3: Deployment and training (2-4 weeks).

The deployment goes live progressively rather than as a single launch event. Staff who completed Skills lead the rollout to additional staff. Movemental provides documentation, role-specific training materials, troubleshooting reference, and active support during the rollout window. The deployment is considered complete when the organization can operate the capability without daily Movemental involvement.

Phase 4: 90-day operational support.

After deployment, Movemental remains engaged for 90 days of operational support — addressing issues as they arise, adjusting configurations based on real-world friction, supporting the organization through the inevitable surprises that emerge once a deployment meets actual operational reality. After 90 days, the engagement transitions to either an ongoing platform subscription (for hosted infrastructure) or a defined annual retainer for periodic optimization (for organization-hosted deployments).

Seven-D: Solutions technological architecture

Movemental's Solutions deployments are built on a consistent technical foundation that adapts to the configuration.

Model layer. Anthropic's Claude is the default model. Other major-provider models can be integrated where the use case warrants, with the corresponding data handling and contractual review. Model choices are documented in the Solutions Brief and revisited as the field evolves.

Application layer. Movemental's existing React, Next.js, and Tailwind front-end stack. Supabase for authentication, database, and storage. Custom agentic orchestration built on Movemental's internal framework. The application layer is consistent across deployments, which is what allows Movemental to build deeply customized solutions in the timelines we name.

Data layer. Each organization's data remains in systems the organization controls (their CRM, member database, donor system, content management system, etc.) with API integrations rather than data migration. This is a deliberate architectural choice: Movemental does not become the custodian of organizational data unless explicitly engaged to do so under a hosted-platform configuration. Most deployments leave data where it lives and integrate around it.

Skill layer. The Claude skills (recipes) the organization developed and adjudicated through Sandbox become the operational layer of the deployment. Solutions does not invent new use cases; it deploys what Sandbox proved. The skills are versioned, documented, and tied to the Safety documentation that governs them.

Governance layer. Every deployment includes a governance layer that ties the technical capability back to the Safety documentation. Acceptable Use Policy violations are detected and surfaced. Data Handling Standards are enforced at the technical level where possible. Disclosure standards are applied to outputs that warrant them. The governance layer is what makes Solutions sustainable: the deployment cannot drift from the foundational decisions because the foundation is wired into the system.

Seven-E: Solutions boilerplate templates

C1. SOLUTIONS BRIEF TEMPLATE

Solutions Brief: [Organization Name]

Date. [Brief signing date.]

Configuration. [A, B, C, D, or E with brief description.]

Upstream readiness. [Confirmation that Safety documentation is ratified, Sandbox Discovery has produced specific adjudicated use cases referenced in this brief, and Skills capacity exists to operate and lead what is being deployed. Reference specific Sandbox use cases by Use Case Portfolio entry numbers.]

Specific deliverables. [Itemized list of what will be built and deployed. Each item references the green-light use cases from Sandbox that motivate it. No deliverable exists in this brief that does not trace back to documented Sandbox findings.]

Integration points. [The systems this deployment integrates with: CRM, member database, donor platform, content management, communication tools, etc. For each, the integration approach: API, file-based, webhook, etc.]

Staff who will operate the deployment. [Named staff and their roles. These staff have completed Skills Development or possess equivalent capacity. The brief names them by role and confirms their availability for the rollout phase.]

Governance integration. [Reference to specific Safety documents that govern this deployment, and how the technical layer enforces or supports them. Acceptable Use Policy alignment, Data Handling Standards adherence, Disclosure standards integration, Care Boundaries enforcement.]

Timeline. [Phase 1 scoping (already underway), Phase 2 build dates, Phase 3 deployment dates, Phase 4 operational support window. Total elapsed time and key milestones.]

Price. [Total engagement price, payment schedule, scope boundaries.]

Out of scope. [Explicit list of items that have been considered and are not part of this engagement. Naming what is out of scope is part of what makes the engagement scoped rather than open-ended.]

Operational support transition. [What happens at the end of Phase 4. Hosted platform subscription, annual retainer, or full handoff. Pricing for ongoing relationship if applicable.]

Signed by. [Organization signatory and Movemental signatory.]

C2. DEPLOYMENT GOVERNANCE INTEGRATION DOCUMENT

Deployment Governance Integration: [Deployment Name]

Purpose. This document records how the technical deployment is integrated with the organization's Safety documentation. It is referenced by the operations staff who run the deployment and the leadership who oversee it.

AI Use Charter alignment. [Brief description of how this deployment expresses the principles in the organization's charter. Specific charter principles referenced.]

Acceptable Use Policy enforcement. [How the deployment enforces or supports the policy. What automated checks are in place. What human reviews are required at what points.]

Data Handling Standards adherence. [What data flows through the deployment, in which categories, with what protections. Specific reference to the organization's data classification.]

Disclosure and Attribution Standard integration. [How the deployment supports the organization's disclosure practices. What outputs receive disclosure annotations automatically. What outputs require human disclosure decisions.]

Care Boundaries enforcement. [What the deployment cannot do, by design. What sacred work the deployment is structurally prevented from entering. How attempts to cross boundaries are surfaced for review.]

Voice cloning and impersonation protections. [How the deployment guards against being used for impersonation. What output authentication is in place. How constituent verification is supported.]

Constituent communication. [How the existence and operation of this deployment is communicated to the organization's constituents, consistent with the organization's prior communications about AI use.]

Review cadence. [Annual review of governance integration, with revision as Safety documentation evolves or technology changes.]

DECISIONS THE ORGANIZATION MUST MAKE

- Have we completed Safety, Sandbox, and Skills (or possess equivalent)? If not, what is the path to readiness, and is Movemental the right partner for that work or is another path more appropriate?
- Which Solutions configuration matches what our Sandbox surfaced? (A through E.) The Discernment Memo from Sandbox should already point clearly to the right answer; if it does not, additional Sandbox work may be appropriate before scoping Solutions.
- What is our timeline expectation, and is it realistic given the configuration? (Configuration A: 8-13 weeks total. Configuration B: 14-19 weeks. Configuration C: 14-19 weeks plus migration. Configuration D: 20-24 weeks.)
- What is our budget envelope, and does it match the configuration? Configurations have real cost floors; under-funded engagements either fail or compromise quality.
- Who in our leadership owns this engagement on the organization's side? Solutions requires sustained leadership attention through the build phase; without an executive sponsor, deployments drift.
- What is our intent for the relationship after the 90-day operational support window? Hosted platform, annual retainer, full handoff?

- What is our internal communication plan for the deployment? Constituents will notice; staff will adjust; the deployment is more sustainable when communicated proactively rather than encountered as fait accompli.

MOVEMENTAL'S ROLE BY LEVEL

At Level 1 (Scoping consultation, \$5,000):

Organizations that have completed Sandbox elsewhere or independently can engage Movemental for a focused 2-3 week scoping engagement that produces a Solutions Brief without proceeding to build. The brief is delivered, the organization can take it to any qualified vendor for build, or return to Movemental. This is the entry tier for organizations that want Movemental's discernment without committing to Movemental's deployment. Pricing is credited toward a subsequent build engagement if the organization returns within 90 days.

At Level 2 (Configuration A — Optimization, \$30,000-\$50,000):

Optimization of existing enterprise AI tools (Claude Enterprise, Microsoft Copilot, or equivalent). Custom skills, prompts, and integrations built within the tool environment. Documentation, staff training, and 90-day operational support included. Suitable for organizations whose Sandbox surfaced strong use cases that do not require new infrastructure.

At Level 3 (Configuration B — Custom Agentic Ecosystem, \$80,000-\$150,000):

Custom agentic infrastructure built on Movemental's technical foundation. Multi-system integration, custom workflow orchestration, hosted on Movemental or organization infrastructure depending on data sovereignty requirements. 90-day operational support plus ongoing platform subscription or retainer.

At Level 4 (Configuration C — Content and Intelligence Platform, \$60,000-\$120,000 plus annual licensing):

Archive-active platform that converts the organization's accumulated content into a queryable, workflow-integrated asset. For institutions, content-heavy churches, and content-heavy nonprofits.

At Level 5 (Configuration D — Network Engagement, from \$60,000 institutional plus per-member scope):

Network deployment for denominational bodies, training networks, and multi-site institutions. Institutional framework plus member-organization deployments under shared governance. The configuration that lets an institution lead its network into AI capability.

Ongoing relationship after deployment:

Hosted platforms include annual platform subscription (typically \$15,000-\$50,000/year depending on configuration). Organization-hosted deployments transition to an annual retainer for periodic optimization (typically \$10,000-\$30,000/year) or full handoff with as-needed engagement. The relationship after deployment is a deliberate choice in the Solutions Brief, not an afterthought.

Part eight: cross-stage productization and integration

This part documents the strategic logic across the four stages of the path and provides the integration view that ties Safety, Sandbox, Skills, and Solutions into a coherent operational and commercial structure.

Eight-A: The full path

The Movemental AI Path is one path, expressed in four stages. Each stage presupposes the work of the prior stage. Each stage produces specific artifacts that the subsequent stage can build on. The path is not a menu where customers pick favorites; it is a sequence with logic. At the same time, organizations enter the path at different points, hold different paces through it, and can pause between stages when their season requires it. The path is firm in structure and accommodating in pace.

- **Safety.** Two-week MVP. Seven artifacts. The non-negotiable foundation. Free Toolkit, AI-Assisted Product (~\$1,997), or Facilitated MVP (\$5,000).
- **Sandbox.** Four weeks. Use Case Portfolio plus Discernment Memo plus Readiness Assessment. The exploratory disciplined discovery. Self-Service Toolkit (~\$497), AI-Assisted Sandbox (~\$3,997), or Facilitated Sandbox Discovery (\$15,000).
- **Skills.** Eight weeks. Eight modules. Practitioner-to-discerner formation. Self-Paced License (\$4,800/year per organization), Cohort Facilitation (\$15,000), or Train-the-Facilitator certification (\$25,000).
- **Solutions.** 4-24 weeks depending on configuration. Five configurations from optimization to network deployment. Pricing from \$30,000 to \$200,000+ depending on scope. The technological deployment built on the upstream foundation.
- **Network Engagement.** From \$60,000. Multi-stage path delivered across an institution and its member organizations. The configuration for denominational bodies, training networks, and multi-site institutions.

Eight-B: The full pathway price

The full Movemental AI Path priced separately is approximately \$65,000 (\$5K Safety + \$15K Sandbox + \$15K Skills + \$30K Solutions floor). Bundled as a complete path engagement, the price is \$58,500 — a 10% discount that reflects engagement efficiency rather than discount-for-discount's-sake.

Movemental does not require organizations to commit to the full path at intake. Most engagements begin at one stage and continue based on what emerges. The organization that completes Safety with us

is more likely to engage Sandbox; the organization that completes Sandbox is overwhelmingly likely to continue to Skills or Solutions. The path sells itself by working.

Eight-C: The productization strategy in summary

Movemental productizes the path at three levels, with one level adapted per stage where the work warrants:

Free or low-cost self-service tier. Field Guide (free) for Safety. Sandbox Toolkit (~\$497) for Sandbox. Self-Paced License (\$4,800/year) for Skills. Scoping consultation (\$5,000) for Solutions. The self-service tier serves the broadest market and builds top-of-funnel credibility.

AI-Assisted product tier. Safety AI-Assisted (~\$1,997). Sandbox AI-Assisted (~\$3,997). Skills Self-Paced (which is the AI-assisted equivalent for this stage). Solutions does not have an AI-assisted tier because deployment work is irreducibly facilitated. The AI-assisted tier captures the middle market and demonstrates Movemental's thesis through product.

Facilitated tier. Safety Facilitated MVP (\$5,000). Sandbox Facilitated Discovery (\$15,000). Skills Cohort Facilitation (\$15,000). Solutions Configurations A through E (\$30,000 to \$200,000+). The facilitated tier serves organizations that want speed, external coordination, and a sustained relationship with Movemental.

Network tier (institutions only). Network Engagements starting at \$60,000 that span multiple stages and multiple member organizations under a shared institutional framework.

Eight-D: The integration view

The four stages produce artifacts that compound. Safety produces the seven governance artifacts plus the constituent communication. Sandbox produces the Use Case Portfolio plus the Discernment Memo plus the Readiness Assessment, all referencing Safety. Skills produces personal commitments, adjudication documentation, and ongoing formation plans, all building on Sandbox findings and Safety documentation. Solutions produces the deployment, the deployment governance integration document, and the operational handoff, all explicitly grounded in everything upstream.

At the end of the full path, the organization holds approximately 18-22 working artifacts that together constitute its AI infrastructure: governance documents, adjudicated use cases, formed staff, deployed capability, and a documented integration that connects all of it. This infrastructure is durable. It survives staff turnover. It survives technology changes. It survives the organization's next season of work. It is, quite literally, the foundation Movemental promised at the beginning of the path.

Eight-E: The build sequence for Movemental over 12 months

If Movemental commits to building all four stages at all three productization levels over the next 12 months, the recommended sequence is:

Months 1-2: Refine Safety. Publish the Field Guide. Build the AI-Assisted Safety Product v1. Continue offering Facilitated MVP. Approximately 110 hours of focused build time given existing infrastructure.

Months 3-5: Build Sandbox infrastructure. Develop the Use Case Portfolio interface. Build the assessment artifact (Claude artifact). Configure the initial recipe library for churches, nonprofits, and institutions. Publish the Self-Service Toolkit. Build the AI-Assisted Sandbox tier. Approximately 200 hours of focused build time.

Months 4-8 (overlapping with Sandbox): Build Skills course. Develop the eight-module course content on the existing LMS infrastructure. Build the agentic tutoring layer. Configure the recipe sandbox environment. Pilot with one cohort. Refine. Open Self-Paced License. Approximately 350 hours of focused build time.

Months 6-12 (overlapping with Skills): Solutions infrastructure maturation. Refine the Solutions architecture based on early engagements. Build the governance integration layer as a reusable component. Codify the five configurations as engagement templates. Train additional facilitators (if growth supports). Approximately 250 hours of focused build time, plus engagement-specific work.

Total approximate focused build time across 12 months: 900 hours, distributed across the calendar with overlap. This is one full-time engineer-equivalent for the year, plus engagement-specific work delivered as engagements occur.

Eight-F: Update protocol for the full path

This document is the single source of truth for all four stages. When boilerplate language is improved, when assessment questions evolve, when course content is refined based on cohort experience, when Solutions configurations are updated based on engagement learning — the update happens here first. Downstream artifacts regenerate from this source.

Quarterly review of the full document by the Movemental team. Annual update of the field paper sections (Parts One and similar) as new sector data becomes available. Continuous update of the working layers as engagements produce new learning.

The path is the path. This document keeps the path coherent.

Sources

Primary sources cited in Part One, with publication dates as of May 2026:

- Exponential AI NEXT and AlforChurchLeaders.com. The 2026 State of AI in the Church Survey Report. Published April 2026.
- Lifeway Research. Pastors and Churchgoers See AI as Concerning and Confusing. Survey of 1,003 Protestant pastors and 1,200 churchgoers conducted September 2025; published April 2026.
- Virtuous and Fundraising.AI. The 2026 Nonprofit AI Adoption Report. Survey of 346 nonprofits conducted late 2025; published February 2026.
- TechSoup. 2025 State of AI in Nonprofits Benchmark Report.
- Forvis Mazars. AI Governance as Fiduciary Responsibility. Published February 2026.
- Grant Thornton. 2026 AI Impact Survey. Survey of 950 organizations.
- Compliance Week and konaAI. AI & Compliance Survey 2026.
- Theology and Science. Generative AI and Theology: A Three-Year Retrospective. Vol. 24, No. 1, 2026.
- WIRED. AI Deepfakes Are Impersonating Pastors to Try to Scam Their Congregations. Published January 2026.
- Baptist News Global. AI is Coming for Your Pastor and Your Trust. Published January 2026.
- FBI public warning on AI-supercharged scams, late 2025.
- AI Incident Database, Incident 1355.
- Saint Meinrad Seminary, Columbia Theological Seminary, Puritan Reformed Theological Seminary, and Austin Presbyterian Theological Seminary published AI policies, 2024-2026.

Compiled by Movemental as the working source-of-truth for the Safety stage of the Movemental AI Path. Citations and statistics in this document are accurate as of May 2026 to the best of our research at the time of writing. Sectors are evolving rapidly; figures should be re-validated annually. Boilerplate text is provided for use by organizations served by Movemental and by the broader sector under the same terms as the published field guide.