

MOVEMENTAL

FIELD GUIDE · VOLUME TWO

It Continues With Exploration.

A field guide for the AI Sandbox.

Brad Brisco · Alan Hirsch · Joshua Shepherd

Edition 1.0 · Movemental · 2026 · movemental.ai

INTRODUCTION

Why we continue here.

This is the second field guide in the Movemental series. The first was *It Starts With Safety* — a guide to building your AI Organizational Guidebook. This one is what continues after.

You cannot run a Sandbox without Safety. The Guidebook gives the Sandbox its boundaries. But Safety also flows naturally into Sandbox, because there is no point developing safety guidelines without a field in which to safely explore. And there is no honest way to determine AI's value to your organization in the abstract. Value gets discovered through actual use, against actual work, by actual people in your context.

The Sandbox answers the question Safety raises: how are we actually going to train people to be safe with AI? And it extends to the question behind every AI conversation in your organization — is this thing really powerful, and can it help us? The Sandbox is where you find out.

The discoveries you are looking for will fall into three categories. **Efficiency** — time saved on work that already exists. **Revenue** — new or sustained funding that AI helps unlock. **Work quality** — the work itself made better. Those are the value axes the Sandbox measures. Nothing else.

The Sandbox is a place to play, learn, and discover — safely. It cannot be turned into more than that. Or less.

How to use this Field Guide.

This guide walks you through the eight phases of the Sandbox: **Boundaries, Assessment, Experimenting, Iteration, Reflection, Ethics Review, Discerning, and Future Plan**. Each phase has a clear purpose and a documented output. By the end of the eight phases, your team has tried AI against real work, documented its value, surfaced the ethical questions, and produced a decision-ready Future Plan your board can act on.

If you want facilitation, Movemental's **SandboxLive** engagement runs the Sandbox with your team — roughly ten hours of in-person teaching per cohort, supported by a learning management system, a library of pre-tested AI use cases customized to your context, and ongoing support throughout the discovery. The two paths section later in this guide compares both options.

AUTHORS' NOTE

Where this work comes from, and what we are inviting you into.

AI will mirror and amplify the humanity it finds sitting before it in any given interaction.

That is the sentence beneath everything in this Field Guide, as it was beneath the first. The question facing mission-driven organizations is whether your organization's already-existing goodness — the convictions you hold, the people you have formed, the trust you have built, the credibility you have earned over decades — will be responsibly translated into the age of AI that has arrived inside your work.

Safety gave you the Guidebook — the standing reference for what AI is for in your work, what it is not for, and what your organization commits to refuse on principle. The Guidebook is necessary, but on its own it is words on paper. Without a place to actually try AI against your work, the Guidebook cannot become practice. That is what the Sandbox is for.

How we got here.

Brad and Alan have spent decades doing missional and ecclesial work — planting churches, training leaders, writing the frameworks that have shaped how a generation of practitioners think about movement, formation, and the recovery of mission. Send Network, Forge, *The Forgotten Ways*, 5Q, APEST, 100Movements, Movement Leaders Collective: that body of work predates AI, and its credibility predates AI.

Josh came to this work from inside it — years as a Methodist pastor, then as founder of a neomonastic, communal nonprofit where he and his wife lived alongside young adults in extended residency — and from inside that practice began building the digital infrastructure for what comes next.

We began collaborating in 2024 around a shared concern: that the present moment in AI would happen to mission-driven organizations rather than be navigated by them. The answer would not come from AI consulting firms, because the answer is not primarily about AI. The answer is about whether organizations whose work depends on human covenant can carry that covenant into an age of machine fluency without compromising it. The Sandbox is the second step of that path.

What we are inviting you into.

A wise and mature exploration of what AI can actually do for your work. Not a debate about AI in the abstract. Not a vendor demo. Not a strategy memo. A real attempt, with real people, against real work, inside the boundaries your Guidebook has already established.

Everything starts with Safety. It continues with Exploration. What you discover here is what will eventually become the formation work of Skills and the integration work of Solutions. But all of that depends on this stage being done honestly. The Sandbox is where AI either earns its place in your organization or doesn't.

CONTENTS

What you'll find here.

	Introduction — Why we continue here	02
	Authors' note	03
	Why Sandbox — the case for bounded exploration	06
	What you're working with — Claude, up close	09
	What this looks like across your work — eighty use cases	13
	How this work gets done	17
	The data — what the research shows	19
	The eight phases of the Sandbox	20
01	Boundaries	21
02	Assessment	23
03	Experimenting	24
04	Iteration	27
05	Reflection	28
06	Ethics Review	29
07	Discerning	31
08	Future Plan	33
	Learning outcomes and course outline	34
	Self-assessment	39
	Two paths into Sandbox — SandboxGuide vs. SandboxLive	41
	About Movemental	46
	Citations	47

WHY SANDBOX

The case for bounded exploration.

AI in the form most of us are now using — large language models you talk with in plain English — has existed for roughly three years. ChatGPT was released in late 2022. The capabilities that followed were not predicted by the field that built them. They emerged. And they emerged in a form none of us were trained to evaluate.

The technology is a black box. Even the people building it do not fully understand why it produces the outputs it produces. The idea of having a conversation with your computer — talking back and forth in writing, asking it to think with you, getting answers that sometimes feel uncanny — is completely new for almost all of us. There is no inherited wisdom about it. All the insight any of us has is available at the point of interaction.

This matters for how you approach AI in your organization. You cannot reason your way to whether AI is valuable for your work. You cannot read your way to it either. The only way to know is to try — carefully, within boundaries, with documentation, against your actual work. That is what a Sandbox is for.

AI is not primarily a technology problem.

This is the second thing to understand. AI is not a technology problem in the way most organizational technology has been a technology problem. There is no IT integration to configure. There is no enterprise rollout. There is no command-line skill required.

If a member of your staff can open an application on a computer, they can use Claude, or ChatGPT, or any major AI tool. The user interface is a text box. You write, the model responds, you write back. That is the entire interaction model. AI at the consumer level is a communication tool, not a technical artifact. The skill of using it is a communication skill, not a technical skill.

This is also why a Sandbox is the right approach. You are not teaching your team a new software platform. You are giving them a structured way to develop a new communication practice with a new kind of partner. The Sandbox is the playground in which that practice develops.

Skill in AI use is mostly about clear communication.

Most of what gets called “prompt engineering” is overcomplicated. The skill of communicating well with an AI model is mostly basic professional communication — the same skills you bring to briefing a thoughtful new hire — with a small amount of specific awareness about what the tool is. Five things matter.

Clear communication. Saying what you actually want. Naming the audience, the format, the constraints. Treating the model as a capable collaborator who needs the same kind of instructions any new hire would

need on their first day.

Context. Telling the model what it needs to know to do the work well. Not all of your context — just the parts that are relevant. Who the audience is. What the tone should be. What's been tried before. What the constraints are. You do not need to fine-tune or engineer this. Say what a thoughtful person would say if they were asking a knowledgeable colleague to help with the same task. If you find yourself wanting to attach a few reference documents or paste in a prior example, that's the right instinct. Most of what gets called “prompting” is just giving the model enough context to do the work.

Definition of done. Knowing what a finished version of the thing you're asking for would look like, and being able to describe it. Most AI failures come from asking for vague things and getting vague answers. “Write me a fundraising appeal” is vague. “Write me a 400-word fundraising appeal to lapsed mid-level donors in the tone of a personal letter from our executive director, focused on our youth program's recent expansion, ending with a specific gift ask” is a definition of done. The second version produces dramatically better output not because the model is smarter but because you have given it more to work with.

Iteration. The first response is rarely the final response. Read what the model produced, identify what's working and what's not, and ask for changes. “More concise.” “Less corporate.” “Add a specific example here.” “This paragraph sounds like a different organization — rewrite it in our voice.” Iteration is where you train your eye to see what to ask for, and where the model gets the chance to actually produce what you want. Treating the first output as final is the most common reason people decide AI “doesn't work for them.” It usually does work — they just stopped before it got good.

Agentic perspective-taking. Understanding what the tool actually is on the inside. A language model has a training cutoff date and doesn't know what happened after it. It has no persistent memory between conversations. It can be confidently wrong. It will guess rather than say it doesn't know unless you give it permission to admit uncertainty. Knowing these things — imagining what it is like to be the model on the other side of the conversation — changes how you communicate with it. The next chapter unpacks what the model actually is, in more depth.

Beyond those five: pretend the AI is a person while remembering it is not. Technique without wisdom produces over-trust; wisdom without technique produces refusal. Both together produce play. It is okay to play. Treat your Sandbox time, temporarily, like a video game — something you are allowed to be curious about, allowed to fail at, allowed to enjoy while you learn.

What is lost if you skip Sandbox.

Without a Sandbox, you cannot honestly assess AI's value to your organization. Without that assessment, you have no real basis for the ethics conversation — there is no point debating the ethics of something whose value has not been established. And without value established and ethics debated, your team is functionally on its own with whatever AI tools they decide to use individually.

The Safety stage produced a Guidebook. Without a Sandbox, that Guidebook is words on paper without a practice to back them up. Your staff cannot actually *do* safety in their work if they have not been given a structured way to learn what AI is for in your specific context. Sandbox is where Safety becomes practice. Without it, Safety is a document. With it, Safety is a working frame your team has internalized.

There is one more piece. The benefit of the doubt belongs to mission-driven organizations to give a good reason for *not* exploring AI's potential value to their work. If AI could save your team five hours a week on routine work that no one wants to do — and redirect those hours to the relational, formative, mission-aligned work your organization exists to do — that is a stewardship question. Refusing to find out is not a neutral position. It is a decision about how the time of the people in your care is spent.

The 95 percent.

MIT's NANDA Initiative published a study in July 2025 finding that 95 percent of enterprise generative AI pilots fail to generate measurable financial returns.⁴ Boston Consulting Group's September 2025 study of 1,250 firms found that 60 percent generate no material value from AI investment at all.⁴

These numbers are widely cited as a warning that AI doesn't work. Read more carefully, they say something different. They say that 95 percent of organizations skipped the discovery work. They bought tools and ran pilots without first knowing what their actual work was, what their actual hypotheses were, or what value would look like if it appeared. Then they failed and concluded the technology was the problem.

The five percent that succeeded did the discovery work. They mapped their actual work. They tried specific use cases against it. They documented value carefully. They iterated. The Sandbox is what that five percent did. The Field Guide you are reading is the structured version of it.

YOUR WORKING PARTNER

What you're working with.

The previous chapter named the basic skills of communicating with an AI assistant. This chapter unpacks what the assistant actually is, so that perspective-taking is grounded in something specific. The Sandbox uses Claude, made by Anthropic, as its default AI assistant. Most of what follows applies to any frontier AI assistant — the model class is broadly similar across providers — but the specifics are for Claude Opus 4.7, the most advanced model Anthropic offers in May 2026.

Imagine a genie.

An imperfect metaphor, but useful as a starting point. Claude is, at the level of interaction, like a genie who answers your wishes — but only through language. You describe what you want; Claude produces what you described. The genie is powerful: enormously knowledgeable, fast, patient, tireless. The genie is also literal-ish. It gives you what you asked for, which sometimes isn't what you meant. The skill is in the wish.

Where the metaphor breaks down: Claude isn't malicious or tricky, doesn't mind being asked twice, and won't hold a wish against you. You can converse with it, refine your wish, ask for alternatives, ask it to think out loud about how it interpreted the wish in the first place. The genie is bounded only by what language can actually do.

Which is, in 2026, a remarkable amount. Language inside a computer can become code that runs. It can become a structured document, a spreadsheet, a slide deck, an image, an audio file, a translation, a summary, a search query against vast knowledge stores. It can control other software when given permission to. Language is the universal interface to digital work. Whatever you can describe carefully, Claude can attempt. Whatever Claude produces, your software can usually consume. That is the underlying capability — not just text in, text out, but language as the bridge between human intention and computational power.

What Claude is, at a basic level.

Claude is the smartest linguist in the history of the world, sitting at your desk as a servile personal assistant with a few additional traits that make the picture work. That is not metaphor; that is roughly the right operational picture. Claude has read an enormous portion of human writing — books, articles, academic papers, code, conversations, poetry, theology, business documents, journalism, law, manuals — across most major languages. It can read, write, summarize, translate, compare, critique, code, reason about, and converse in any of them.

Beyond linguistic ability, four traits matter. **Massive working memory.** Claude can hold roughly 500 pages of text in a single conversation — enough for a book-length document, several hours of meeting transcripts, or a stack of policy documents — and reason across all of it at once. **Speed.** A page-length response arrives

in seconds; a complex multi-page analysis arrives in under a minute. **Tireless patience.** The thousandth question is treated with the same care as the first. Claude doesn't get tired, doesn't get offended, doesn't hold grudges. **Compliance.** Within ethical limits, Claude does what you ask. It is not a colleague with its own agenda; it is a remarkably capable assistant whose job is to be useful to you.

What Claude isn't: alive, conscious, present, or persistent. Claude has no memory of you between conversations. Claude doesn't experience your work. Claude can be confidently wrong about facts, especially recent ones. Claude can be subtly sycophantic, agreeing with framings it should push back on, unless you explicitly ask for honest critique. Claude is not your friend, your advisor, your therapist, or your colleague. It is an instrument. A powerful one.

A note on Claude's training.

What we can know publicly: Claude is a large language model trained on a very large corpus of text. Its training involves two main phases. First, learning patterns across the entire corpus — this is what gives Claude its linguistic and factual breadth. Second, training to be helpful, harmless, and honest based on Anthropic's published principles. The result is a system that has read enormously and has been shaped to be useful and safe in conversation.

What we can't know publicly: exactly what was in the training corpus, exactly how the model came to behave the way it does, or precisely why it produces a particular output in a particular case. The model is, in this sense, a black box even to the people who built it. This is part of why the Sandbox is the right approach — you cannot reason your way to what Claude can do for your work. You can only try.

Claude Opus 4.7's training cutoff is the end of January 2026. That means Claude knows the world up to that date and doesn't know what has happened since unless you tell it, or unless it uses a tool like web search to find out. When you ask Claude about something recent, it may guess, it may say it doesn't know, or it may suggest you check current sources — depending on how the question is framed.

What Claude can do beyond chat.

The text-in-text-out picture is true at the simplest level but undersells what Claude has become. As of May 2026, Claude operates as the intelligence inside a number of surfaces, each of which extends what "talking to Claude" means in practice.

Claude.ai (the chat interface). The web and mobile chat where most people first meet Claude. Conversation, document analysis, drafting, research, light coding, image analysis.

Artifacts. Interactive code or documents that Claude generates and renders inline in the chat — web pages, calculators, data visualizations, dashboards, prototypes, formatted documents. You can iterate on them in conversation. Artifacts can store data that persists across sessions, which means you can build small functional tools, not just static outputs.

Projects. Organized workspaces inside Claude where you upload reference files, set custom instructions about how Claude should behave in this project, and reuse the setup across many conversations. A Project is the right home for ongoing work — a long writing project, an organizational knowledge base, a recurring task that always needs the same context.

Connectors. Plug-ins that let Claude access your other tools through a standard called the Model Context Protocol, or MCP. Connectors exist for Google Drive, Gmail, Google Calendar, Slack, Notion, Salesforce, Asana, Linear, and a steadily growing list of third-party services. Through connectors, Claude can read your files, draft your messages, schedule your meetings, query your databases, or take other actions in your existing tools — with whatever permissions you grant.

Skills. Specialized capabilities packaged as files Claude can use. Examples include creating proper Microsoft Word documents, filling PDF forms, building Excel spreadsheets, producing print-quality PDFs, and rendering high-design web interfaces. Skills extend what Claude can produce beyond what raw text generation alone would allow.

Web search and web fetch. Claude can query the public web for current information and read specific URLs you point it to. This addresses the training-cutoff problem for any task involving recent events or live information.

Code execution. In appropriate contexts Claude can actually run code, not just write it. This is what powers data analysis on uploaded spreadsheets, document creation, image manipulation, and similar tasks that require computation rather than just description.

Claude in Chrome, Claude in Excel, Cowork, Claude Code. Companion products that put Claude into specific work environments — your browser, your spreadsheet, your desktop, your terminal. Each one is a different surface for the same underlying intelligence.

Context, files, and the window.

Claude works inside what is called a **context window** — the amount of text Claude can hold in working memory at once during a conversation. Opus 4.7's context window is roughly 200,000 tokens, which corresponds to about 150,000 words or 500 pages of typical text. Everything in a conversation — your messages, Claude's responses, any files you have uploaded, any tool results — fills the window.

When the window fills, Claude starts to lose track of earlier parts of the conversation. There are two ways to manage this. Start a fresh conversation when you switch to a new topic, which clears the window. Or, in longer conversations, summarize what has been established so far and continue with a tighter context. For most Sandbox work, one conversation per experiment is plenty.

Files. You can attach PDFs, Word documents, spreadsheets, images, and most common file types to a conversation. Claude reads them and uses them as context. There are size limits — typically thirty megabytes per file, with reasonable limits on total uploads per conversation — but for almost all

organizational work those limits are well beyond what you will hit. If you have a hundred-page PDF, Claude can read it. If you have a long meeting transcript, Claude can summarize it. If you have a spreadsheet, Claude can analyze it.

Tokens. The basic unit Claude works in. A token is roughly three-quarters of a word in English. You do not need to count tokens yourself; the interface manages this for you. But knowing the rough conversion — 200,000 tokens $\approx$ 500 pages — helps you judge what fits in a conversation. Most everyday tasks use a small fraction of the window.

Claude in full use.

Imagine a senior staff member at your organization who can read any document you give them in seconds, draft any kind of writing in your voice once you have shown them what your voice is, translate fluently between English and most major languages, research topics by searching the web and reading the sources, write code to automate routine tasks, build a small interactive tool you can hand to a colleague, fill out a form, summarize a long meeting, propose an analysis of your donor data, critique its own work honestly when you ask it to, and pick up exactly where you left off any time you bring it back to a project. Now imagine that staff member does not get tired, costs almost nothing, and is available at nine on a Sunday night when you need to draft something before Monday morning. That is roughly what Claude is, in 2026, when you use it fully. The Sandbox is where your team learns what that means for your actual work.

WHERE THIS COULD GO

What this looks like across your work.

The Sandbox tests AI use cases against your specific work. Below is a representative catalog of the kinds of use cases that surface in mission-driven organizations — eight standard functions, ten sample use cases per function. These are not prompts. They are descriptions of the kinds of work AI can plausibly assist with at the current state of the technology. The Sandbox is where you find out which of these — and which others, specific to your context — actually produce value in your organization.

Use cases here assume Claude’s full capabilities: Artifacts, Projects, Connectors, Skills, web search, code execution, file analysis. Not just chat. Each one would pass through the eight phases of the Sandbox before any operational deployment.

FUNCTION 01

Executive leadership and strategy.

- 01 Annual strategic plan first draft synthesized from current-state assessment, prior plans, board direction, and program outcome data.
- 02 Board meeting prep brief that synthesizes recent program data, financial position, and policy history into a senior-leader-readable document.
- 03 Strategic memo drafting with reference to existing organizational documents, with citations back to source policies.
- 04 Senior leadership communication drafts — internal all-staff, donor newsletters, partner updates — with voice calibration against past examples.
- 05 Decision-map navigation that answers “who decides this, and against which policy” for any organizational question.
- 06 Annual report narrative drafting from program data, with the communications team retaining final editorial.
- 07 Executive search candidate profile synthesis from public sources and stated criteria.
- 08 Public statement drafting on emerging issues, checked against the organization’s theological or ethical commitments.
- 09 Sector landscape analysis: what comparable organizations are doing on a given strategic question.
- 10 Annual planning retreat materials preparation, including pre-reads, framing documents, and discussion guides.

FUNCTION 02

Development and fundraising.

- 01 Major donor profile research synthesis from public sources only, with explicit privacy boundaries.
- 02 Grant proposal first draft generated from program outcome data and prior winning proposals.
- 03 Foundation prospect research that summarizes funding history, stated priorities, and recent grants.
- 04 Donor segmentation analysis from CRM data with insights about patterns and outliers.
- 05 Personalized stewardship letter drafts with relationship context and explicit human-review-required workflow.
- 06 Campaign messaging variations for A/B testing across appeal channels.

- 07 Acknowledgment letter pipeline with strict human-review gate before any send.
- 08 Donor survey response clustering and theme identification from open-text answers.
- 09 Year-over-year fundraising analysis with insights about what changed and why.
- 10 Stewardship communication calendar planning across the donor lifecycle.

FUNCTION 03

Programs and direct service.

- 01 Beneficiary intake summarization (sanitized for AI use; never raw sensitive data) for caseload visibility.
- 02 Program manual maintenance and version updates as policies and practices evolve.
- 03 Cultural and contextual review of program materials before they reach community participants.
- 04 Translation of program documents into the languages your community actually speaks, with human translator review.
- 05 Outcome measurement framework drafting against your theory of change.
- 06 Case study writing from program data, with consent-aware redaction.
- 07 Curriculum adaptation for different program contexts (in-person, online, urban, rural).
- 08 Plain-language reformatting of dense materials for accessibility.
- 09 Program evaluation report drafting that consolidates data, narratives, and reflections.
- 10 Volunteer training material development tailored to specific program roles.

FUNCTION 04

Communications and marketing.

- 01 Voice-calibrated style-checker that reads any outbound communication and flags drift from your organizational voice.
- 02 Social media content variations adapted to each platform from a single source idea.
- 03 Newsletter drafting from your organizational activity log.
- 04 Press release first drafts with quotes shaped to your spokespeople's actual voices.
- 05 Website copy refreshes that maintain voice consistency across pages.
- 06 Email campaign A/B test variation development.
- 07 Brand guideline application checking across any new content.
- 08 Crisis communication template drafting before the crisis (not during).
- 09 Photo and image alt-text generation for accessibility, reviewed for accuracy.
- 10 Constituent FAQ generation and maintenance from incoming questions.

FUNCTION 05

Operations and administration.

- 01 Process documentation generated from staff interviews about how work actually happens.
- 02 Meeting minute synthesis with action item extraction across recurring meetings.

- 03 Internal decision-map assistant that answers “who decides what” from documented authority structures.
- 04 Vendor evaluation against documented organizational criteria.
- 05 Project management status synthesis across teams for senior visibility.
- 06 Office policy and procedure updates with consistency checking.
- 07 Compliance checklist maintenance and tracking.
- 08 Internal knowledge base maintenance with content freshness checks.
- 09 Standard operating procedure drafting from observed workflows.
- 10 Cross-departmental coordination summaries that reduce duplicative meetings.

FUNCTION 06

Finance and accounting.

- 01 Budget narrative writing for grant proposals from financial data and program plans.
- 02 Audit preparation document organization and gap-finding.
- 03 Financial report executive summary drafting from underlying statements.
- 04 Expense categorization assistance with explicit human-review-required workflow.
- 05 Board financial committee briefing preparation with context and trend analysis.
- 06 Grant compliance reporting drafts tied to specific funder requirements.
- 07 Policy interpretation memos that reference organizational financial policy.
- 08 Donor acknowledgment wording for tax purposes that meets IRS requirements.
- 09 Financial dashboard narrative generation that explains what the numbers mean.
- 10 Year-end reporting first drafts that consolidate financials, programs, and outcomes.

FUNCTION 07

Human resources and volunteer management.

- 01 Job description drafting from role requirements and existing role examples.
- 02 Volunteer onboarding companion that answers questions from your handbook and FAQs.
- 03 Performance review draft language as starting points, with human authorship final and required.
- 04 Policy handbook updates with consistency checking across sections.
- 05 Recruitment messaging variations targeted to specific candidate audiences.
- 06 New hire onboarding curriculum development with role-specific tracks.
- 07 Internal communication drafts for sensitive announcements (drafts only — the actual announcement is human-authored).
- 08 Training material development for staff professional growth.
- 09 Exit interview synthesis with appropriate consent and redaction.
- 10 Compensation research from publicly available salary benchmarks.

FUNCTION 08*Board relations and governance.*

- 01 Board meeting minute synthesis with policy cross-reference for any decisions made.
 - 02 Bylaws and policy consistency checking as policies are updated.
 - 03 Committee report drafts from underlying program and financial data.
 - 04 Board orientation material maintenance for new directors.
 - 05 Governance question answering against bylaws, policies, and prior decisions.
 - 06 Annual board self-assessment synthesis from individual responses.
 - 07 Strategic alignment checks of proposed actions against your current plan.
 - 08 Conflict of interest declarations review and tracking.
 - 09 Board recruitment candidate research from public sources.
 - 10 Policy revision drafts with annotations showing what changed and why.
-

Eighty use cases. Few organizations should attempt all of them. Most organizations should Sandbox three to seven that match their actual friction points, learn what AI can and can't do in their context, and graduate the ones that earn their place. The catalog exists to show the surface area, not to prescribe the work. The Sandbox is where the work actually gets chosen.

PROCESS

How this work gets done.

The Sandbox is a discovery process. Eight phases, each with a clear purpose, each producing a documented output that feeds the next. The phases are sequential. You cannot run Experimenting before Boundaries; you cannot run Discerning before Ethics Review. The order is the methodology.

Underneath the phases, the work is built around three categories of value the Sandbox is looking for. Everything you try in the Sandbox is evaluated against these three. Nothing else.

01***Efficiency***

Time saved on work that already exists.

02***Revenue***

New revenue or sustained funding that AI helps unlock.

03***Quality***

Work made better — clearer, more accurate, more accessible.

Cohorts and teams.

Sandbox work runs in cohorts. A cohort is usually built around an existing team or function in your organization — the development team, the program staff, the pastoral team, the operations group. But the exact construction of cohorts is part of the flexible design of each engagement. We adapt to each organization, working with the team structures you already have rather than imposing new ones.

Whatever the configuration, everyone in the cohort is a learner. There is no AI expert in the room and there is no AI skeptic in the room. There is a group of people doing the work of finding out, together.

What the SandboxLive engagement provides.

SandboxLive is the facilitated version of the Sandbox. Roughly ten hours of in-person teaching per cohort, delivered by Movemental in your space or remotely, structured around the eight phases. The in-person time is accompanied by:

A learning management system that hosts the materials each cohort member uses between sessions — video modules, the recipe library, the assessment tools, the recording templates, and the dashboard surfaces for the Ethics Review and the Future Plan. The LMS is the cohort's permanent reference. Nothing has to be memorized during the in-person teaching time, because everything taught is also documented in the LMS and accessible whenever the cohort member needs it — during the engagement, after the engagement, and

across future Sandbox cycles.

A library of pre-tested AI recipes for churches, nonprofits, and theological institutions, customized to your specific context based on the Assessment phase. Each recipe is a documented AI use case that has been tested across organizations like yours, with the hypothesis, the workflow, the data tier, the failure modes, and the projected value already documented. Cohort members try the recipes against their actual work.

Customized resources and supports throughout the engagement, including templated experiment briefs adapted to your context, dashboard tools that integrate with your existing Guidebook, and ongoing access to Joshua Shepherd and the Movemental team for questions as they arise.

Scheduling is coordinated through Calendly to fit your existing team rhythms. The cohort meets when it makes sense for the cohort, not on a schedule imposed from outside.

What your team brings.

The Sandbox is yours to run. Movemental facilitates. Your team does the discovery. The work that is yours and cannot be delegated:

Identifying who is in the cohort. The people who will actually try AI against their work.

Engaging honestly during the Assessment phase — telling the truth about what your work actually is, what frustrates you about it, what you wish were different, what you would protect from AI involvement.

Trying the recipes. Recording the value. Iterating with curiosity. Reflecting honestly.

Inviting trusted voices into the Ethics Review when the time comes.

Making the call during Discerning — Green, Yellow, or Red for each use case.

Putting your name on the Future Plan.

THE DATA

What the research shows.

Three findings from the last twelve months frame the case for a disciplined Sandbox. They are independently sourced and they converge on the same pattern.

The pilot failure rate.

MIT's most-cited finding on enterprise AI in 2025. The 95 percent that failed did not fail because the technology cannot work. They failed because they tried to deploy without first discovering what would produce value in their specific context.

95%



of enterprise generative AI pilots fail to generate measurable financial returns.

MIT NANDA INITIATIVE · JULY 2025

60%



of companies generate no material value from AI investment.

BCG WIDENING AI VALUE GAP · SEPTEMBER 2025

6%



of organizations qualify as AI high performers in McKinsey's 2025 study.

McKINSEY STATE OF AI · NOVEMBER 2025

The nonprofit pattern.

The same pattern holds in the nonprofit sector, where 92 percent of organizations use AI in some form but only 7 percent report major capability improvements. The gap is the discovery discipline. Sandbox closes it.

92%



of nonprofits report using AI in some capacity.

VIRTUOUS · 2026 NONPROFIT AI ADOPTION REPORT

81%



describe their AI use as individual and ad hoc, without shared workflows.

VIRTUOUS · 2026

7%



report major improvements in organizational capability.

VIRTUOUS · 2026

THE ARCHITECTURE

The eight phases of the Sandbox.

The Sandbox is a sequence of eight phases. Each phase depends on the one before it. The process maps to a simple narrative: map current reality, build and try AI recipes, document potential value, discern ethical issues, decide.

- | | | |
|----|----------------------|--|
| 01 | Boundaries | What this Sandbox protects: publication, privacy, personal. |
| 02 | Assessment | Mapping the current work and context of each cohort member. |
| 03 | Experimenting | Trying Movemental recipes against real work; recording value. |
| 04 | Iteration | Refining what worked. Trying variations. Dropping what didn't. |
| 05 | Reflection | Stepping back. What we learned about AI, our work, ourselves. |
| 06 | Ethics Review | Opening the discoveries to outside voices for input. |
| 07 | Discerning | Green, Yellow, or Red for every use case the Sandbox surfaced. |
| 08 | Future Plan | What the organization commits to next. |

The rest of this guide walks each phase in order. Each phase chapter names the purpose of the phase, what a complete version produces, and what is at stake if the phase is skipped or done poorly.

PHASE 01

Boundaries.

What this Sandbox protects.

Before the cohort tries anything with AI, it agrees on three categories of boundaries that define what the Sandbox is and is not. The boundaries come from your AI Organizational Guidebook (Safety) and are adapted to the specific scope of this Sandbox engagement.

There are three categories. Publication, privacy, and personal.

Publication boundaries.

What may leave the Sandbox during exploration. The default is nothing. Sandbox outputs are internal. They are not posted to social media, sent to donors or constituents, included in newsletters, shared with partners, or used in any work product that reaches outside the cohort. The Sandbox is a closed environment until a use case passes through the Discerning phase and becomes approved for actual deployment.

A complete Publication clause contains:

- A default rule that nothing produced in the Sandbox is published outside the cohort
- A specific list of exceptions, if any (typically none in a first Sandbox)
- A clear consequence for crossing the boundary
- Sign-off from each cohort member

Privacy boundaries.

What data may enter the Sandbox. Drawn directly from the Data Handling Rules in your AI Organizational Guidebook. Tier 1 (public) and Tier 2 (internal operational) data are typically permitted in Sandbox experiments. Tier 3 (sensitive) requires additional protocols or synthetic substitutes. Tier 4 (confidential) is excluded entirely. Specific examples: real donor names do not enter the Sandbox; synthetic donor data or anonymized fragments may. Real pastoral notes do not enter; sanitized teaching cases may.

A complete Privacy clause contains:

- A reference to the Data Classification in your AI Organizational Guidebook
- A specific list of data types permitted in this Sandbox
- A specific list of data types excluded from this Sandbox
- Named workflows for using synthetic substitutes when sensitive categories would otherwise be needed
- Sign-off from each cohort member

Personal boundaries.

The category most often missed. Prolonged conversational AI work affects the person doing it in ways that are different from other knowledge work. The cohort should name what they are watching for and what safeguards they will hold.

Three patterns to watch for. **Cognitive fatigue** — hours of back-and-forth with a language model can produce a different kind of attention exhaustion than other tasks. **Parasocial attachment** — regular AI use can produce a sense of working with someone, which warrants honest reflection. **Reality recalibration** — prolonged immersion in AI-generated text can subtly shift the cohort member's sense of what good writing, sound reasoning, or true insight looks like. None of these is a reason not to use AI. All of them are reasons to use it with awareness.

A complete Personal clause contains:

- Time-boxed Sandbox sessions (no all-day AI immersion without breaks)
- In-person cohort debrief between sessions, not chat-only check-ins
- Permission for any cohort member to name when AI use is leaving them feeling differently than they expect
- Permission to stop and step back from a Sandbox session if needed
- A culture that treats those decisions as wisdom, not weakness

What is at stake without it: the Sandbox stops being a Sandbox. Without publication boundaries, experiments leak into real work before they have been adjudicated. Without privacy boundaries, sensitive data flows into tools that do not protect it. Without personal boundaries, the people doing the discovery work are subtly harmed by it.

PHASE 02

Assessment.

Mapping the current work and context of each cohort member.

Before AI is brought to bear on anyone's work, the cohort maps the current reality of that work. Each cohort member completes a structured assessment that documents what they actually do, how their time is spent, where the friction lives, what they wish were different, and what they would protect from AI involvement.

The assessment is itself AI-assisted. An interview module in the learning management system walks each cohort member through structured questions, and a language model helps generate a first draft of their Current Reality Map. The cohort member reviews and revises the draft. The Map is then published to the cohort dashboard for shared visibility.

This is intentional. The Assessment phase is the cohort member's first structured encounter with AI in the Sandbox. The Map they end up with is itself an artifact of AI-assisted work — paired with clear human input and a clear definition of done. It is the first object lesson in what AI can do well when those conditions are met.

A complete Assessment contains, per cohort member:

- Work categories with rough time allocation across a typical week
- Friction points — the categories of work that take longer than they should
- Quality concerns — the work that is good enough but not as good as it could be
- Stewardship concerns — the work that feels wasteful or low-leverage
- Categories the cohort member would protect from AI involvement, and why
- Categories the cohort member is curious to test AI against
- Sign-off by the cohort member, with the Map made visible to the cohort

What is at stake without it: experiments target the wrong work. The cohort tries flashy AI use cases that do not map to actual organizational pain. Or worse, the cohort tries use cases that an individual member privately considers sacred work that should not be automated, and discovers that view only after harm has been done. The Map prevents both failure modes.

PHASE 03

Experimenting.

Trying Movemental recipes against real work; recording value.

This is the heart of the Sandbox. The cohort tries AI use cases against their actual work and records the value those use cases generate. The phase has two activities that work together: Recipes and Recording.

Recipes.

Movemental maintains a library of pre-tested AI use cases built for churches, nonprofits, and theological institutions. Each recipe is a documented use case with a hypothesis, the AI workflow, the tools required, the data tier permitted, the known failure modes, and the expected categories of value. The recipes are tested across many organizations like yours before they reach your cohort.

In a SandboxLive engagement, the recipes are customized to your context based on the Assessment phase outputs. Each cohort member is given a starting set of recipes to try against the friction points and curiosity areas they named in their Map. The cohort tries the recipes during in-person teaching modules and asynchronously through the LMS.

The recipes do double work. They give the cohort member a structured use case to try, which would have taken them weeks to design from scratch. And they teach the cohort member about AI in general, because working through a guided example is the fastest way to internalize how AI actually behaves. By the end of the recipes a cohort member has tried, they understand AI more deeply than any reading or video could have produced.

Recording.

After each recipe is tried, the cohort member records the projected value of the use case. Three categories only:

Time saved. Specifically: how long did this work take with AI, and how long would it have taken without? The difference, projected across a typical month or year.

Revenue generated or sustained. Specifically: is there a plausible new revenue source this use case opens up, or a current revenue source it helps sustain? Honest projections only — not vanity claims.

Work quality improved. Specifically: did the work get better in a way someone outside the cohort would notice? Better in what way — clearer, more accurate, more accessible, more consistent?

An AI assistant built into the dashboard helps with the recording task. The cohort member describes what they did, the AI assistant produces a structured Value Record, and the cohort member reviews and revises. The Value Record is the documentation that travels with the use case through every subsequent phase.

Notes on AI use technique inside the recipes.

Every recipe is built around the same three skills named earlier. Clear communication (saying what you actually want, with the audience and the format named). Definition of done (knowing what a finished version of the thing looks like). Agentic perspective-taking (remembering what the model is, what it knows, what it doesn't).

Cohort members are also reminded that it is okay to play. The recipes are designed for experimentation. Mistakes during the Sandbox are not failures; they are data. A recipe that produces a strange or wrong output the first time is showing the cohort member something useful about how the model behaves with that kind of input. Iterate, learn, move on.

How we design experiments inside the recipes.

Movemental brings a specific experimental discipline to the recipes. Every recipe a cohort member tries comes with a small set of design features that distinguish it from casual experimentation. The same features apply when your cohort designs its own recipes in the self-directed path.

Hypothesis. Every recipe states, in advance, what the experiment expects to learn. Not a vague “let’s see what happens” but a specific claim: *This recipe will reduce the time required to draft a major-donor thank-you letter from approximately forty-five minutes to under ten minutes, while preserving the personal voice of the development director.* That hypothesis is testable. Either it holds or it doesn't.

Time baseline. Before AI touches the work, the cohort member documents how long the task currently takes. Without a baseline, time-savings claims are guesses. With a baseline, they are assessments.

Quality baseline. What does the work currently look like at the organization’s best? A sample of recent human-only output is documented for comparison. AI-assisted output is then compared against that sample, not against an imagined gold standard.

Success criteria. What would tell us the recipe is working? Usually specific: time reduced by a stated percentage; quality matches or exceeds baseline as judged by a named reviewer; revenue plausibility supported by a concrete projection.

Failure modes anticipated in advance. What could go wrong? The recipe names the most likely failure modes before the experiment runs. “The AI may produce a letter that feels too generic.” “The AI may invent facts about the donor.” “The AI may use language that’s technically correct but theologically off-key.” Naming these in advance helps the cohort member watch for them rather than be surprised by them.

Multiple trials. One recipe attempt is not an experiment. Each recipe is run at least three times, on three different real instances of the work, so the cohort member sees variation. Three letters drafted, not one. Three social posts, not one. The pattern across the trials is what produces real evidence.

A/B comparison where it makes sense. When the work can be split, the cohort member runs parallel cases — some with AI assistance, some without — and compares the outcomes honestly. Not every recipe permits this, but where it does, A/B is the cleanest form of evidence.

Cross-cohort validation. A recipe that works for one cohort member is tried by another to see whether the value generalizes. A recipe that works only for the most AI-skilled cohort member is not yet a recipe; it is a personal workflow.

These design features are built into every Movemental recipe in the SandboxLive library. In the SandboxGuide path, the cohort designs its own recipes and applies these same features as it goes — which is most of why the SandboxGuide path takes longer. The discipline is the work.

A complete Experimenting phase output contains:

- A list of recipes tried by each cohort member
- A Value Record per recipe — time, revenue, work quality, with honest projections
- Notes on what worked, what didn't, what surprised the cohort member
- A draft list of use cases the cohort member believes are worth pursuing further
- Sign-off per cohort member, with the records visible to the cohort

What is at stake without it: without recipes, cohort members are asked to invent their own experiments from scratch, which they cannot do because they do not yet know AI well enough to design good experiments. Without recording, the experiments produce vibes-based opinions rather than assessed value. Either failure mode produces the 95 percent.

PHASE 04

Iteration.

Refining what worked. Trying variations. Dropping what didn't.

After the first round of recipes, the cohort iterates. This is the phase where first impressions get tested against second tries. Recipes that worked are refined and extended. Recipes that almost worked are modified and retried. Recipes that clearly didn't fit the cohort member's work are dropped.

Iteration is also where cohort members share with each other. A recipe one person tried successfully against one category of work might be exactly what another cohort member needs for a different category. The shared dashboard makes this visible. The in-person teaching time supports it.

A complete Iteration phase output adds to the Value Record:

- Variations tried per recipe (different inputs, different framings, different definitions of done)
- Updated value assessments after iteration
- Recipes the cohort decides to drop, with a brief reason
- Recipes the cohort decides to pursue further, with the path forward named
- Cross-pollination: recipes shared across cohort members and tried by new people

What is at stake without it: the cohort treats first impressions as final. AI use cases that look bad on first attempt often improve substantially with refinement. AI use cases that look great on first attempt sometimes do not hold up under repeated use. Iteration is where the cohort develops actual judgment rather than first reactions.

PHASE 05

Reflection.

Stepping back. What we learned about AI, our work, ourselves.

Before the cohort moves into the harder phases of ethics and decision, it pauses to consolidate what it has learned. Reflection is not ethics yet. It is learning consolidation. What surprised us? What confirmed what we already believed? What changed our minds? What tensions surfaced between cohort members?

The Reflection phase is structured around a cohort session, in person, facilitated, where each cohort member shares what they have learned and what they are still working out. The session is documented in a shared Reflection Output that becomes the input for the Ethics Review and Discerning phases that follow.

A complete Reflection output contains:

- What the cohort learned about AI's actual capabilities and limitations in their work
- What the cohort learned about its own work — patterns that became visible only when AI was applied to them
- What the cohort learned about itself — the people in the room, their reactions, their growth
- Themes that emerged across multiple cohort members
- Tensions that surfaced — places where cohort members disagree
- Surprises, good and bad

What is at stake without it: the experiments stay individual. The cohort does not develop the shared vocabulary it will need to communicate what it found to the rest of the organization. The Future Plan ends up reading like a list of disconnected use cases rather than a coherent account of what the Sandbox revealed.

PHASE 06

Ethics Review.

Opening the discoveries to outside voices for input.

After Reflection, the Sandbox surfaces to outside voices. The cohort dashboard opens to a selected group of reviewers — board members, key donors, senior leaders not in the cohort, ministry partners, and other trusted constituents the organization names. The reviewers see the use cases the Sandbox has surfaced, with descriptions readable by people who were not in the room, and they are invited to weigh in.

This phase is about weighing in, not deciding. The decision comes next, in the Discerning phase. Ethics Review is the phase where the cohort makes sure it has not become so familiar with the use cases under consideration that it has lost the perspective of people who have not spent the past month inside the Sandbox.

The reviewers are asked specific questions. Does this use case fit the kind of organization we are? Does it sit comfortably with the constituents the organization serves? Are there subtle risks the cohort might be missing? Would the reviewer be comfortable if the organization's public communication named this as something the organization does? Their input is documented in the dashboard for use in the next phase.

Why public commentary matters.

The people who have spent the most time with AI use cases in a Sandbox are also the people most likely to have grown comfortable with use cases that should disturb the broader organization. This is not a failure of the cohort. It is a feature of being deeply involved in any new practice. The Ethics Review phase exists to make sure the cohort's familiarity does not become a blind spot.

It is also a moment of organizational transparency. The reviewers see what the cohort has been doing. The act of inviting their input is itself a sign of organizational seriousness about AI — a signal that the work is not being done in private, that decisions are not being made without consultation, that the people closest to the organization have voice in what comes next.

A complete Ethics Review output contains:

- A list of use cases publicly named for review, with descriptions readable by non-cohort members
- A list of reviewers invited, and a record of who weighed in
- All comments received, organized by use case
- Themes in the feedback — where multiple reviewers raised the same concern
- Surprises — where the cohort expected concern and received encouragement, or vice versa
- Open questions that emerged from the review and need attention before Discerning

What is at stake without it: the cohort makes decisions in isolation. Use cases that would disturb major donors, congregants, or board members get approved without anyone outside the cohort knowing they were

being considered. The first time the broader organization learns about the use case is when it shows up in actual work, which is the worst time to surface a problem.

PHASE 07

Discerning.

Green, Yellow, or Red for every use case the Sandbox surfaced.

This is where decisions get made. Every use case the Sandbox surfaced — from the Experimenting phase, refined through Iteration, reviewed externally in Ethics Review — is adjudicated as Green Light, Yellow Light, or Red Light. The cohort, with the senior leader who sponsored the Sandbox, makes the call together.

Each adjudication weighs five inputs. The Value Records from Experimenting (time, revenue, quality). The boundary alignment from Phase 01 (does the use case fit the Sandbox's publication, privacy, and personal boundaries?). The Guidebook alignment from Safety (does the use case fit your AI Organizational Policy and Rules?). The Ethics Review input (what did outside voices say?). The Reflection insights (what does the cohort believe after sitting with the work?).

The three verdicts.

Green Light. The use case is approved for operational deployment. The cohort and the senior leader agree it produces value the organization wants, fits within the Guidebook without strain, and survives outside review. Greens move into Skills (where staff are formed around the use case) or directly into Solutions (where the use case is integrated into a real workflow) depending on readiness.

Yellow Light. The use case is conditionally approved. The cohort believes it has promise but it needs more work — specific modifications, additional safeguards, scope limitations, or a future re-review. Yellows do not deploy yet. The conditions and the path forward are named.

Red Light. The use case is refused. It either failed to produce real value, crossed a boundary that the cohort or the reviewers found significant, or surfaced a tension with the organization's identity that the cohort cannot resolve. Reds do not deploy. The reason for refusal is added to the Named Refusals section of your AI Organizational Guidebook so that the decision is durable beyond the cohort's memory.

A complete Discerning output contains:

- Every Sandbox use case adjudicated with a Green, Yellow, or Red verdict
- Reasoning per verdict, in writing
- For each Green: the deployment path (Skills or Solutions) and the conditions
- For each Yellow: the modifications required, the timeline, and the re-review date
- For each Red: the named reason and the proposed addition to the Guidebook's Named Refusals
- Sign-off from the senior leader who sponsored the Sandbox

What is at stake without it: the Sandbox erases itself. Successful experiments quietly become production workflows without review. Problematic experiments get forgotten until they cause harm. The whole purpose of the Sandbox — generating evidence the organization can act on — fails at the final step. Reds, in particular,

must be named and recorded; the act of refusing publicly is what makes future vendor pitches resistible.

PHASE 08

Future Plan.

What the organization commits to next.

The final phase is the board-facing output. The Future Plan is a short document — two to four pages — that summarizes what the Sandbox did, what it learned, what it commits to, and what it refuses. It is the document the cohort hands to the senior leadership and the board to make the case for what comes next on the Movemental Path.

A complete Future Plan contains:

- A summary of the Sandbox: cohort composition, scope, recipes tried, work tested
- A summary of what the cohort learned (the Reflection themes)
- The Green Light use cases the organization commits to deploying, with timeline
- The Yellow Light use cases that need more work, with the path forward
- The Red Light use cases the organization commits to publicly refusing, with reasoning
- Governance updates the AI Organizational Guidebook needs based on what was learned
- The case to the board for next steps: Skills, Solutions on specific use cases, or another Sandbox cycle
- Signatures from senior leadership

What is at stake without it: the Sandbox produces learning that does not translate to action. The board never sees what was tried. The next stage of the Movemental Path gets purchased rather than chosen. The work the cohort did is owned only by the cohort.

LEARNING DESIGN

Learning outcomes and course outline.

The Sandbox is a structured learning experience. Below are the ultimate learning outcomes of the engagement and a module-by-module course outline that maps to them. The outline is organized by module, not by week or hour. Your cohort moves through it at the pace the work requires, supported by a learning management system your team has permanent access to. Cohort members do not need to memorize anything; everything taught in person is also documented in the LMS and available to revisit whenever they need it — during the engagement, after the engagement, and across future Sandbox cycles.

Ultimate learning outcomes.

By the end of the Sandbox, every cohort member can:

- Engage Claude (or any frontier AI assistant) as a working partner using basic communication skills, with practical understanding of what the tool is and how it behaves.
- Identify, design, and execute a small experiment that tests an AI use case against their actual work, with hypothesis, baselines, success criteria, anticipated failure modes, and recorded value across time, revenue, and quality.
- Iterate intelligently rather than treating first results as final.
- Surface ethical considerations and weigh them honestly against documented value.
- Contribute to an organizational decision about which AI use cases the organization adopts, modifies, or refuses — with Green, Yellow, and Red Light reasoning the board can ratify.
- Communicate the Sandbox's findings to leadership, board, and constituents in language they can act on.

Sample course outline.

Eleven modules. Each is delivered through a combination of in-person teaching time and asynchronous LMS work. The in-person teaching is roughly ten hours total across the engagement; the LMS work expands the in-person teaching with video modules, examples, templates, recipe library access, and dashboard tools. The order below is the default, but the structure is adapted to your cohort's context.

MODULE 01

Understanding your working partner.

What Claude is, how it was trained, what it can do beyond chat, and the genie metaphor in practice.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Articulate, at a basic level, what Claude is and what its training cutoff means for the work.
- Distinguish Claude's core chat capability from Artifacts, Projects, Connectors, Skills, web search, and code execution.
- Manage the context window deliberately — knowing when to start fresh and when to attach files.
- Upload and use files within token limits without anxiety about hitting them.

MODULE 02

How to talk to it.

The five core AI interaction skills: clear communication, context, definition of done, iteration, agentic perspective-taking.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Write a request that names audience, format, constraints, and definition of done.
- Provide useful context without over-engineering it.
- Iterate productively rather than abandoning after the first attempt.
- Take agentic perspective on the model — imagining what it is like to be the model on the other side of the conversation.
- Hold technique and wisdom together: pretending the model is a person while remembering it is not.

MODULE 03

Sandbox boundaries.

The three categories of boundaries that define what the Sandbox protects: publication, privacy, personal.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Apply the three boundaries to a real experiment in the cohort member's work.
- Connect Sandbox privacy boundaries to the Data Handling Rules in the AI Organizational Guidebook.
- Recognize the personal patterns to watch for: cognitive fatigue, parasocial attachment, reality recalibration.
- Hold the cohort's permission to stop and step back.

MODULE 04

Mapping your work.

The Assessment phase. Producing a Current Reality Map of the cohort member's actual work using an AI-assisted interview process.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Produce an accurate Current Reality Map of own work, including time allocation and friction points.
- Identify high-leverage work categories where AI could plausibly help.
- Articulate work categories the cohort member would protect from AI involvement, and why.
- Experience the Assessment process as the first structured encounter with AI in the Sandbox — an object lesson in AI-assisted work.

MODULE 05

Recipes.

Trying Movemental-built AI recipes against the cohort member's actual work. Recipes from the library are customized to context based on the Assessment.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Try at least three recipes from the library against real work.
- Recognize when a recipe fits the cohort member's context and when it doesn't.
- Adapt a recipe to a slightly different category of work than it was originally written for.
- Begin developing the personal sense of when AI is the right tool and when it isn't.

MODULE 06

Recording.

Documenting projected value across the three axes: time saved, revenue generated, work quality improved. The AI assistant in the dashboard helps produce the Value Record.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Establish a time baseline before AI touches a task, so time-savings claims are assessments rather than guesses.
- Establish a quality baseline from recent human-only work, for honest comparison.
- Measure projected time savings with appropriate confidence.
- Identify plausible revenue impacts where they exist, and decline to claim them where they don't.
- Articulate quality improvements with specifics, not adjectives.

MODULE 07

Iteration.

Refining what works, dropping what doesn't, sharing across the cohort.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Refine a recipe based on results from multiple trials.
- Test variations systematically rather than randomly.
- Drop a use case honestly when iteration shows it isn't earning its place.
- Cross-pollinate — try a recipe another cohort member developed and report back honestly.

MODULE 08

Reflection.

Stepping back. What we learned about AI, about our work, about ourselves. Not ethics yet — learning consolidation.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Articulate what was learned about AI's actual capabilities and limitations.
- Articulate what was learned about the cohort member's own work.
- Name tensions that surfaced — places where cohort members disagree.
- Contribute to shared cohort vocabulary that will travel into the rest of the organization.

MODULE 09

Ethics review.

Opening the Sandbox's discoveries to outside voices — board, donors, partners, senior leaders not in the cohort — for input before decisions are made.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Prepare use cases for outside review with descriptions readable by non-cohort members.
- Invite and integrate honest feedback from trusted constituents.
- Recognize the cohort's own familiarity bias and use the Ethics Review to correct for it.
- Document themes, surprises, and open questions for the Discerning phase.

MODULE 10

Discernment.

Green, Yellow, or Red Light for every use case the Sandbox surfaced.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Adjudicate a use case against documented criteria — value, boundary alignment, Guidebook alignment, Ethics Review input, Reflection insights.
- Articulate reasoning for each verdict in writing.
- Name Red Light refusals substantively, with reasoning robust enough to add to the organization's Named Refusals.
- Plan deployment paths for Green Lights and modification paths for Yellow Lights.

MODULE 11

Future plan.

Producing the board-facing output. Communicating Sandbox findings to the rest of the organization.

BY THE END OF THIS MODULE, COHORT MEMBERS CAN:

- Contribute to a board-ready Future Plan that names commitments, refusals, and next steps.
- Communicate Sandbox results to leadership, board, and constituents in language they can act on.
- Identify governance updates the AI Organizational Guidebook needs based on what was learned.
- Make the case for what comes next on the Movemental Path: Skills, Solutions on specific use cases, or another Sandbox cycle.

Modules can be combined, sequenced, or repeated to fit your cohort. The LMS structure follows the same eleven-module shape so that the in-person teaching and the asynchronous work track together. A cohort member who wants to revisit any module — during the engagement, after the engagement, or two years from now — can do so in the LMS at any time.

TAKE STOCK

Self-assessment.

Eight questions a senior leader can answer in twenty minutes. One question per phase. Each answer is **yes**, **partial**, or **no**. The pattern of answers tells you where you are honestly standing.

Take this assessment before committing to a Sandbox engagement, and again after each phase as the work progresses.

PHASE 01 — BOUNDARIES

Have you named, in writing, what your Sandbox protects — publication, privacy, and personal — with sign-off from each cohort member?

A strong answer is a signed Boundaries document drawn from your AI Organizational Guidebook. A weak answer is informal agreement. Without written boundaries the Sandbox is not a Sandbox; it is experimentation in disguise.

PHASE 02 — ASSESSMENT

Has each cohort member produced a Current Reality Map of their actual work, with friction points, quality concerns, and protected categories named?

A strong answer is a published Map per cohort member, visible to the cohort. A weak answer is verbal description. Without Maps the experiments target the wrong work.

PHASE 03 — EXPERIMENTING

Have cohort members tried structured AI recipes against their actual work and recorded value in time, revenue, and quality?

A strong answer is a Value Record per recipe, with honest projections. A weak answer is anecdote. The 95 percent who fail are the organizations that stop here without the records.

PHASE 04 — ITERATION

Has the cohort iterated — refined what worked, dropped what didn't, shared across members?

A strong answer is documented variations with updated Value Records. A weak answer is first impressions treated as final.

PHASE 05 — REFLECTION

Has the cohort sat together to consolidate what it learned — about AI, about its work, about itself?

A strong answer is a documented Reflection Output the cohort agrees on. A weak answer is individual takeaways that never become shared knowledge.

PHASE 06 — ETHICS REVIEW

Have you opened the Sandbox discoveries to outside voices — board, donors, partners, senior leaders not in the cohort — and documented their input?

A strong answer is a list of named reviewers, their comments, and the themes that emerged. A weak answer is the cohort making decisions in isolation.

PHASE 07 — DISCERNING

Has every use case the Sandbox surfaced been adjudicated as Green, Yellow, or Red, with reasoning and senior sign-off?

A strong answer is a complete verdict portfolio. A weak answer is some use cases adjudicated and others left in limbo. Use cases left in limbo become shadow production.

PHASE 08 — FUTURE PLAN

Have you produced a Future Plan your board has reviewed — naming what the organization will deploy, what it refuses, and what comes next?

A strong answer is a signed Future Plan with timeline. A weak answer is a Sandbox that ended without a handoff to the rest of the organization.

How to read your pattern.

No to three or more phases. Begin the Sandbox now. The exploration is already happening informally across your staff; the cost of formalizing it through the eight phases is much smaller than the cost of continuing without discipline.

Yes or partial on most phases. Finish what you started this quarter. The most common gap is Phase 07 — Discerning — because making the Green / Yellow / Red call is harder than the rest of the work. SandboxLive provides the facilitator support to hold the discipline at this step.

Yes on all eight. You have completed a Sandbox cycle. The question facing you is what comes next on the Movemental Path — Skills (formation around your Green Light use cases) or Solutions (operational deployment of specific Greens that are ready) — and whether another Sandbox cycle should run alongside it.

YOUR TWO PATHS

Two paths into Sandbox.

There are two paths through the Sandbox. Both produce the same eight-phase discovery and the same Future Plan. The difference is who facilitates the work, what the in-person experience is, and what supports are provided along the way.

How the two paths compare.

	<i>SandboxGuide</i> <i>Self-directed</i>	<i>SandboxLive</i> <i>Facilitated by Movemental</i>
Cost	Free.	Approximately \$15,000.
Teaching time	Self-paced through the Field Guide and your own design. Your team facilitates its own learning.	Roughly ten hours of in-person teaching per cohort, scheduled with your team through Calendly.
Difficulty	High. Requires that your team design its own recipes from scratch and hold the discipline of eight phases without outside facilitation.	Moderate. Movemental brings the recipes, the LMS, the dashboard, and the structure. Your team brings the work and the discernment.
Recipes	Your team designs its own AI use cases to try. The Field Guide provides the framework but no specific recipes.	A library of pre-tested AI use cases for churches, nonprofits, and theological institutions, customized to your context based on the Assessment phase.
Learning support	The Field Guide. Whatever supplementary resources your team finds on its own.	A learning management system with video modules, the recipe library, assessment tools, recording templates, and dashboard surfaces for Ethics Review and the Future Plan.
Finished artifact	Whatever format your team produces. Quality depends entirely on your team's capacity to hold the discipline.	A published Future Plan inside your private organizational dashboard, integrated with your Safety Guidebook, plus a print-quality PDF for board presentation.

Likely challenges	Designing recipes without prior AI experience. Holding cadence across eight phases. Making the Green / Yellow / Red call without external facilitation.	None at the Sandbox level. Some organizations underestimate the post-Sandbox work of operationalizing Green use cases — which the Skills stage addresses.
Best fit	Teams with AI experience already in-house, strong internal facilitation capacity, and discipline to finish what they start.	Teams that would benefit from a pre-built library of recipes, structured in-person teaching, dashboard-delivered tools, and external pacing.

What SandboxLive delivers.

When we say SandboxLive delivers the Sandbox better than the self-directed path in every measurable way, we mean specifically these things. Each one is included in the engagement fee. No upsell. No add-on tier.

Approximately ten hours of in-person teaching per cohort.

Joshua Shepherd and the Movemental team deliver structured teaching modules across the eight phases. The modules walk cohort members through the boundaries work, the assessment, the recipes, the recording practice, the iteration sessions, the reflection, the ethics review setup, the discerning session, and the future plan drafting. The teaching is in-person or by video conference depending on what fits your team.

A learning management system that holds everything between sessions.

Video modules cohort members can revisit. The recipe library customized to your context. Assessment tools for each cohort member to complete asynchronously. Recording templates with AI assistance built in. Reflection prompts. Ethics review surfaces. Future Plan drafting tools. Everything in one place, accessible from anywhere your cohort members work.

A customized library of AI recipes.

The single most valuable asset of SandboxLive. The recipe library is the work of Movemental across many engagements with churches, nonprofits, and theological institutions — every recipe pre-tested, documented, scored against an eight-dimension rubric, and customized to your specific context based on the Assessment phase outputs. Cohort members do not have to invent experiments from scratch. The recipes that match their friction points are waiting.

Dashboard integration with your Safety Guidebook.

The Boundaries phase of the Sandbox reads directly from your AI Organizational Guidebook. Your Data Handling Rules screen experiments before they launch. Your Named Refusals automatically flag any use case that would cross them. The Ethics Review and the Future Plan publish back to the same dashboard your Guidebook lives in. The Sandbox is integrated with the rest of the Movemental Path from day one.

Ongoing support throughout the engagement.

Access to Joshua and the Movemental team between teaching sessions for questions, troubleshooting, and discernment. Customized resources produced for your context as the engagement progresses. Templated experiment briefs, Value Records, Reflection Outputs, and Future Plan drafts adapted to your organization.

If you do this on your own.

If you read this Field Guide and run your Sandbox with your own team, we hope it serves you. The framework is genuinely sufficient if your team has the experience, the discipline, and the time. The risk of SandboxGuide is not that the framework is missing something. The risk is that designing the recipes from scratch — without a library to draw from — takes longer than most teams expect, and the discipline of finishing eight phases without external facilitation is harder than it looks.

SandboxLive exists for teams that would rather have the library, the teaching, and the supports in place from day one. The price is the price. There is no enterprise tier above it and no discount tier below it.

WHO WE ARE

About Movemental.

Movemental is the company that built the path that walks mission-driven organizations through AI governance, formation, and deployment without compromising the trust their work depends on. We are a small team of practitioners and engineers who put together the framework, the infrastructure, and the facilitated engagements that make the four-stage Movemental Path — Safety, Sandbox, Skills, Solutions — runnable inside an organization on a real Tuesday afternoon.

The four-stage Movemental Path.

Safety is the foundation. It produces your AI Organizational Guidebook — Statement, Policy, Context, Rules, and Response Plans — ratified by your board. The companion field guide is *It Starts With Safety*.

Sandbox is what this Field Guide is about. Roughly ten hours of in-person teaching per cohort across eight phases. The output is a Future Plan with Green, Yellow, and Red Light use cases discerned through real experimentation against real work.

Skills is the formation stage. Cohorts of leaders from comparable organizations form the capacities of discernment, authorship, and stewardship around AI. Not training on tools — formation of judgment that travels with the leaders into their work.

Solutions is where AI is integrated into how your organization actually runs, owned by formed people, governed by your working Guidebook. Scoped per engagement, ranging from tool optimization through network-scale deployments for institutions and denominations.

How to start a conversation.

If you read the Field Guide and want to talk about your organization, email Joshua Shepherd at josh@movemental.ai or visit movemental.ai/contact. We respond personally to every inbound message.

If you decide to run the Sandbox yourselves, that is a good outcome. The Field Guide is freely shareable in its entirety with attribution. Passing it to a peer at another organization is welcomed.

SOURCES

Citations.

Every factual claim in this document is sourced. Where possible the source is the primary publication; where the primary source is paywalled or restricted, an authoritative secondary citation is provided.

- [1] Movemental. *It Starts With Safety: A Field Guide for Building Your AI Organizational Guidebook*. Edition 1.1, 2026. The companion field guide that establishes the AI Organizational Guidebook this Sandbox depends on.
- [2] Movemental. *AI Has Already Arrived: A Documentation of the AI Reality Facing Churches, Nonprofits, and Institutions in 2026*. Field paper, May 2026. The grounding research for the urgency of AI governance and discovery work in mission-driven organizations.
- [3] Movemental. *The Movemental Stack, Plainly: A Framework, A Rubric, and 20 Nonprofit Use Cases*. Operating document. The eight-dimension rubric used in SandboxLive's Discerning phase and the pre-tested use case library that informs the recipe customization.
- [4] MIT NANDA Initiative. *The GenAI Divide: State of AI in Business 2025*. July 2025. The widely-cited finding that 95 percent of enterprise generative AI pilots fail to generate measurable financial returns. Boston Consulting Group. *The Widening AI Value Gap — Build for the Future 2025*. September 2025. Global study of 1,250 firms; 60 percent generate no material value from AI investment. McKinsey and Company. *The State of AI in 2025: Agents, Innovation, and Transformation*. November 2025. 1,993 respondents across 105 nations; 88 percent organizational adoption, 6 percent qualify as AI high performers.
- [5] Virtuous and Fundraising.AI. *The 2026 Nonprofit AI Adoption Report*. Released February 16, 2026. Benchmark survey of 346 nonprofit organizations. 92 percent of nonprofits use AI in some capacity; 7 percent report major improvements in organizational capability; 81 percent report using AI individually and on an ad hoc basis.
- [6] TechSoup. *2025 State of AI in Nonprofits Benchmark Report*. Documents the formal AI strategy gap across the nonprofit sector and supports the case for disciplined discovery work.
- [7] Forvis Mazars. *AI Governance for Nonprofit Boards*. By Karen Bodach and Dan Prater. February 13, 2026. Frames AI governance as a fiduciary expansion of the board's duties of Care, Loyalty, and Obedience — the same fiduciary frame that applies to the Sandbox stage of the Movemental Path.
- [8] Movemental. *The AI Stewardship Sequence — Field Guide for Organizational Leaders*. Canon document for the Safety, Sandbox, Skills, Solutions sequence. Describes the costs of inverting the sequence and the diagnostics for naming which stage an organization is honestly in.
- [9] Movemental. *Sandbox Discovery: Where Learning Actually Happens*. Methodology document. The supersession of the earlier Discovery Lab framing into the Sandbox stage of the Movemental Path.
- [10] Movemental. *A Guide to AI and Credibility in 2026*. Companion guide on the posture, boundaries, and practices for living with AI in a way that protects credibility. The principles informing the Personal boundaries clause in Phase 01 and the Reflection phase.

COLOPHON

It Continues With Exploration.

A Field Guide for the AI Sandbox.

Edition 1.0 · Movemental · 2026
Brad Brisco · Alan Hirsch · Joshua Shepherd

© 2026 Movemental. This document may be freely shared in its entirety with attribution. For permission to excerpt, contact josh@movemental.ai.

movemental.ai